
Lecture 7

CUDA Programming 2

张奇
复旦大学

Outline

CUDA Memory hierarchy

Efficient Shared Memory Use

GPU Computing with CUDA

Lecture 2 - CUDA Memories

*Christopher Cooper
Boston University*

*August, 2011
UTFSM, Valparaíso, Chile*

Recap

- ▶ GPU
 - A real alternative in scientific computing
 - Massively parallel commodity hardware
 - ▶ Market driven development
 - ▶ Cheap!
- ▶ CUDA
 - NVIDIA technology capable of programming massively parallel processors with general purpose
 - C/C++ with extensions

Recap

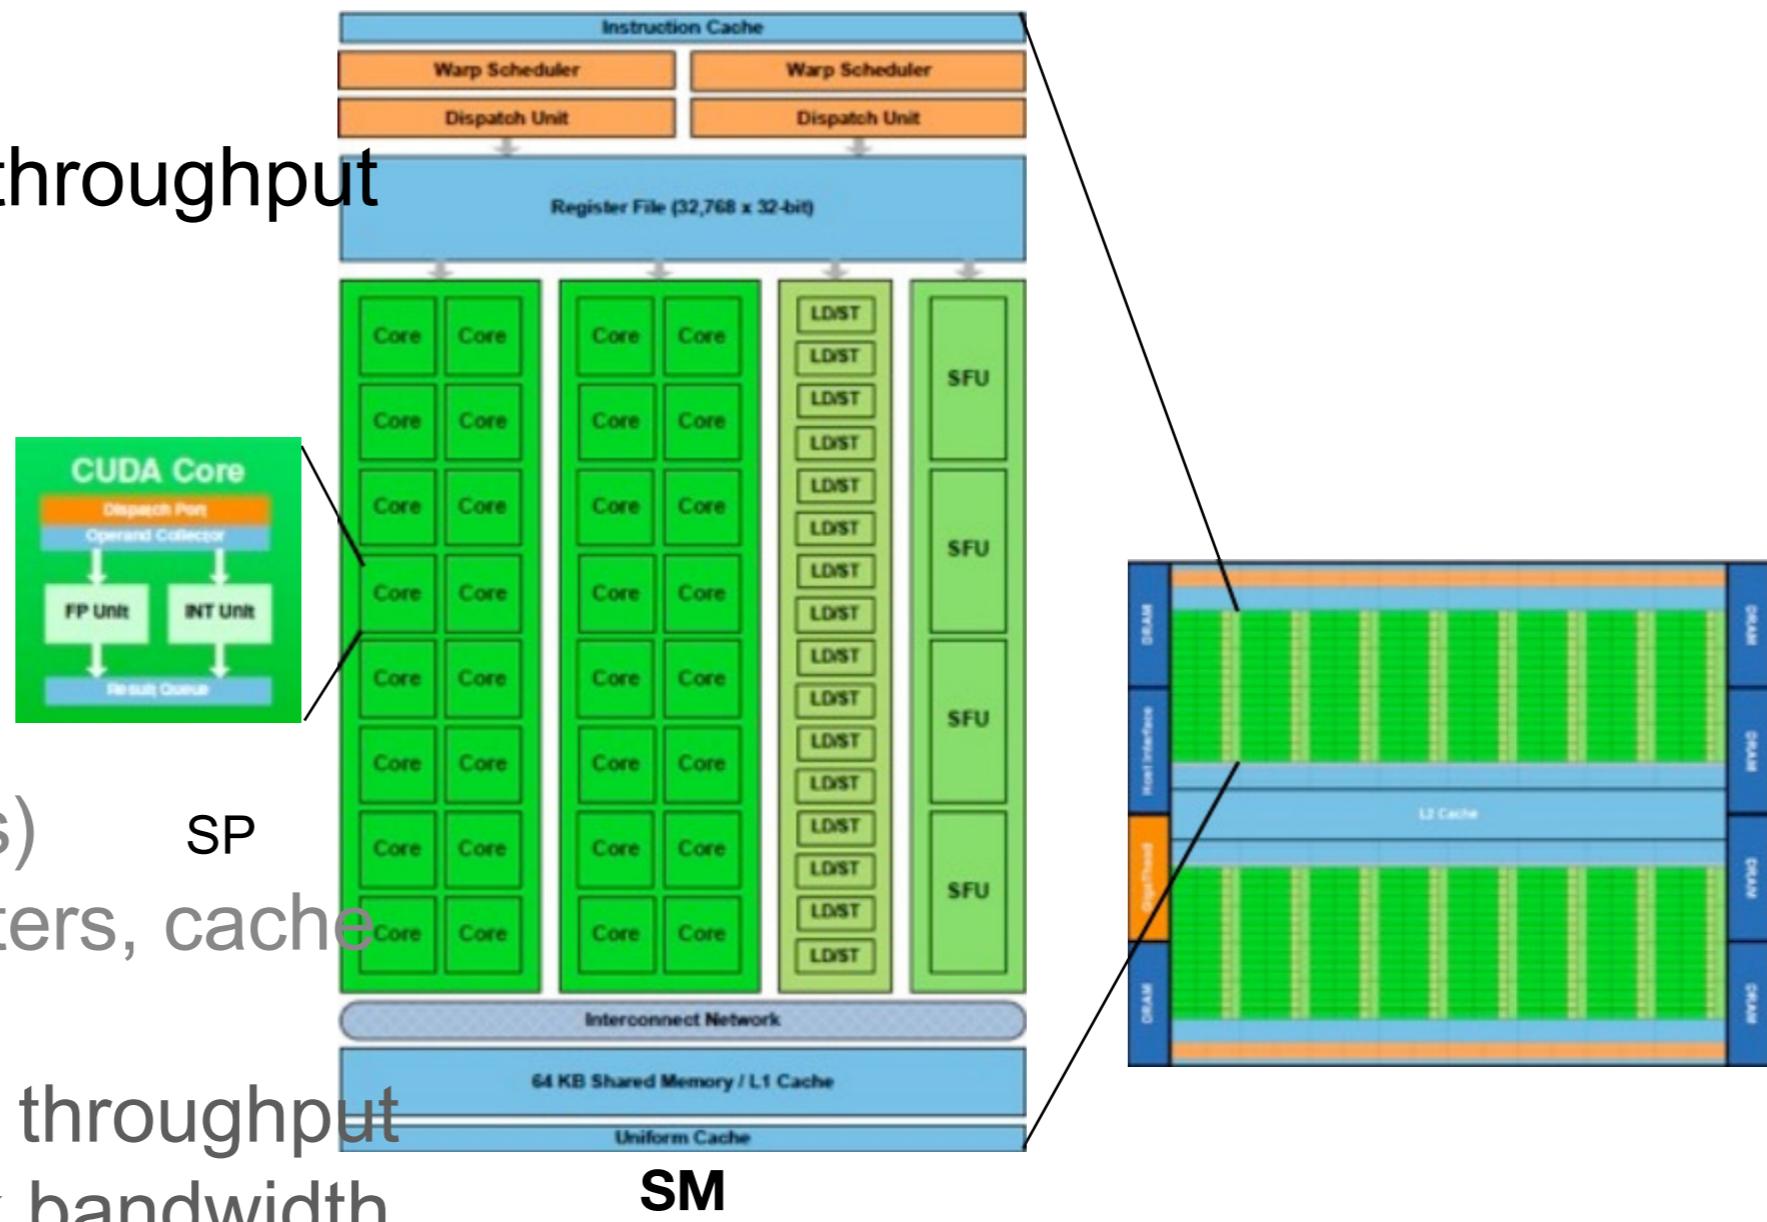
► GPU chip

- Designed for high throughput tasks

- Many simple cores

- Fermi:

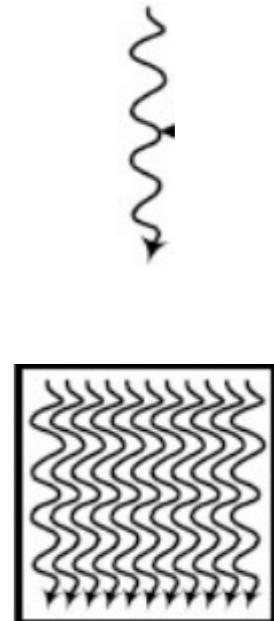
- ▶ 16 SMs
 - 32 SPs (cores)
 - Shared, registers, cache
 - SFU, CU
- ▶ 1TFLOP peak throughput
- ▶ 144GB/s peak bandwidth



Recap

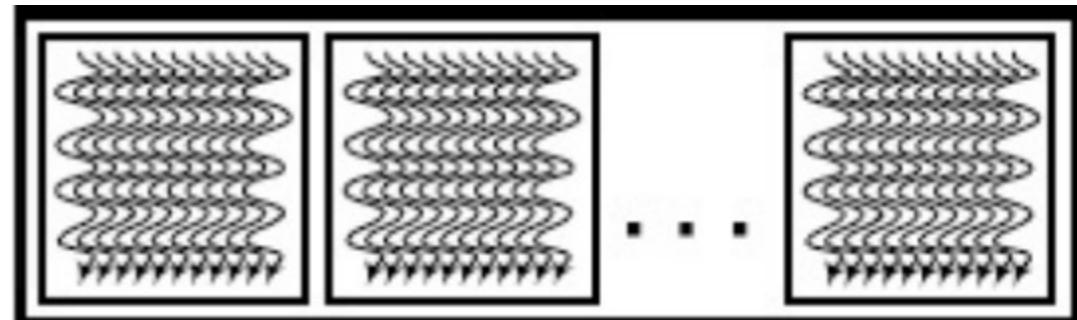
- ▶ Programming model

- Based on threads
- Thread hierarchy: grouped in thread blocks
- Threads in a thread block can communicate



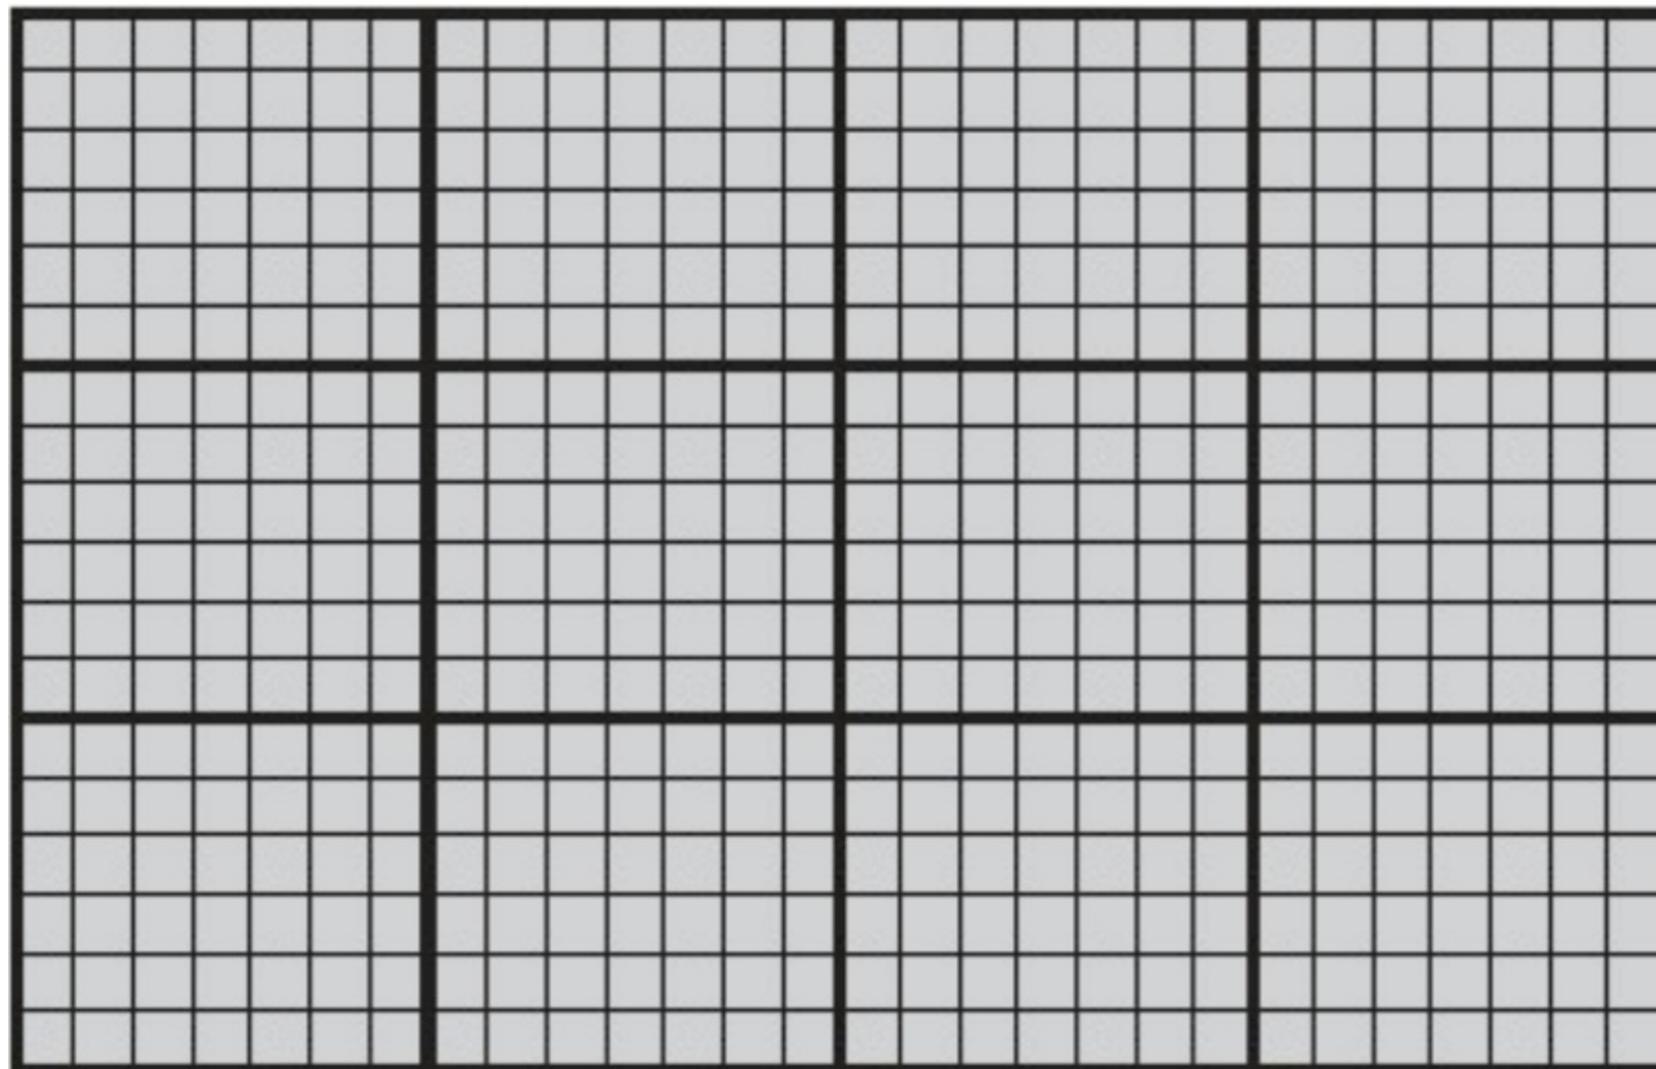
- ▶ Challenge: parallel thinking

- Data parallelism
- Load balancing
- Conflicts
- Latency hiding



CUDA - Programming model

- ▶ Connecting the hardware and the software



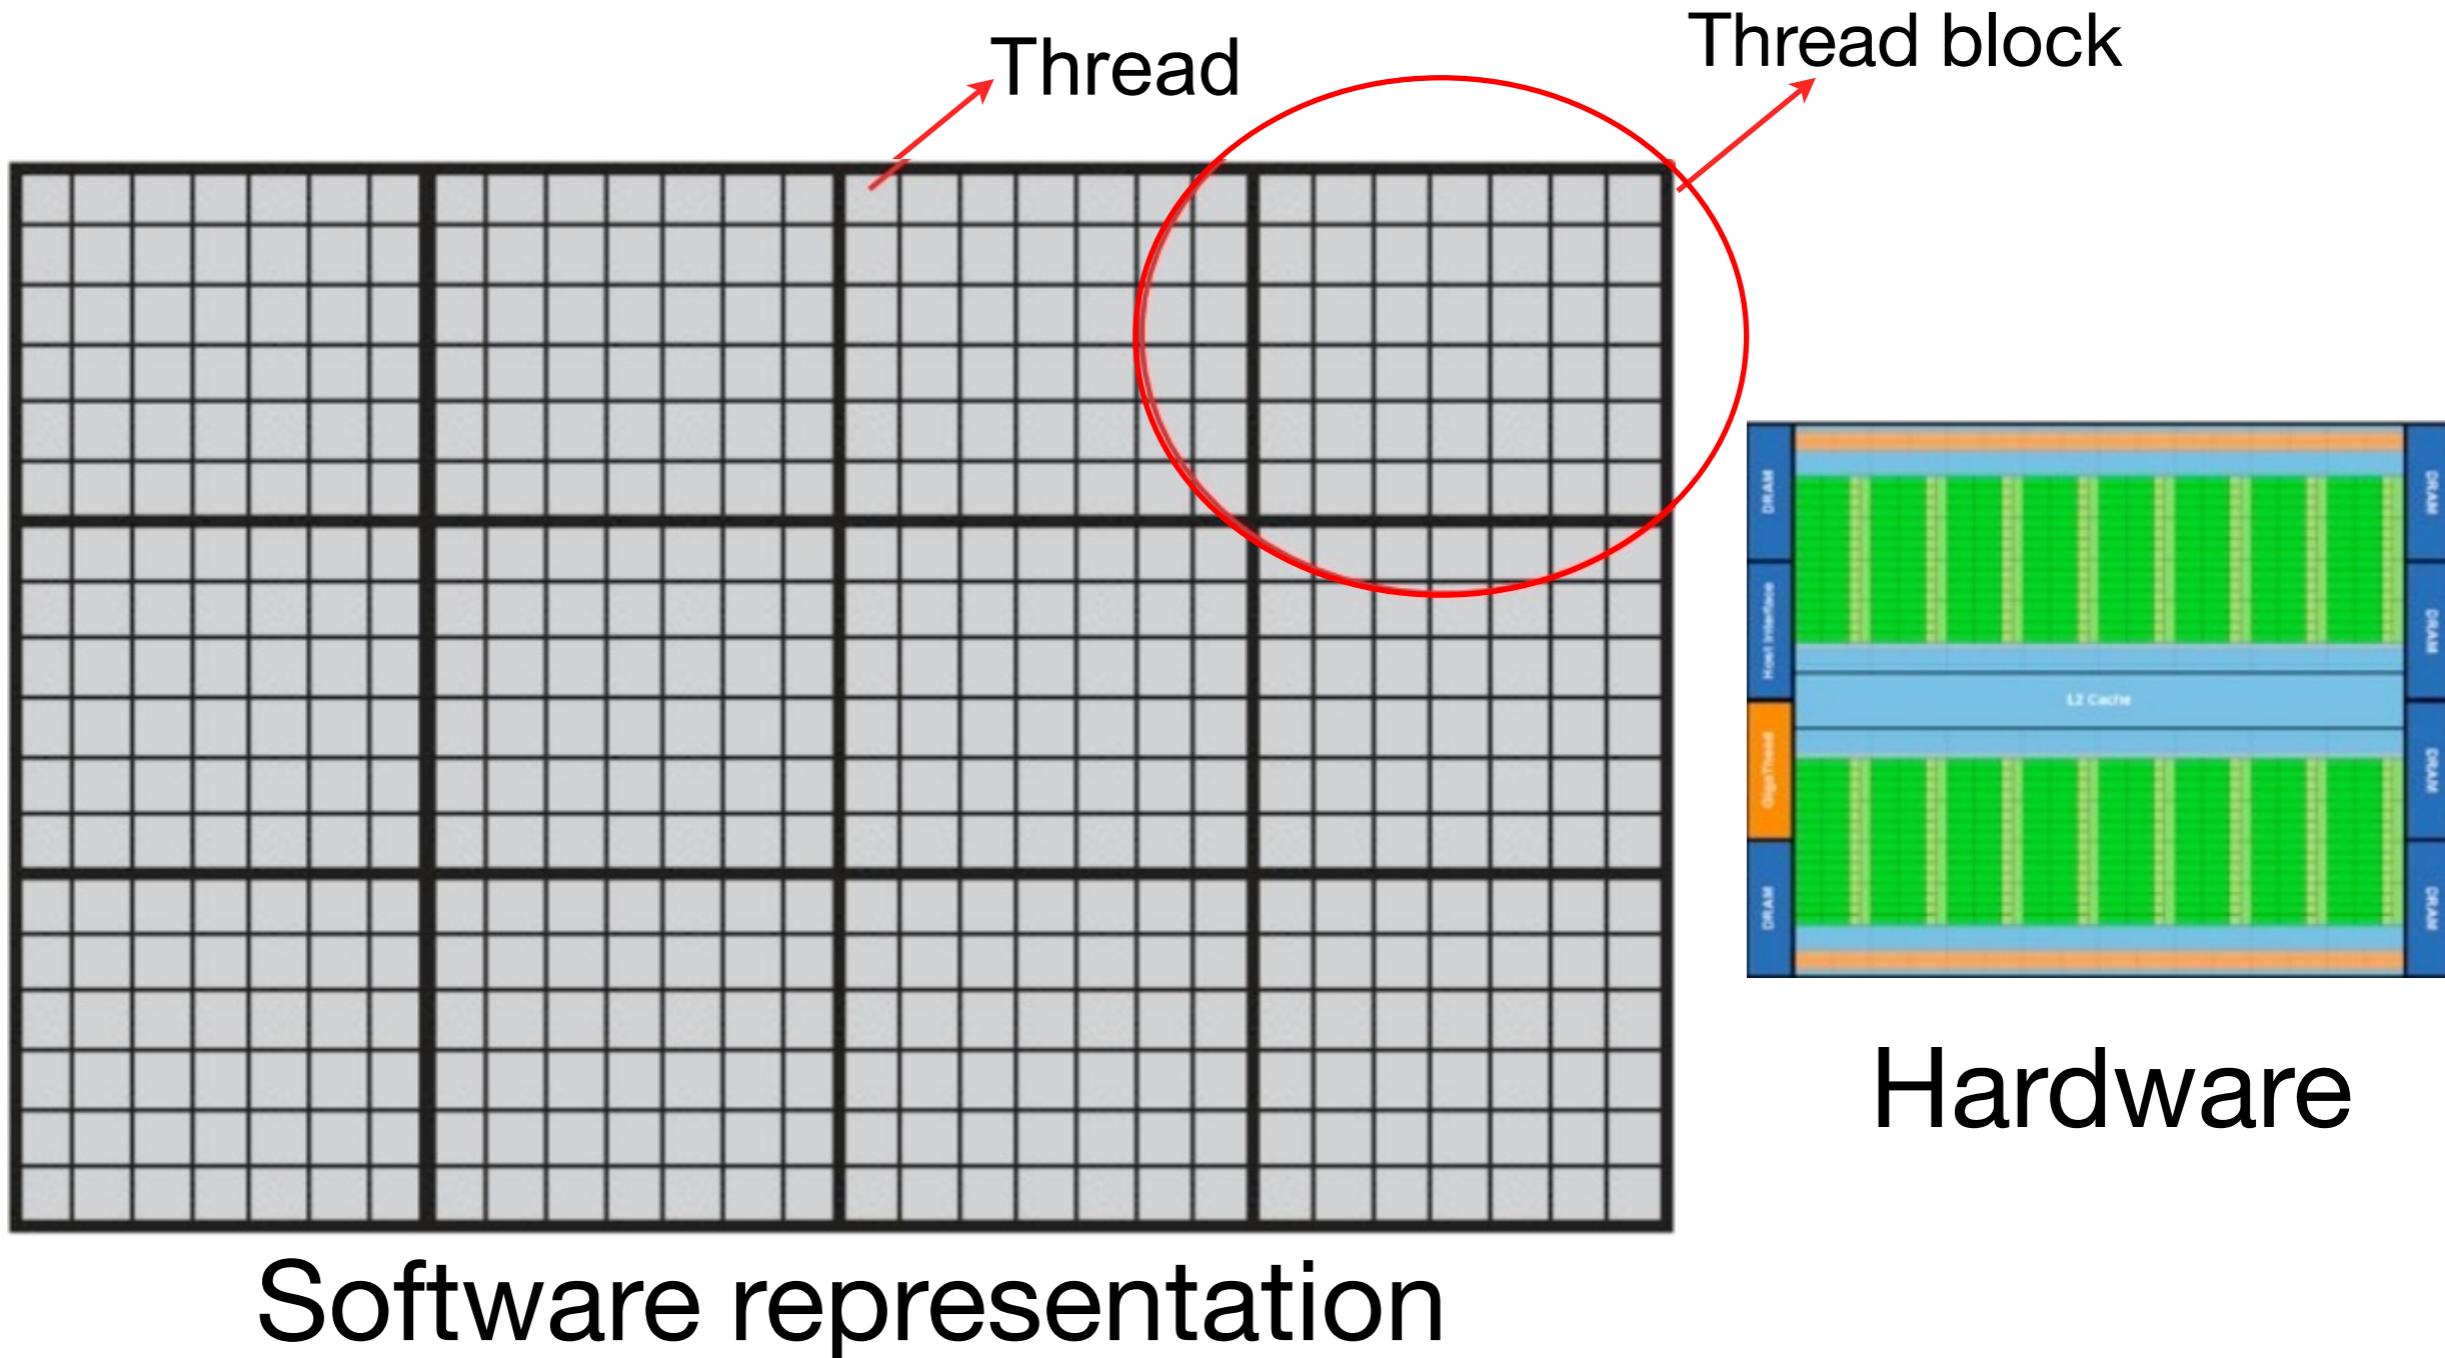
Software representation



Hardware

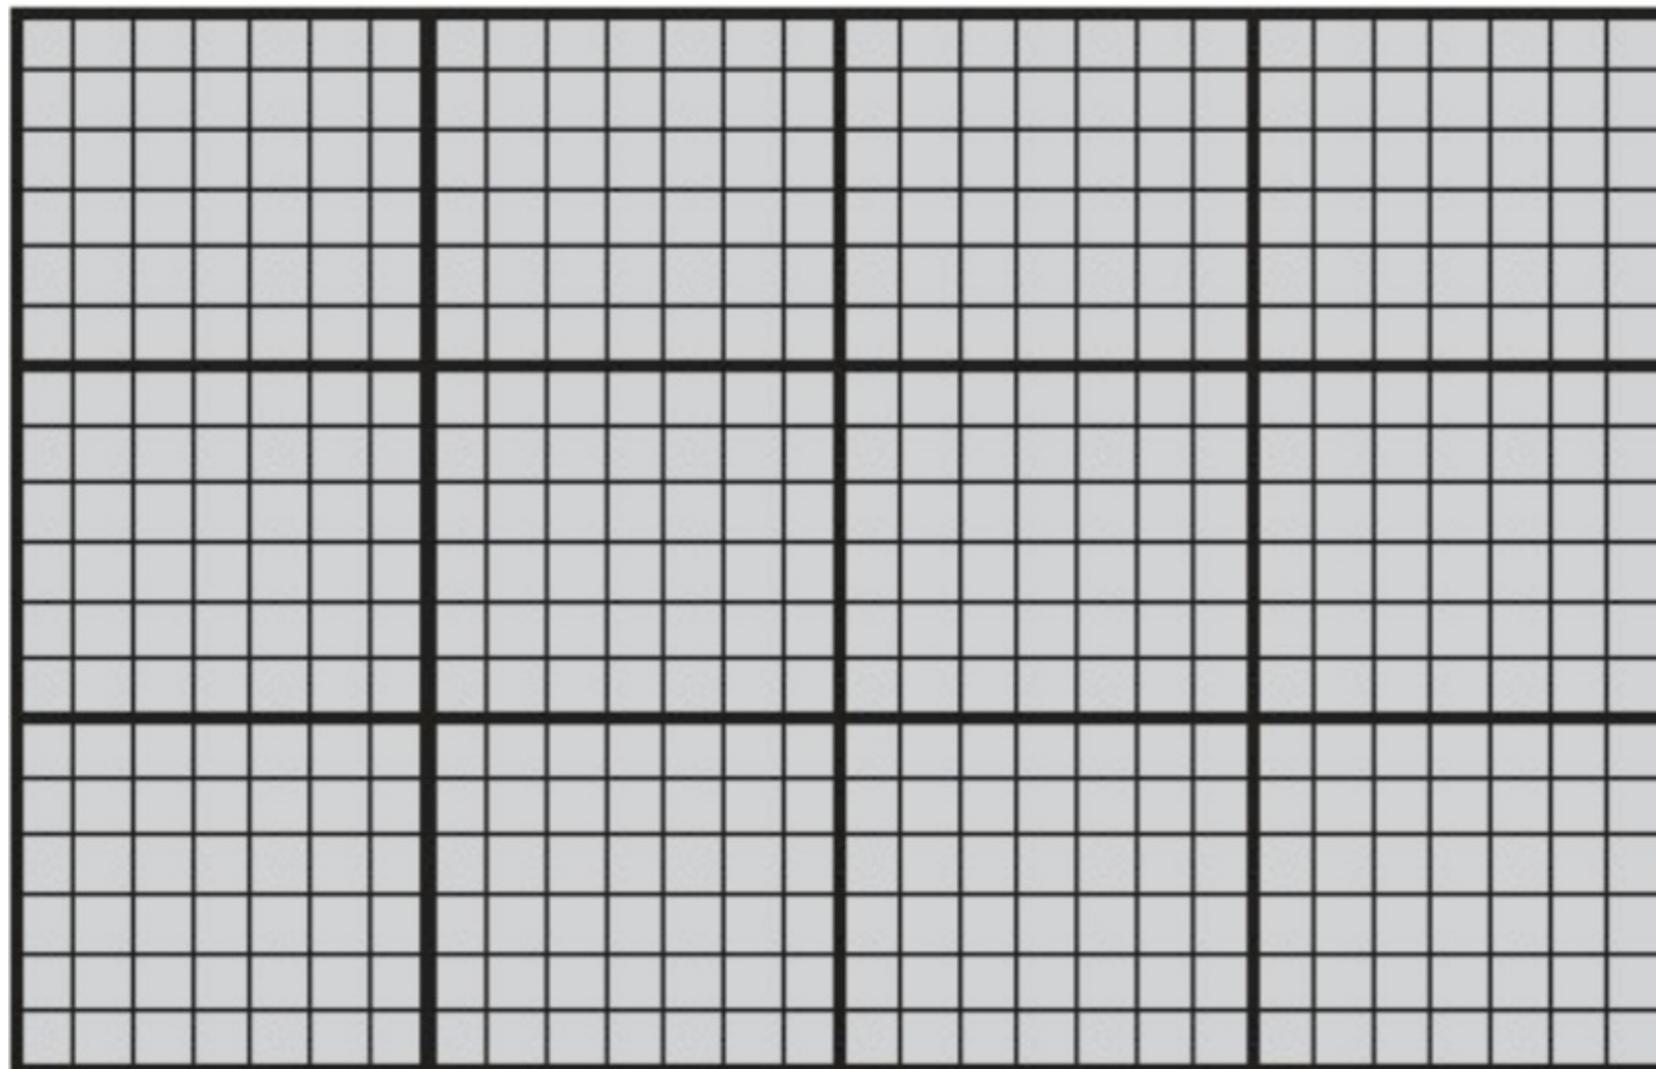
CUDA - Programming model

- ▶ Connecting the hardware and the software

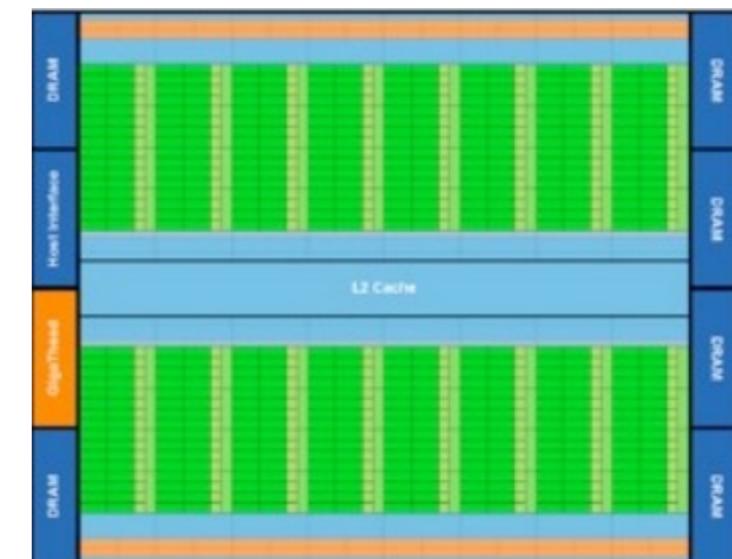


CUDA - Programming model

- ▶ Connecting the hardware *and the software*



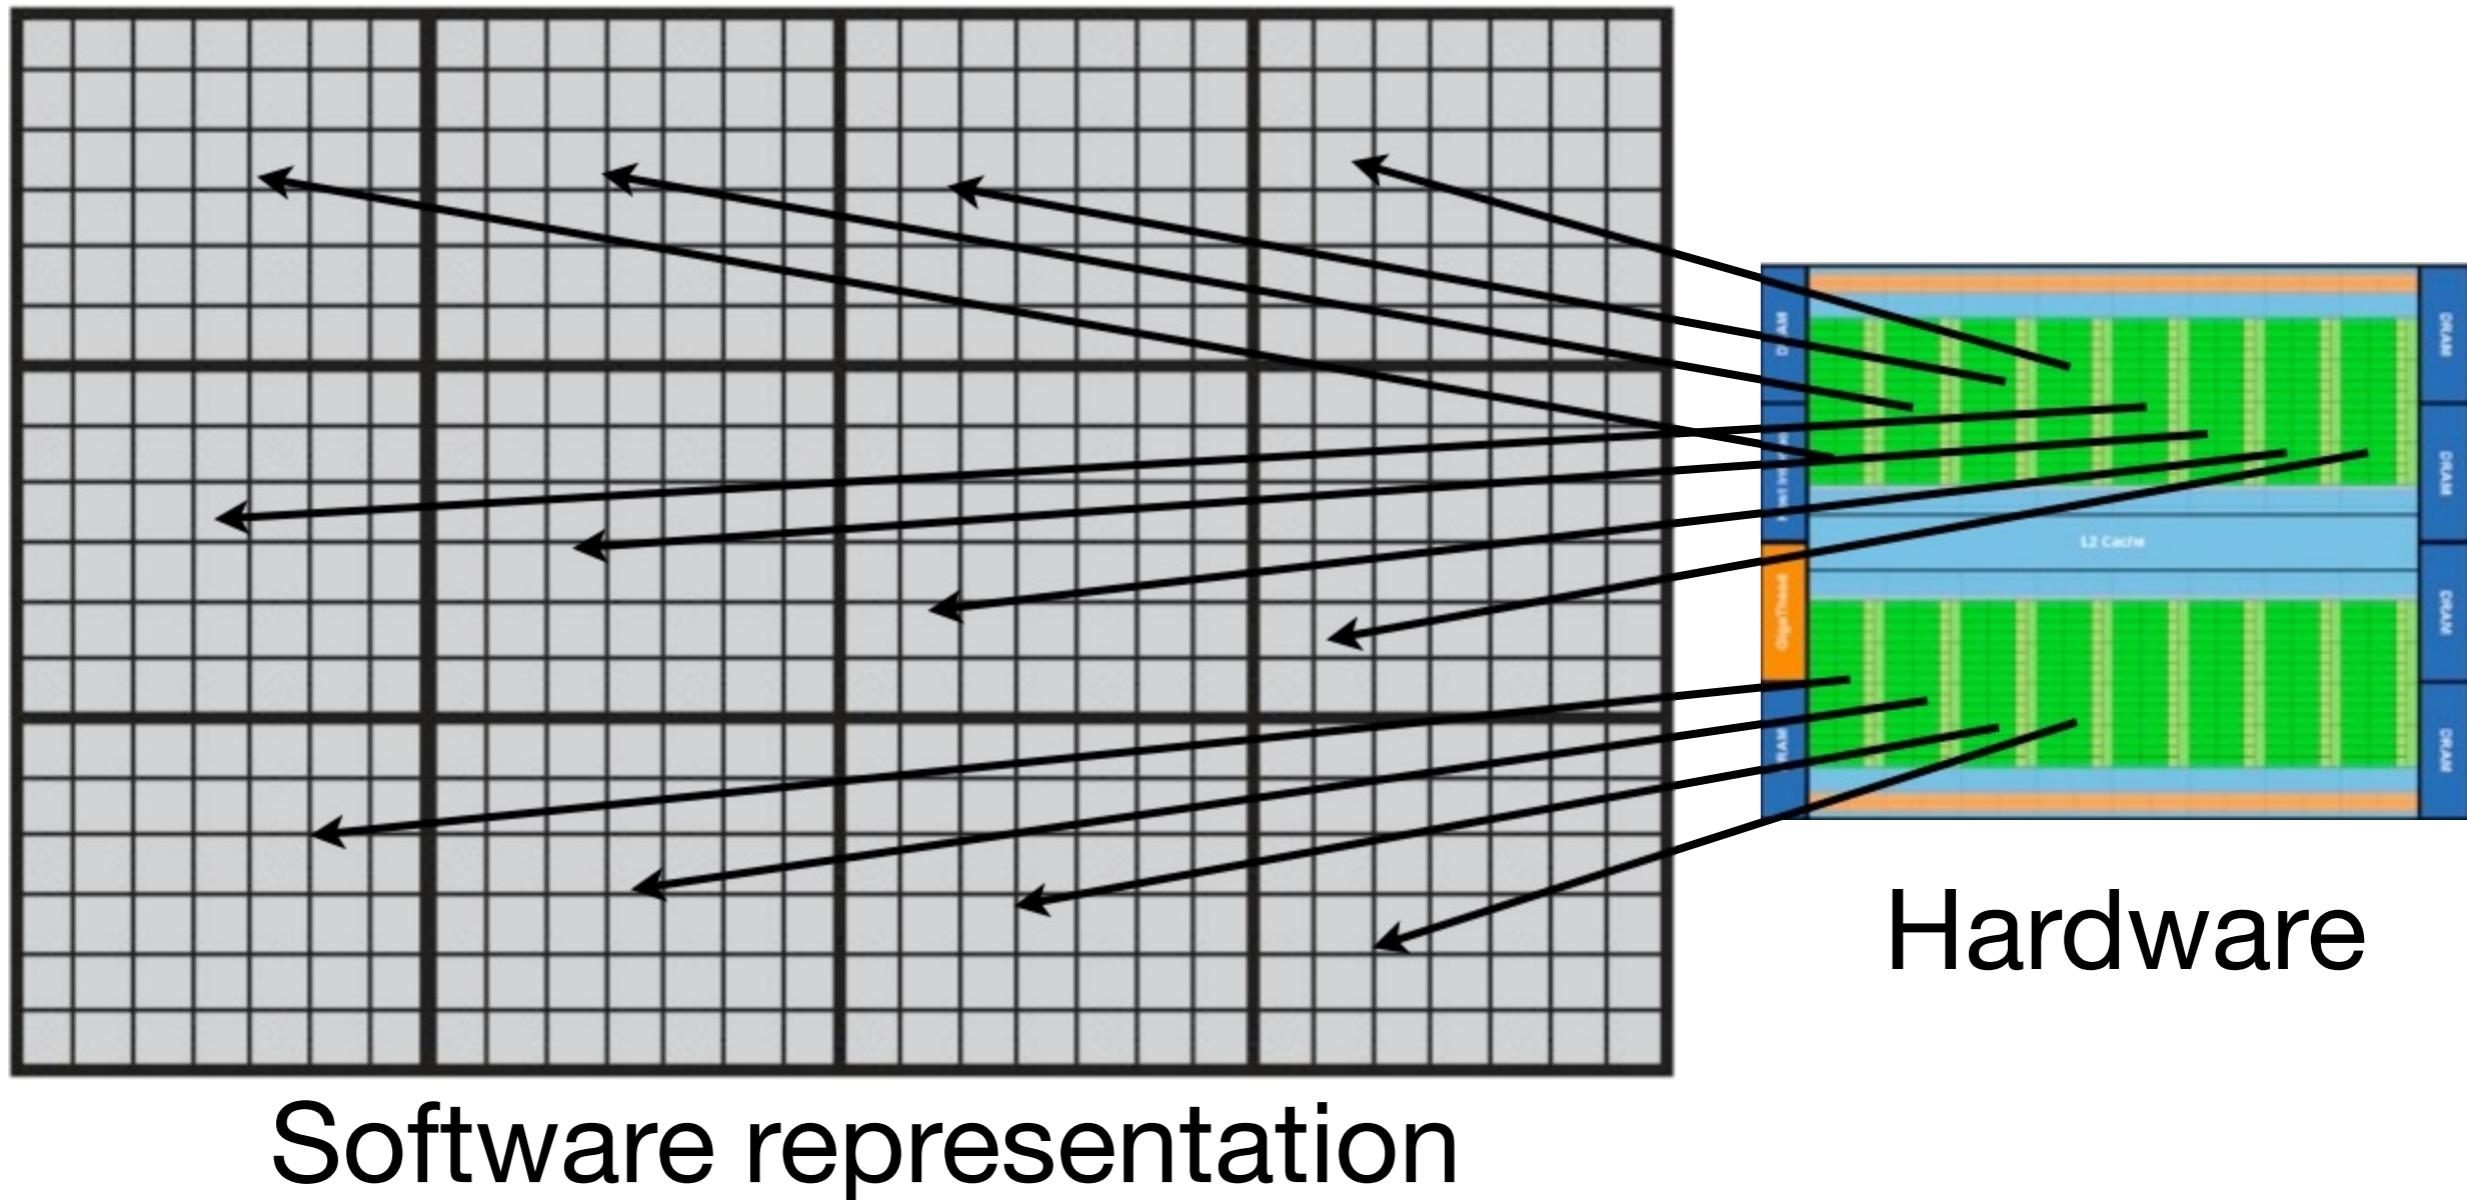
Software representation



Hardware

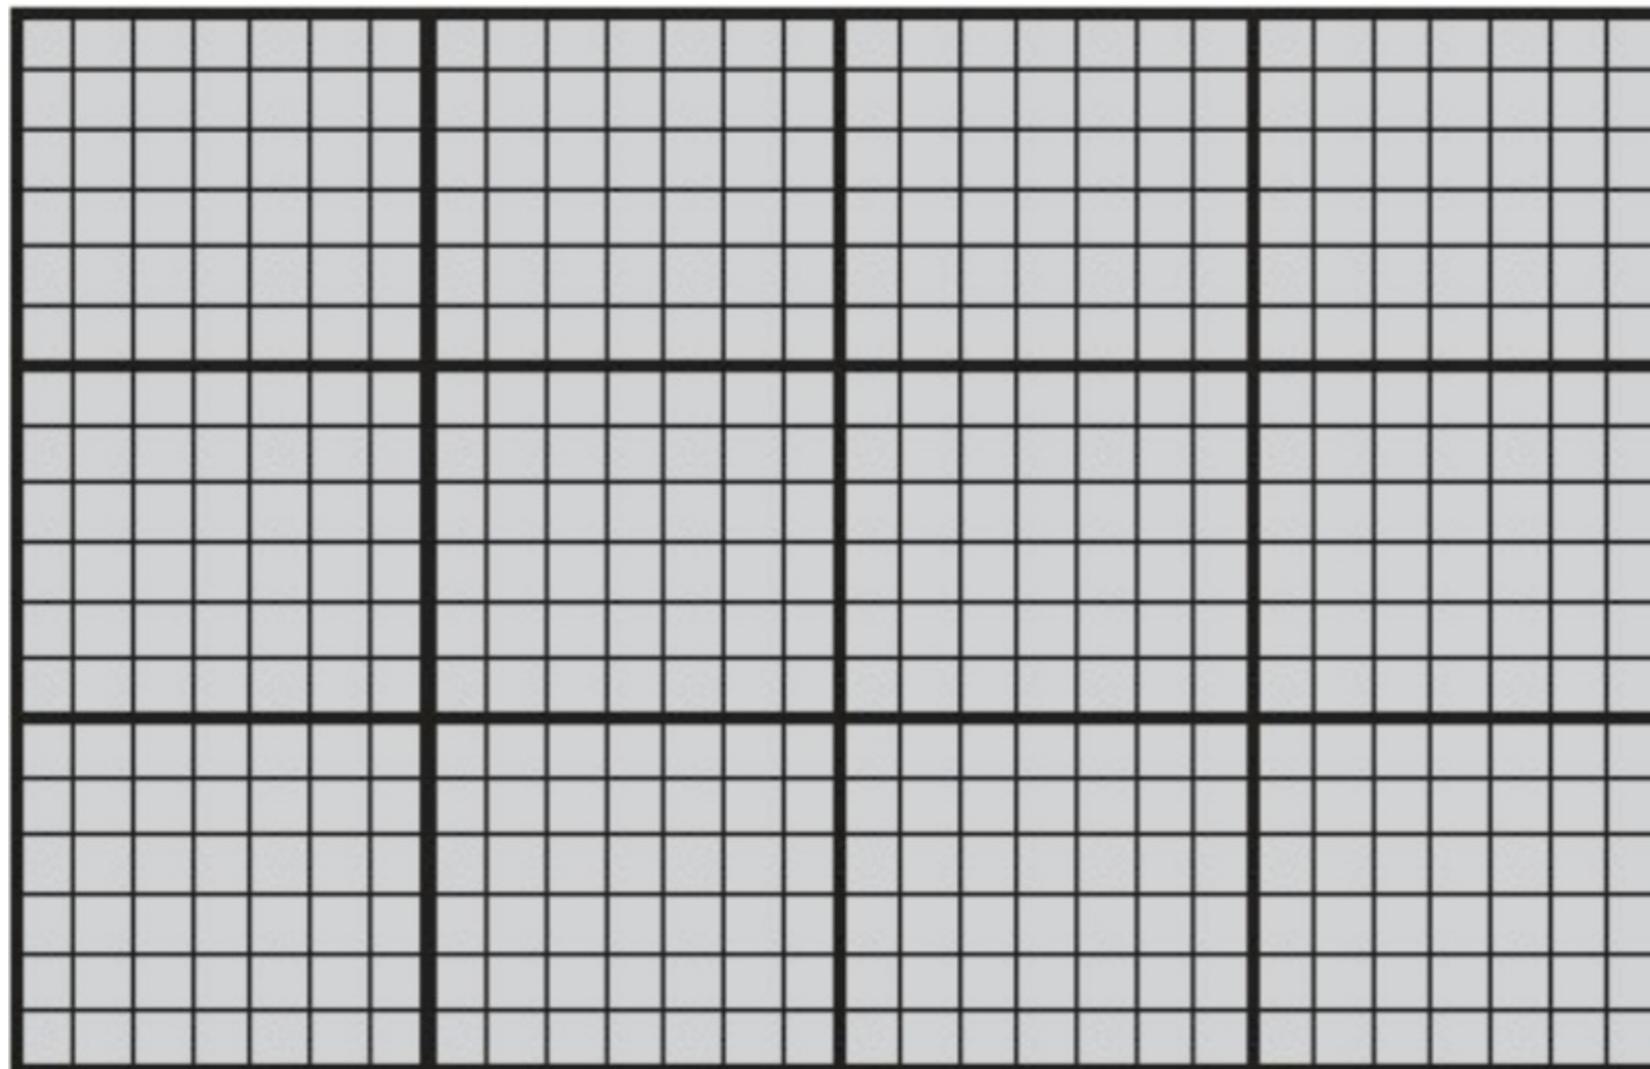
CUDA - Programming model

- ▶ Connecting the hardware and the software

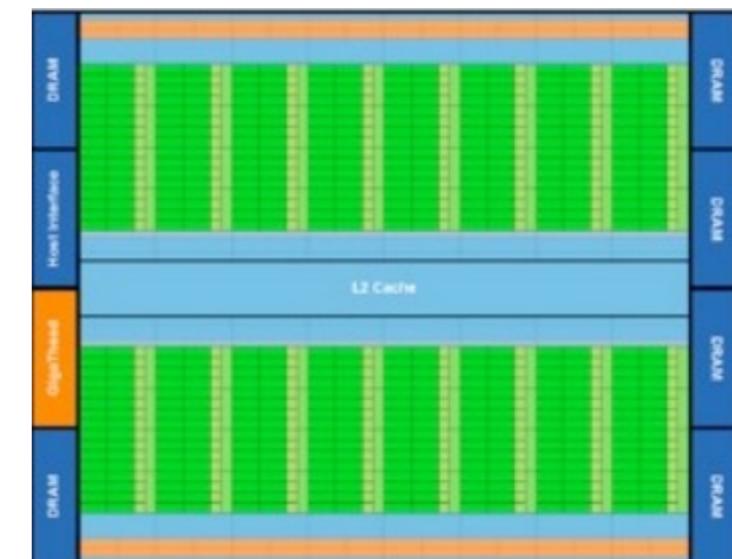


CUDA - Programming model

- ▶ Connecting the hardware and the software



Software representation



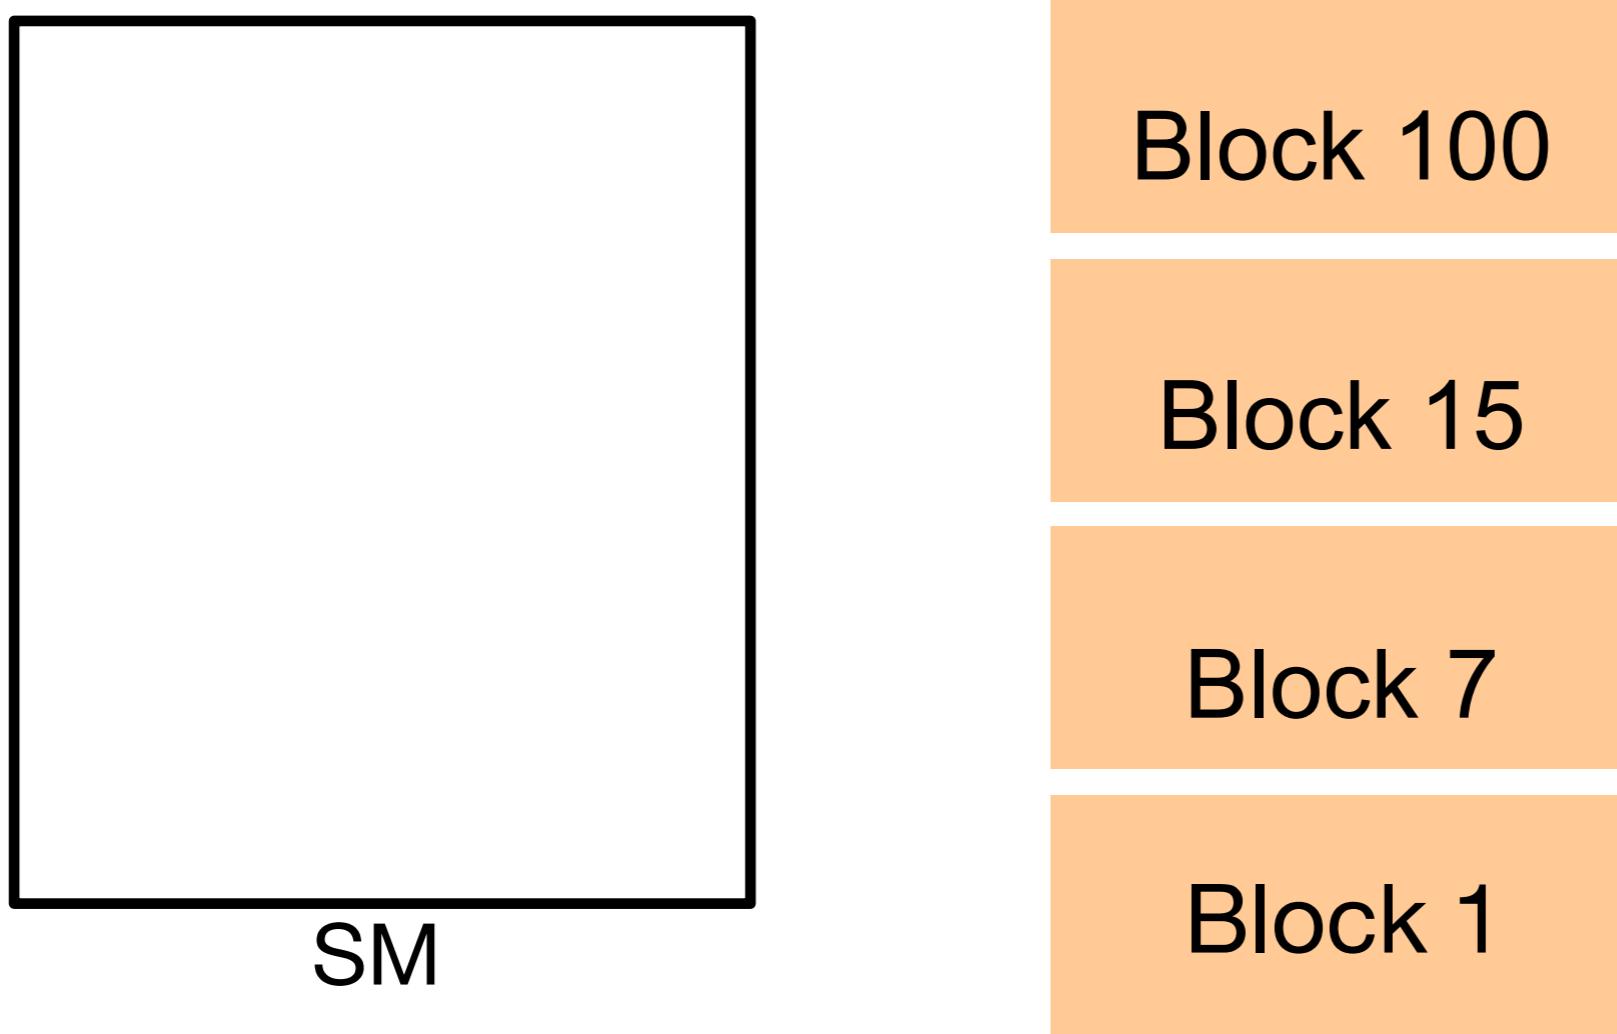
Hardware

How a Streaming Multiprocessor (SM) works

- ▶ CUDA Threads are grouped in thread blocks
 - All threads of the same thread block are **EXECUTED** in the same SM at the same time
 - ▶ SMs have shared memory, then threads within a thread block can communicate
 - The entirety of the threads of a thread block must be executed before there is space to schedule another thread block

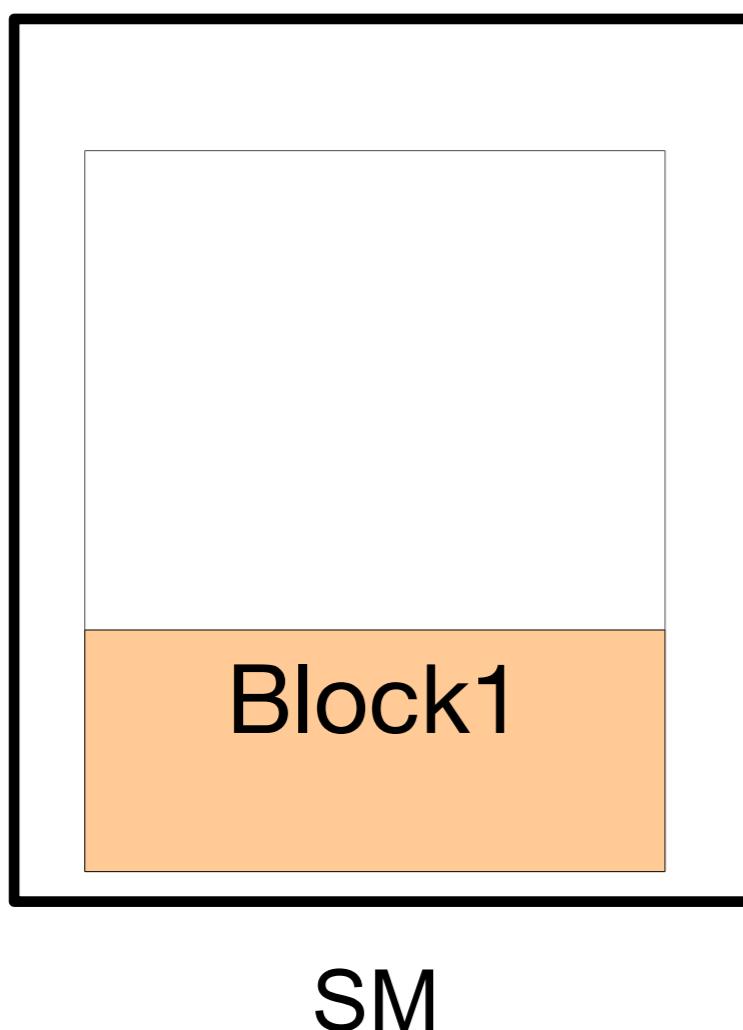
How a Streaming Multiprocessor (SM) works

- ▶ Hardware schedules thread blocks onto available SMs
 - No guarantee of order of execution
 - If a SM has more resources the hardware will schedule more blocks



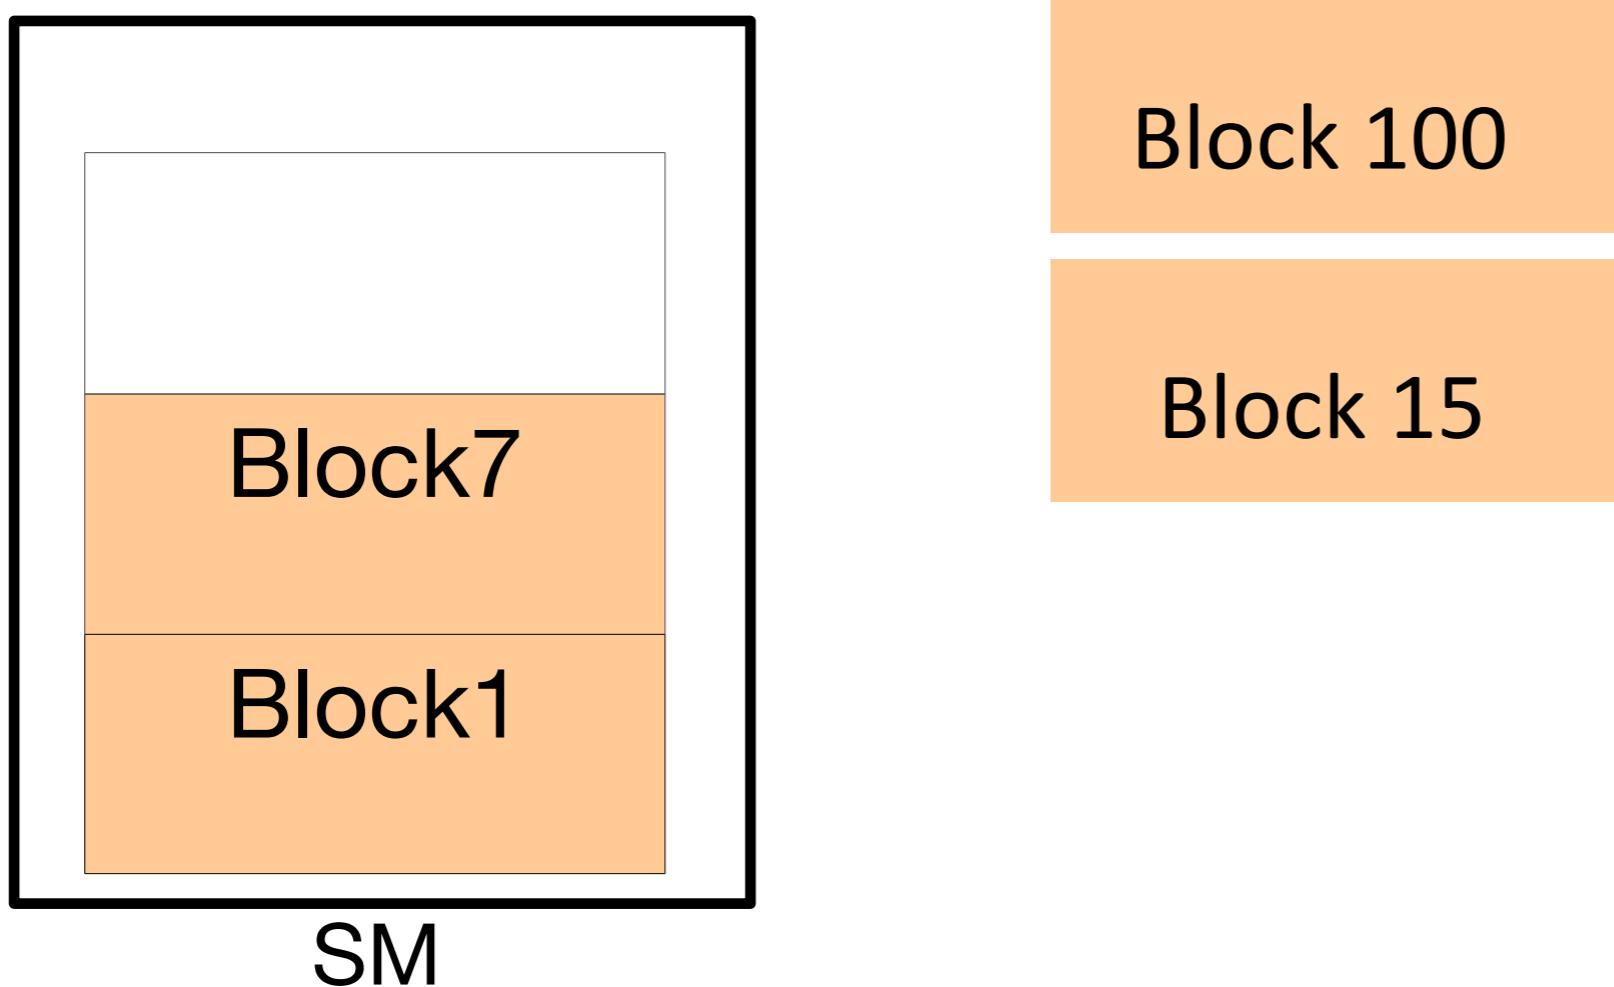
How a Streaming Multiprocessor (SM) works

- ▶ Hardware schedules thread blocks onto available SMs
 - No guarantee of order of execution
 - If a SM has more resources the hardware will schedule more blocks



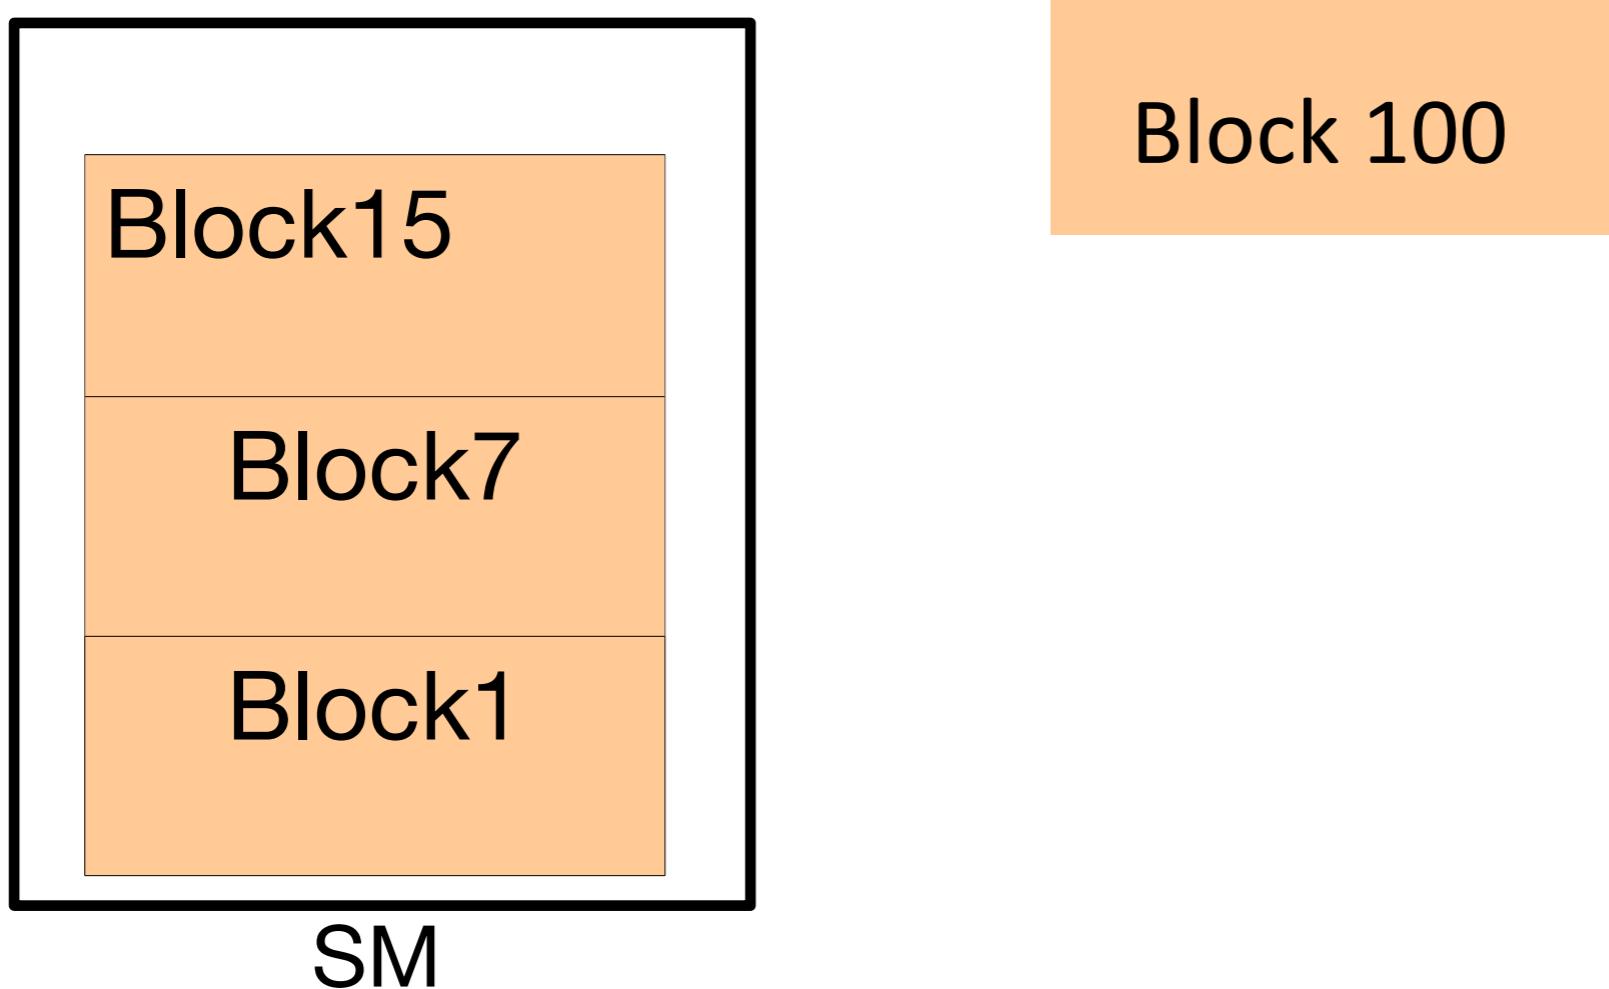
How a Streaming Multiprocessor (SM) works

- ▶ Hardware schedules thread blocks onto available SMs
 - No guarantee of order of execution
 - If a SM has more resources the hardware will schedule more blocks



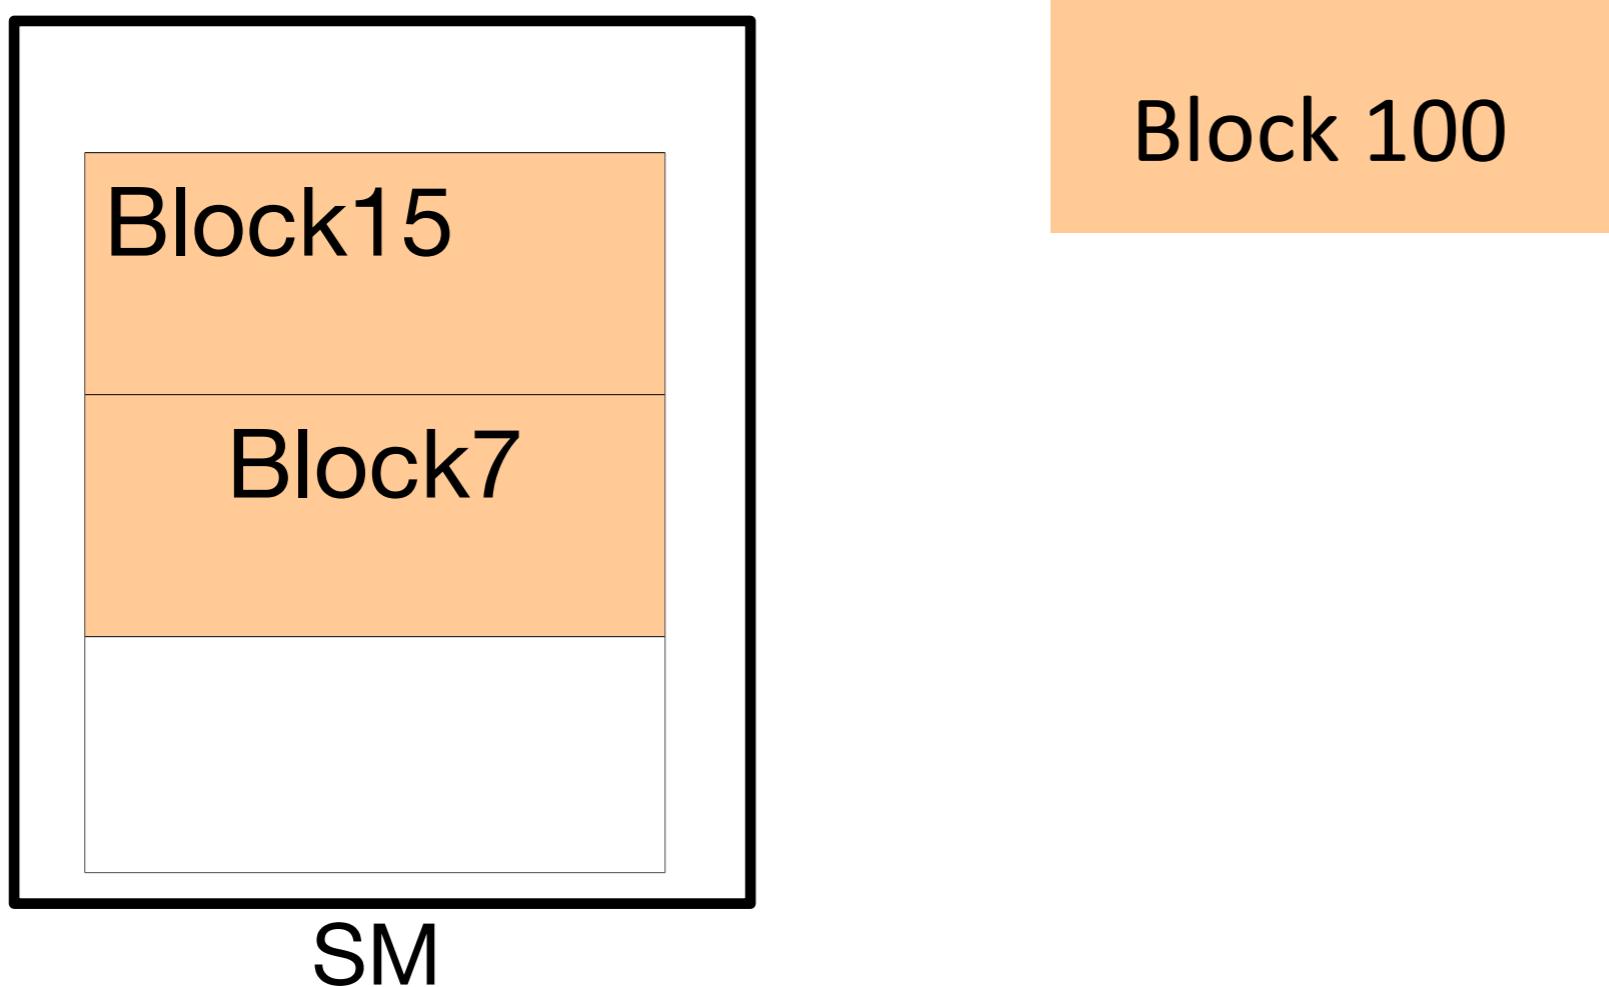
How a Streaming Multiprocessor (SM) works

- ▶ Hardware schedules thread blocks onto available SMs
 - No guarantee of order of execution
 - If a SM has more resources the hardware will schedule more blocks



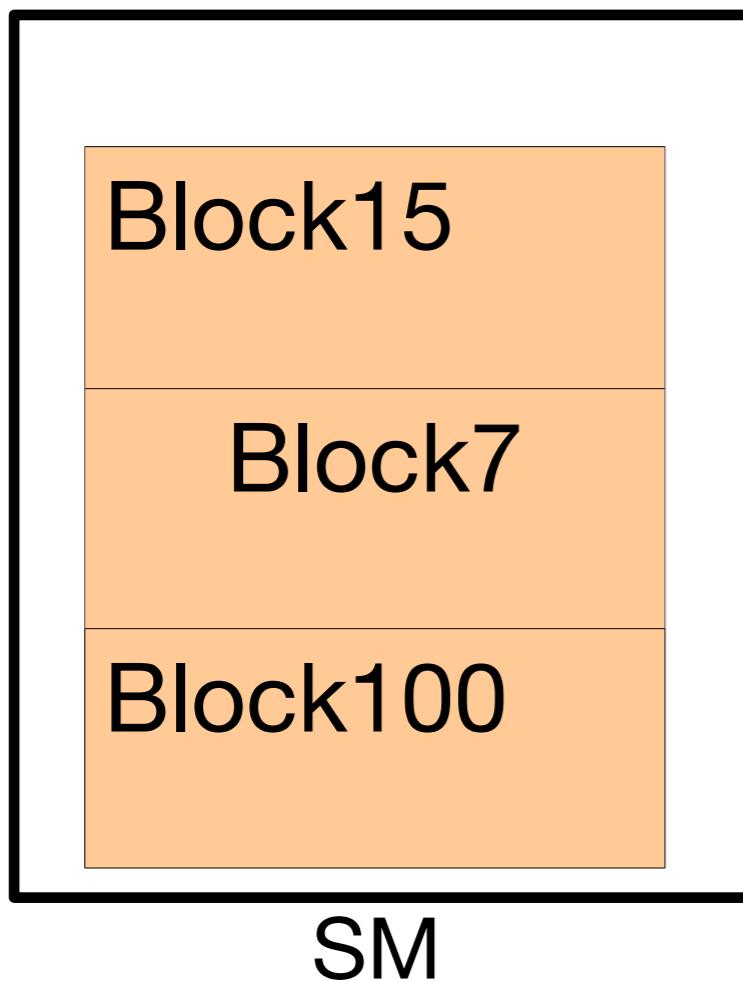
How a Streaming Multiprocessor (SM) works

- ▶ Hardware schedules thread blocks onto available SMs
 - No guarantee of order of execution
 - If a SM has more resources the hardware will schedule more blocks



How a Streaming Multiprocessor (SM) works

- ▶ Hardware schedules thread blocks onto available SMs
 - No guarantee of order of execution
 - If a SM has more resources the hardware will schedule more blocks



Warps

- ▶ Inside the SM, threads are launched in groups of 32 called warps
 - Warps share the control part (warp scheduler)
 - At any time, only one warp is executed per SM
 - Threads in a warp will be executing the same instruction
 - Half warps for compute capability 1.X
- Fermi:
 - ▶ Maximum number of active threads $1024 * 8 * 32 = 262144$

Warps

- ▶ Inside the SM, threads are launched in groups of 32 called warps
 - Warps share the control part (warp scheduler)
 - At any time, only one warp is executed per SM
 - Threads in a warp will be executing the same instruction
 - Half warps for compute capability 1.X

- Fermi:
 - ▶ Maximum number of active threads $1024 * 8 * 32 = 262144$

Max threads
per block



Warps

- ▶ Inside the SM, threads are launched in groups of 32 called warps
 - Warps share the control part (warp scheduler)
 - At any time, only one warp is executed per SM
 - Threads in a warp will be executing the same instruction
 - Half warps for compute capability 1.X

- Fermi:

- ▶ Maximum number of active threads $1024 * 8 * 32 = 262144$

Max threads per block Max blocks per SM

```
graph TD; A[Max threads per block] --> C["1024 * 8 * 32 = 262144"]; B[Max blocks per SM] --> C;
```

Warps

- ▶ Inside the SM, threads are launched in groups of 32 called warps
 - Warps share the control part (warp scheduler)
 - At any time, only one warp is executed per SM
 - Threads in a warp will be executing *the same instruction*
 - Half warps for compute capability 1.X

- Fermi:

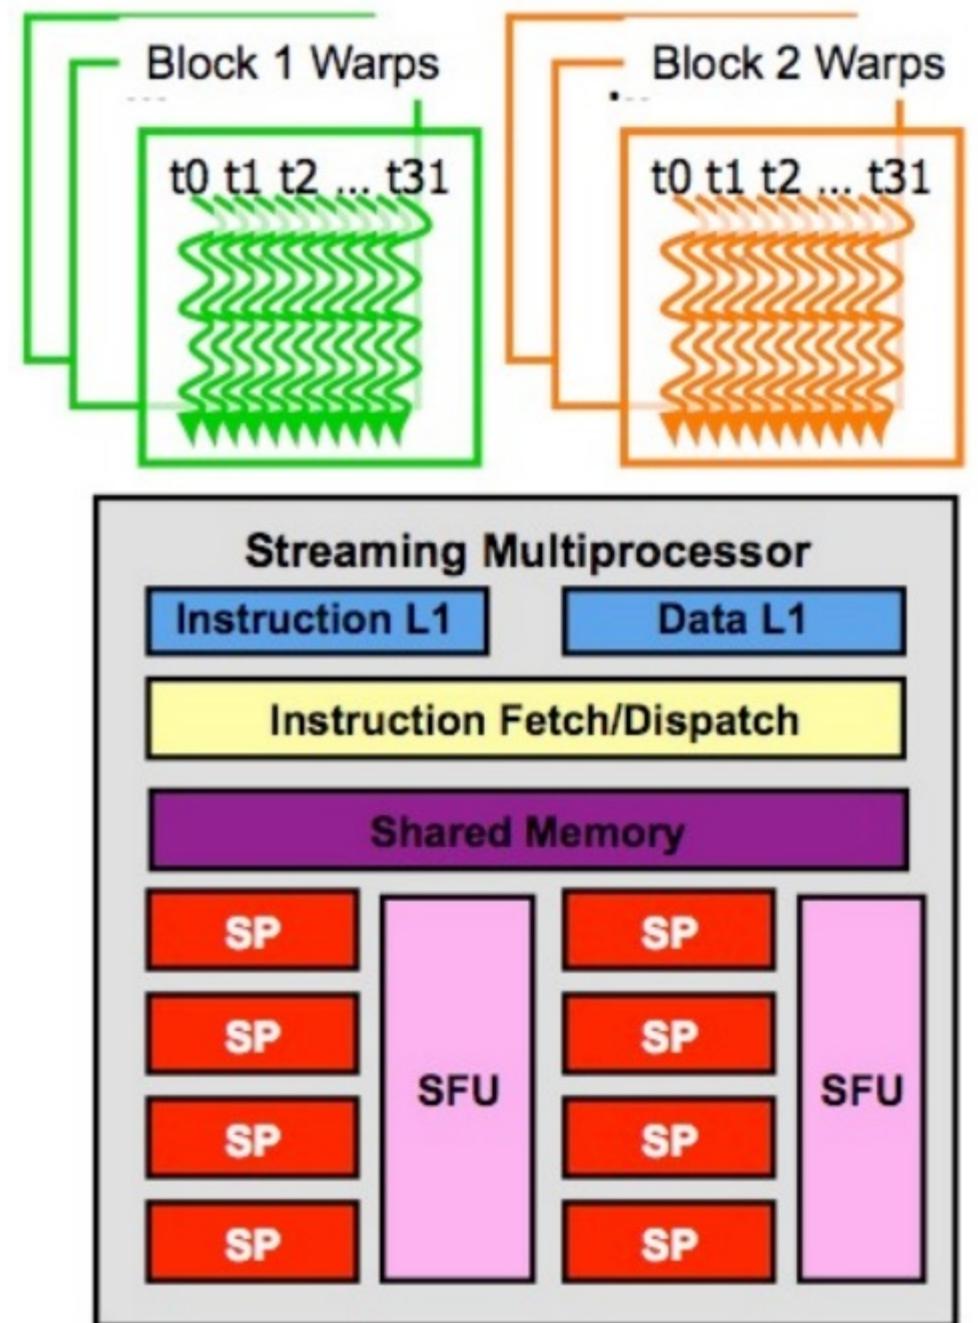
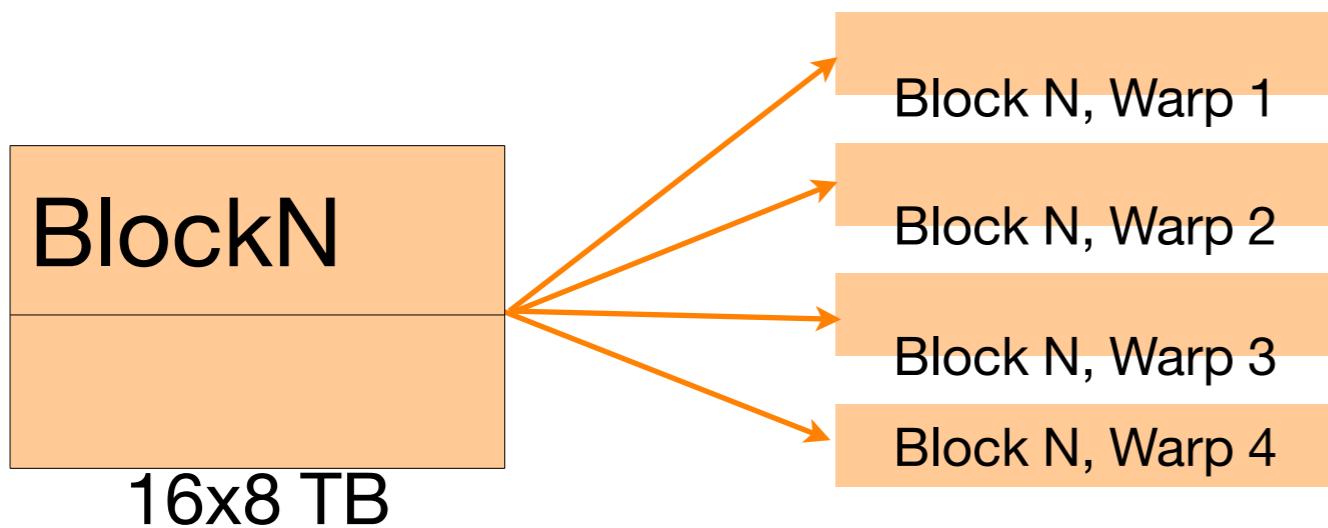
- ▶ Maximum number of active threads $1024 * 8 * 32 = 262144$

Max threads per block	Max blocks per SM	SMs
1024	8	32

$$1024 * 8 * 32 = 262144$$

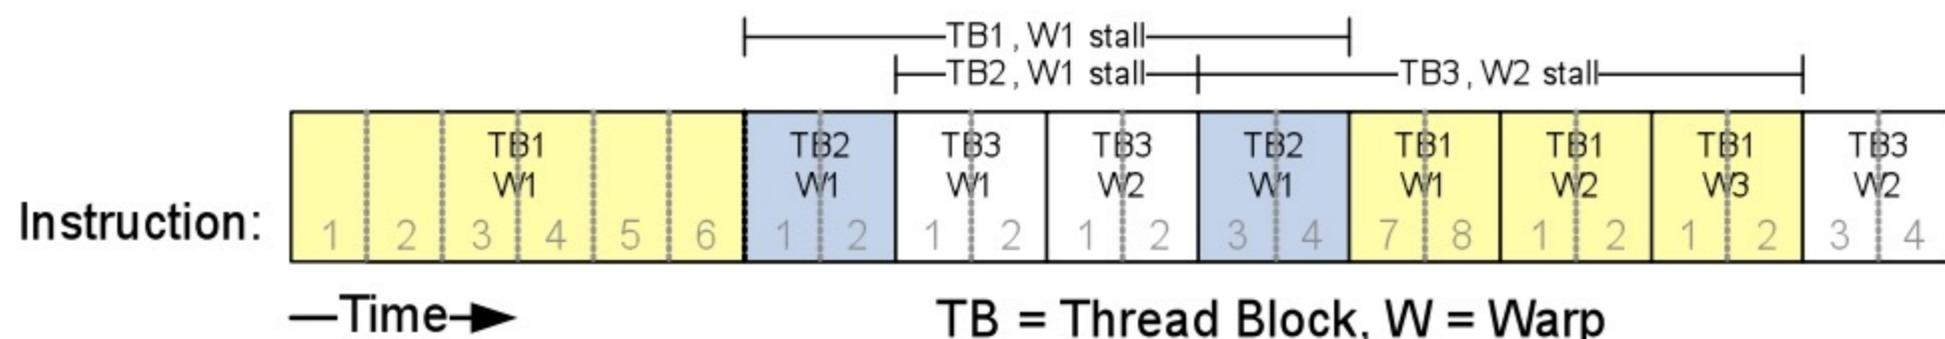
Warp designation

- ▶ Hardware separates threads of a block into warps
 - All threads in a warp correspond to the same thread block
 - Threads are placed in a warp sequentially
 - Threads are scheduled in warps



Warp scheduling

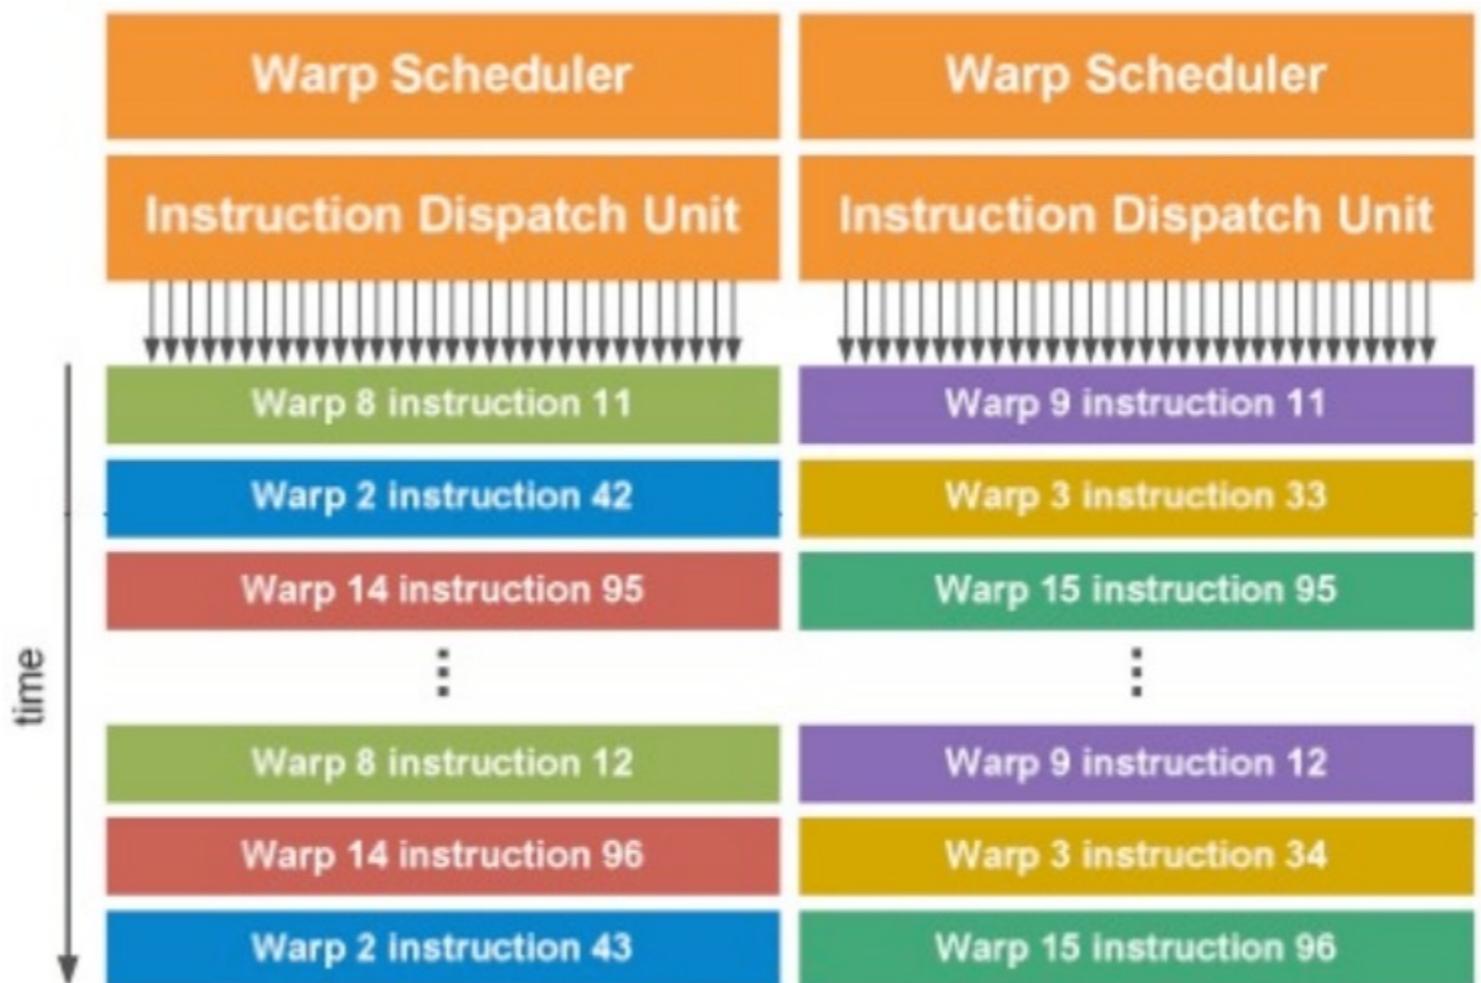
- ▶ The SM implements a zero-overhead warp scheduling
 - Warps whose next instruction has its operands ready for consumption are eligible for execution
 - Eligible warps are selected for execution on a prioritized scheduling policy
 - All threads in a warp execute the same instruction



Warp scheduling

► Fermi

- Double warp scheduler
- Each SM has two warp schedulers and two instruction units



Memory hierarchy

- ▶ CUDA works in both the CPU and GPU
 - One has to keep track of which memory is operating on (host - device)
 - Within the GPU there are also different memory spaces

- Each thread can:
 - Read/write per-thread **registers**
 - Read/write per-thread **local memory**
 - Read/write per-block **shared memory**
 - Read/write per-grid **global memory**
 - Read/only per-grid **constant memory**

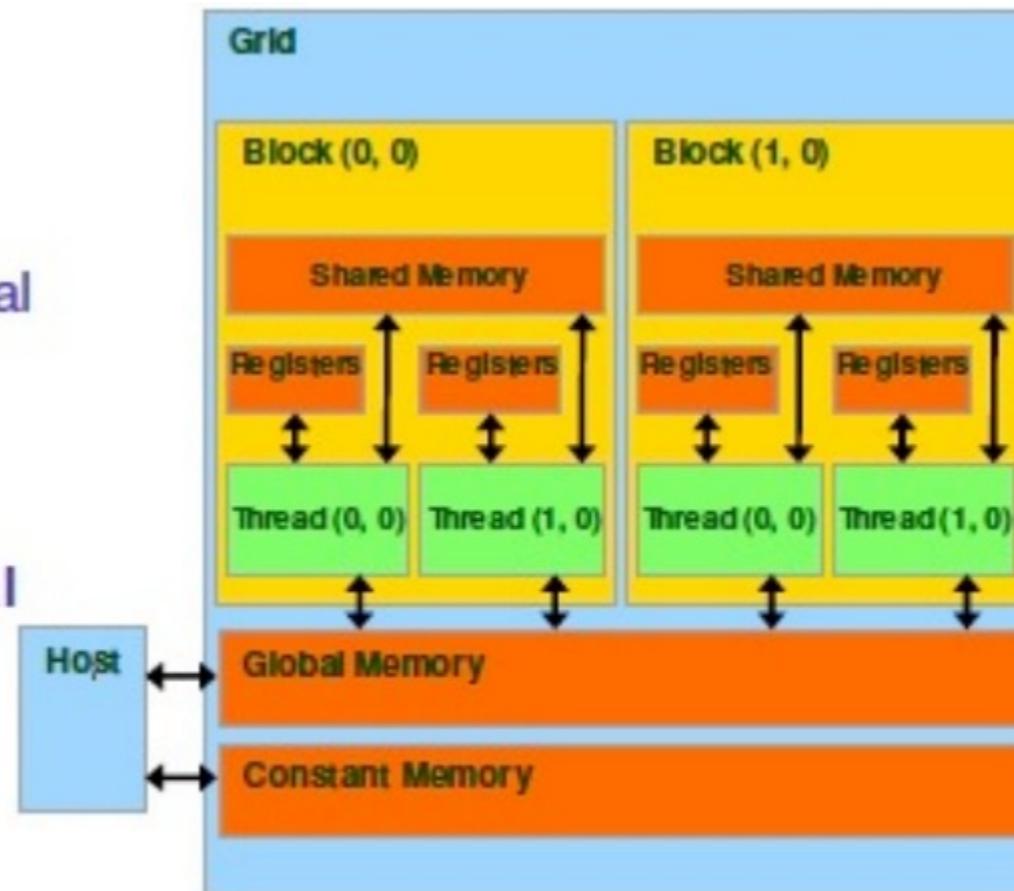
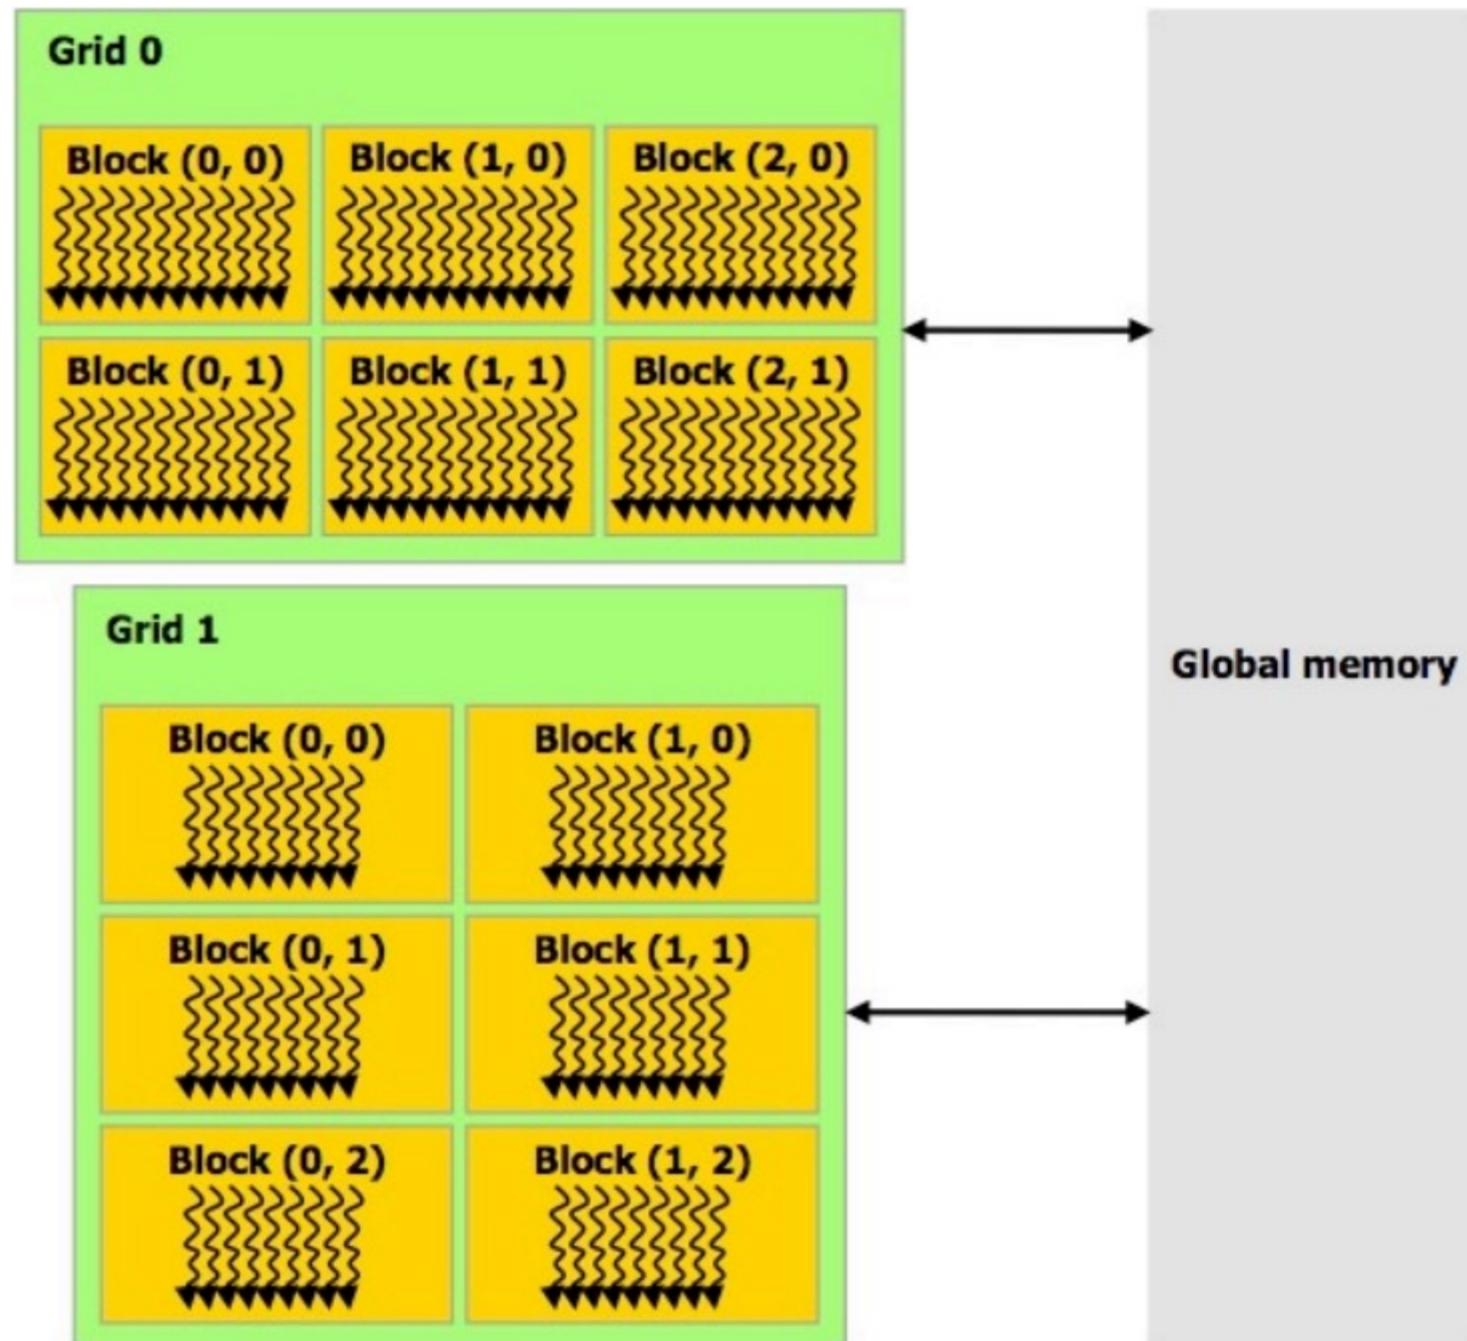


Figure 4.2 GeForce 8800GTX Implementation of CUDA Memories

Memory hierarchy

► Global memory

- Main GPU memory
- Communicates with host
- Can be seen by all threads
- Order of GB
- Off chip, slow (~400 cycles)



`__device__ float variable;`

Memory hierarchy

▶ Shared memory

- Per SM

- Seen by threads of the same thread block

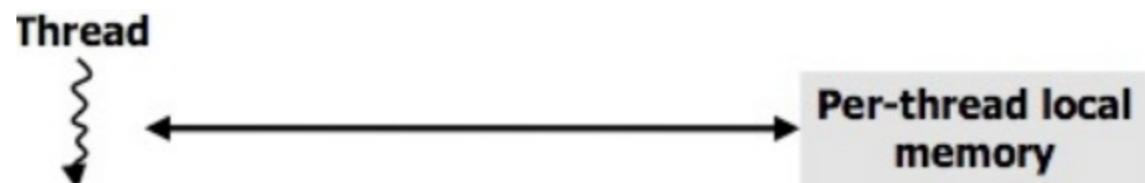
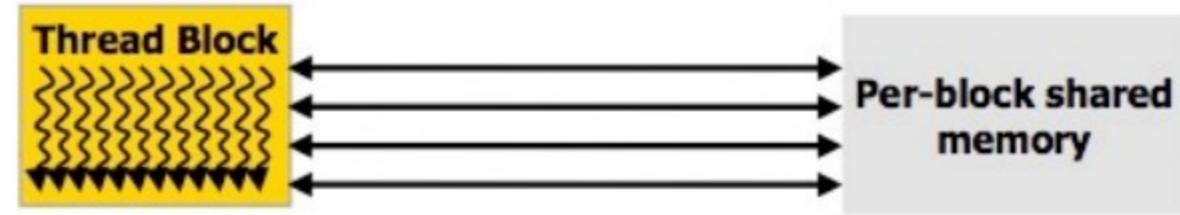
- Order of kB

- On chip, fast (~4 cycles) `__shared__ float variable;`

▶ Registers

- Private to each thread

- On chip, fast



`float variable;`

Memory hierarchy

► Local memory

- Private to each thread
- Off chip, slow
- Register overflows



```
float variable [10];
```

► Constant memory

- Read only
- Off chip, but fast (cached)
- Seen by all threads
- 64kB with 8kB cache

```
__constant__ float variable;
```

Memory hierarchy

- ▶ Texture memory
 - Seen by all threads
 - Read only
 - Off chip, but fast (cached) if cache hit
 - Cache optimized for 2D locality
 - Binds to global

```
texture<type, dim> tex_var;
cudaChannelFormatDesc();
cudaBindTexture2D(...);
tex2D(tex_var, x_index, y_index);
```

Memory hierarchy

► Texture memory

- Seen by all threads
- Read only
- Off chip, but fast (cached) if cache hit
- Cache optimized for 2D locality
- Binds to global

```
texture<type, dim> tex_var;  
cudaChannelFormatDesc();  
cudaBindTexture2D(...);  
tex2D(tex_var, x_index, y_index);
```

Initialize



Memory hierarchy

► Texture memory

- Seen by all threads
- Read only
- Off chip, but fast (cached) if cache hit
- Cache optimized for 2D locality

- Binds to global

```
texture<type, dim> tex_var;  
cudaChannelFormatDesc();  
cudaBindTexture2D(...);  
tex2D(tex_var, x_index, y_index);
```

Initialize

Options

Memory hierarchy

► Texture memory

- Seen by all threads
- Read only
- Off chip, but fast (cached) if cache hit
- Cache optimized for 2D locality

- Binds to global

```
texture<type, dim> tex_var;  
cudaChannelFormatDesc();  
cudaBindTexture2D(...);  
tex2D(tex_var, x_index, y_index);
```

Initialize

Options



Bind

Memory hierarchy

- ▶ Texture memory

- Seen by all threads
- Read only
- Off chip, but fast (cached) if cache hit
- Cache optimized for 2D locality

- Binds to global

```
texture<type, dim> tex_var;  
cudaChannelFormatDesc();  
cudaBindTexture2D(...);  
tex2D(tex_var, x_index, y_index);
```

Initialize

Options

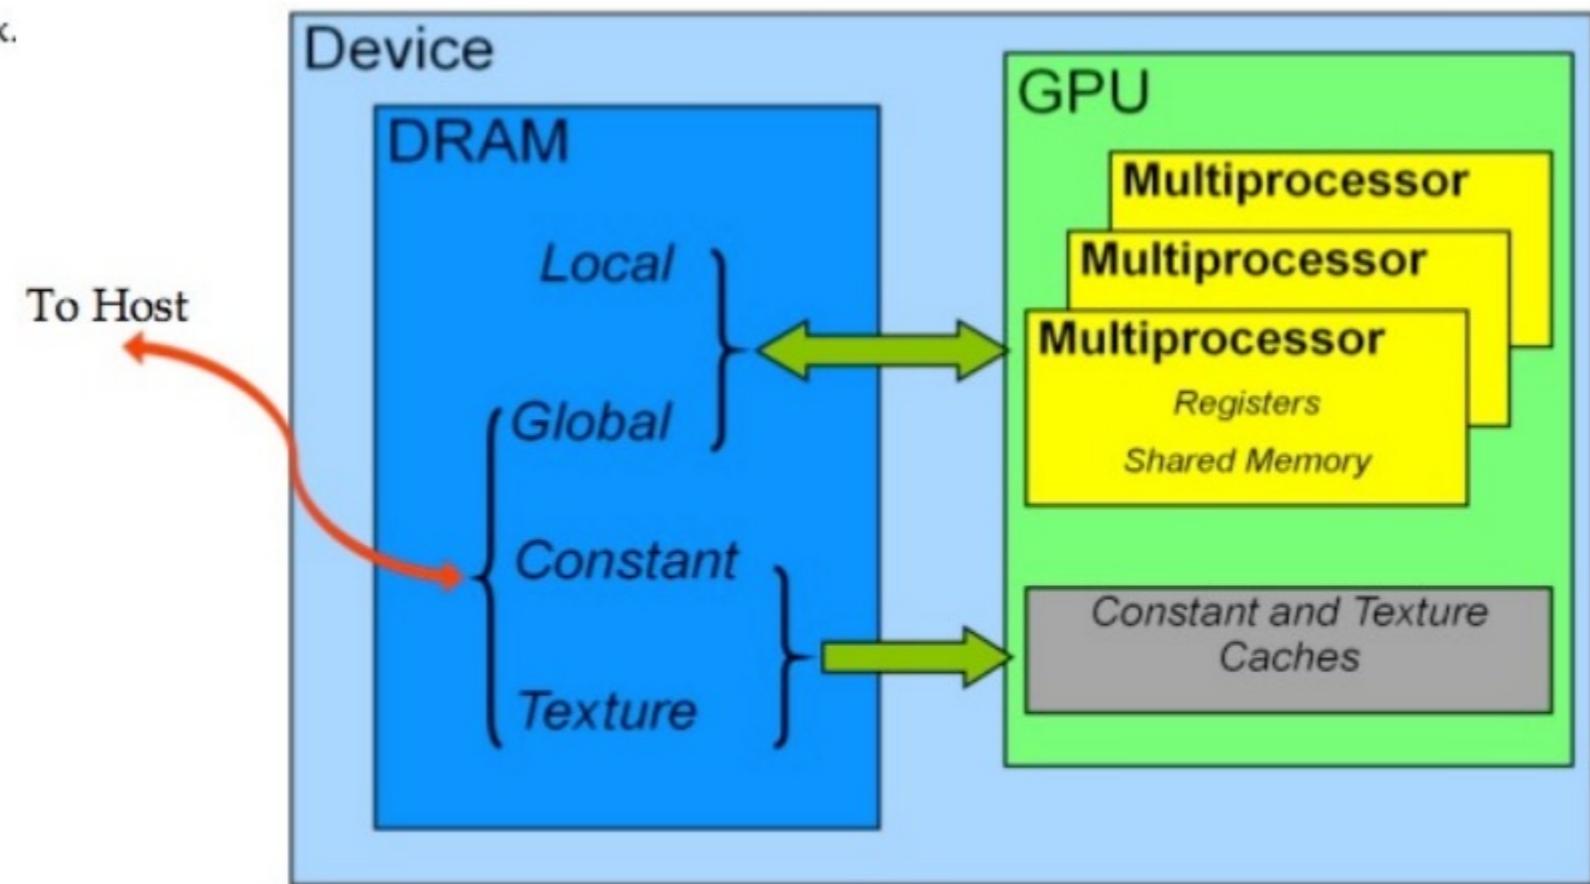
Bind

Fetch

Memory hierarchy

Memory	Location on/off chip	Cached	Access	Scope	Lifetime
Register	On	n/a	R/W	1 thread	Thread
Local	Off	†	R/W	1 thread	Thread
Shared	On	n/a	R/W	All threads in block	Block
Global	Off	†	R/W	All threads + host	Host allocation
Constant	Off	Yes	R	All threads + host	Host allocation
Texture	Off	Yes	R	All threads + host	Host allocation

[†]Cached only on devices of compute capability 2.x.

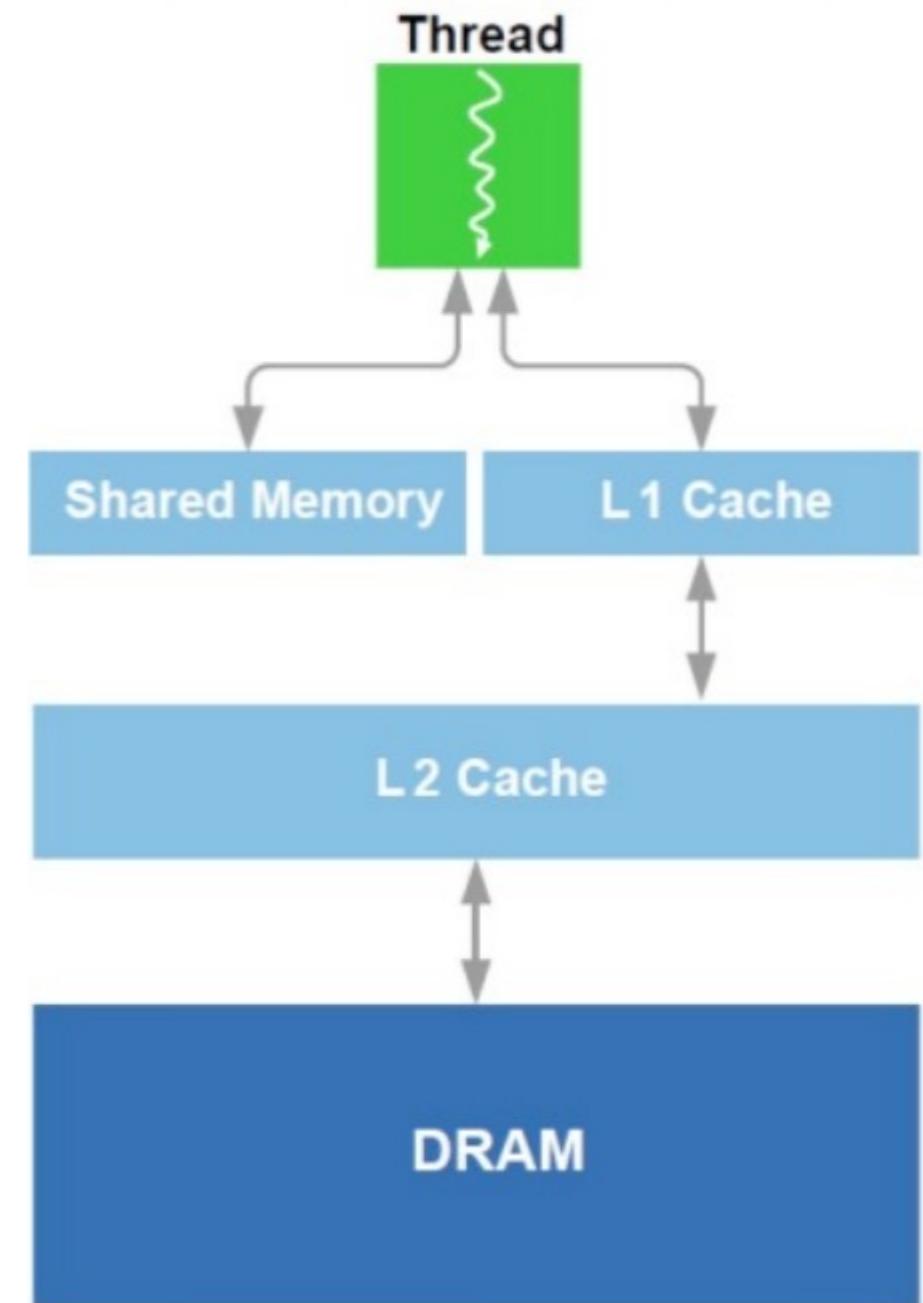


Memory hierarchy

► Case of Fermi

- Added an L1 cache to each SM
- Shared + cache = 64kB:
 - Shared = 48kB, cache = 16kB
 - Shared = 16kB, cache = 48kB

Fermi Memory Hierarchy



Memory hierarchy

- ▶ Use your memory strategically
 - Read only: `__constant__` (fast)
 - Read/write and communicate within a block: `__shared__` (fast and communication)
 - Read/write inside thread: registers (fast)
 - Data locality: texture

Resource limits

- ▶ Number of thread blocks per SM at the *same* time is limited by
 - Shared memory usage
 - Registers
 - No more than 8 thread blocks per SM
 - Number of threads

Resource limits - Examples

Number of blocks

How big should my blocks be? 8x8, 16x16 or 64x64?

- ▶ 8x8 Max threads per block 
 $8 \times 8 = 64$ threads per block, $1024/64 = 16$ blocks. An SM can have up to 8 blocks: only 512 threads will be active at the same time
- ▶ 16x16
 $16 \times 16 = 256$ threads per block, $1024/256 = 4$ blocks. An SM can take all blocks, then all 1024 threads will be active and achieve full capacity unless other resource overrule
- ▶ 64x64
 $64 \times 64 = 4096$ threads per block: doesn't fit in a SM

Resource limits - Examples

Registers

We have a kernel that uses 10 registers. With 16x16 block, how many blocks can run in G80 (max 8192 registers per SM, 768 threads per SM)?

$10 \times 16 \times 16 = 2560$. SM can hold $8192 / 2560 = 3$ blocks, meaning we will use $3 \times 16 \times 16 = 768$ threads, which is within limits.

If we add one more register, the number of registers grows to $11 \times 16 \times 16 = 2816$. SM can hold $8192 / 2816 = 2$ blocks, meaning we will use $2 \times 16 \times 15 = 512$ threads. Now as less threads are running per SM is more difficult to have enough warps to have the GPU always busy!

GPU Computing with CUDA

Lecture 3 - Efficient Shared Memory Use

*Christopher Cooper
Boston University*

*August, 2011
UTFSM, Valparaíso, Chile*

Shared Memory - Making use of it

- ▶ Looking at a 1D FDM example (similar to lab)

$$\frac{\partial u}{\partial t} = c \frac{\partial u}{\partial x} \longrightarrow u_i^{n+1} = u_i^n - \frac{c\Delta t}{\Delta x} (u_i^n - u_{i-1}^n)$$

```
__global__ void update (float *u, float *u_prev, int N, float dx, float dt, float c, int BLOCKSIZE)
{
    // Each thread will load one element
    int i = threadIdx.x + BLOCKSIZE * blockIdx.x;

    if (i>=N){return;}

    // u_prev[i] = u[i] is done in separate kernel

    if (i>0)
    {
        u[i] = u_prev[i] - c*dt/dx*(u_prev[i] - u_prev[i-1]);
    }
}
```

Shared Memory - Making use of it

- ▶ Looking at a 1D FDM example

$$\frac{\partial u}{\partial t} = c \frac{\partial u}{\partial x} \longrightarrow u_i^{n+1} = u_i^n - \frac{c\Delta t}{\Delta x} (u_i^n - u_{i-1}^n)$$

```
__global__ void update (float *u, float *u_prev, int N, float dx, float dt, float c, int  
BLOCKSIZE)  
{  
    // Each thread will load one element  
    int i = threadIdx.x + BLOCKSIZE * blockIdx.x;  
  
    if (i>=N){return;}  
  
    // u_prev[i] = u[i] is done in separate kernel  
  
    if (i>0)  
    {        u[i] = u_prev[i] - c*dt/dx*(u_prev[i] - u_prev[i-1]);  
    }  
}
```

Thread i

Shared Memory - Making use of it

- ▶ Looking at a 1D FDM example

$$\frac{\partial u}{\partial t} = c \frac{\partial u}{\partial x} \longrightarrow u_i^{n+1} = u_i^n - \frac{c\Delta t}{\Delta x} (u_i^n - u_{i-1}^n)$$

```
__global__ void update (float *u, float *u_prev, int N, float dx, float dt, float c, int BLOCKSIZE)
{
    // Each thread will load one element
    int i = threadIdx.x + BLOCKSIZE * blockIdx.x;
    Thread i
    Loads
    // u_prev[i] = u[i] is done in separate kernel element i
    if (i>0)
    {
        u[i] = u_prev[i] - c*dt/dx*(u_prev[i] - u_prev[i-1]);
    }
}
```

Shared Memory - Making use of it

- ▶ Looking at a 1D FDM example

$$\frac{\partial u}{\partial t} = c \frac{\partial u}{\partial x} \longrightarrow u_i^{n+1} = u_i^n - \frac{c\Delta t}{\Delta x} (u_i^n - u_{i-1}^n)$$

```
__global__ void update (float *u, float *u_prev, int N, float dx, float dt, float c, int BLOCKSIZE)
{
    // Each thread will load one element
    int i = threadIdx.x + BLOCKSIZE * blockIdx.x;
    // u_prev[i] = u[i] is done in separate kernel
    if (i>0)
    {
        u[i] = u_prev[i] - c*dt/dx*(u_prev[i] - u_prev[i-1]);
    }
}
```

The diagram illustrates the redundant loads in the CUDA kernel. It shows two overlapping ovals, each containing the text "Loads element i". Red arrows point from these ovals to the two occurrences of `u_prev[i]` in the assignment statement `u[i] = u_prev[i] - c*dt/dx*(u_prev[i] - u_prev[i-1]);`. This highlights that each thread performs two redundant memory loads for every element `i > 0`.

Order N redundant loads!

Shared Memory - Making use of it

- ▶ Looking at a 1D FDM example

$$\frac{\partial u}{\partial t} = c \frac{\partial u}{\partial x} \longrightarrow u_i^{n+1} = u_i^n - \frac{c\Delta t}{\Delta x} (u_i^n - u_{i-1}^n)$$

```
__global__ void update (float *u, float *u_prev, int N, float dx, float dt, float c, int  
BLOCKSIZE)  
{  
    // Each thread will load one element  
    int i = threadIdx.x + BLOCKSIZE * blockIdx.x;  
  
    if (i>=N){return;}  
  
    // u_prev[i] = u[i] is done in separate kernel  
  
    if (i>0)  
    {        u[i] = u_prev[i] - c*dt/dx*(u_prev[i] - u_prev[i-1]);  
    }  
}
```

Thread i+1

Shared Memory - Making use of it

- ▶ Looking at a 1D FDM example

$$\frac{\partial u}{\partial t} = c \frac{\partial u}{\partial x} \longrightarrow u_i^{n+1} = u_i^n - \frac{c\Delta t}{\Delta x} (u_i^n - u_{i-1}^n)$$

```
__global__ void update (float *u, float *u_prev, int N, float dx, float dt, float c, int BLOCKSIZE)
{
    // Each thread will load one element
    int i = threadIdx.x + BLOCKSIZE * blockIdx.x; Thread i+1
    // u_prev[i] = u[i] is done in separate kernel
    if (i>0)
    {
        u[i] = u_prev[i] - c*dt/dx*(u_prev[i] - u_prev[i-1]);
    }
}
```

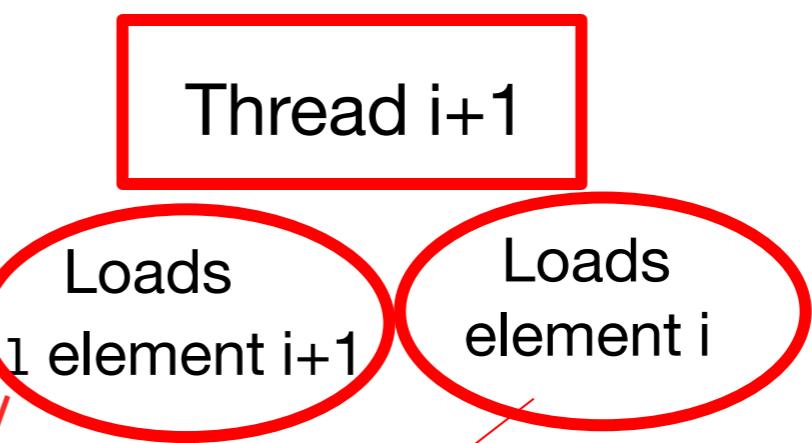
Loads element $i+1$

Shared Memory - Making use of it

- ▶ Looking at a 1D FDM example

$$\frac{\partial u}{\partial t} = c \frac{\partial u}{\partial x} \longrightarrow u_i^{n+1} = u_i^n - \frac{c\Delta t}{\Delta x} (u_i^n - u_{i-1}^n)$$

```
__global__ void update (float *u, float *u_prev, int N, float dx, float dt, float c, int BLOCKSIZE)
{
    // Each thread will load one element
    int i = threadIdx.x + BLOCKSIZE * blockIdx.x;
    // u_prev[i] = u[i] is done in separate kernel
    if (i>0)
    {
        u[i] = u_prev[i] - c*dt/dx*(u_prev[i] - u_prev[i-1]);
    }
}
```



Order N redundant loads!

Shared Memory - Making use of it

- ▶ Idea: We could load only once to shared memory, and operate there

```
__global__ void update (float *u, float *u_prev, int N, float dx, float dt, float c)
{
    // Each thread will load one element
    int i = threadIdx.x;
    int I = threadIdx.x + BLOCKSIZE * blockIdx.x;
    __shared__ float u_shared[BLOCKSIZE];

    if (I>=N){return;}

    u_shared[i] = u[I];
    __syncthreads();
    if (I>0)
    {
        u[I] = u_shared[i] - c*dt/dx*(u_shared[i] - u_shared[i-1]);
    }
}
```

Shared Memory - Making use of it

- Idea: We could load only once to shared memory, and operate there

```
__global__ void update (float *u, float *u_prev, int N, float dx, float dt, float c)
{
    // Each thread will load one element
    int i = threadIdx.x;
    int I = threadIdx.x + BLOCKSIZE * blockIdx.x;
    __shared__ float u_shared[BLOCKSIZE]; ←
    if (I>=N){return;}
    u_shared[i] = u[I]; ←
    __syncthreads(); ←
    if (I>0)
    {
        u[I] = u_shared[i] - c*dt/dx*(u_shared[i] - u_shared[i-1]);
    }
}
```

The code is annotated with three red boxes and arrows:

- A red box labeled "Allocate shared array" surrounds the declaration of the shared memory array `__shared__ float u_shared[BLOCKSIZE];`. An arrow points from the left margin to the start of the declaration.
- A red box labeled "Load to shared mem" surrounds the assignment `u_shared[i] = u[I];`. An arrow points from the left margin to the start of the assignment.
- A red box labeled "Fetch shared mem" surrounds the term `u_shared[i-1]` in the update assignment. An arrow points from the left margin to the start of the term.

Shared Memory - Making use of it

- Idea: We could load only once to shared memory, and operate there

```
__global__ void update (float *u, float *u_prev, int N, float dx, float dt, float c)
{
    // Each thread will load one element
    int i = threadIdx.x;
    int I = threadIdx.x + BLOCKSIZE * blockIdx.x;
    __shared__ float u_shared[BLOCKSIZE]; ← Allocate shared array

    if (I>=N){return;}

    u_shared[i] = u[I]; ← Load to shared mem
    __syncthreads();      51
    if (I>0)
    {      u[I] = u_shared[i] - c*dt/dx*(u_shared[i] - u_shared[i-1]);      }
}
```

Allocate shared array

Load to shared mem

Fetch shared mem

Works if $N \leq$ Block size... What if not?

Boundary!

Shared Memory - Making use of it

```
__global__ void update (float *u, float *u_prev, int N, float dx, float dt, float c)
{
    // Each thread will load one element
    int i = threadIdx.x;
    int I = threadIdx.x + BLOCKSIZE * blockIdx.x;
    __shared__ float u_shared[BLOCKSIZE];

    if (I>=N){return;}

    u_shared[i] = u[I];
    __syncthreads();

    if (i>0 && i<BLOCKSIZE-1)
    {
        u[I] = u_shared[i] - c*dt/dx*(u_shared[i] - u_shared[i-1]);
    }
    else
    {
        u[I] = u_prev[I] - c*dt/dx*(u_prev[I] - u_prev[I-1]);
    }
}
```

52

Reduced loads from 2^*N to $N+2^*N/BLOCKSIZE$

Using shared memory as cache

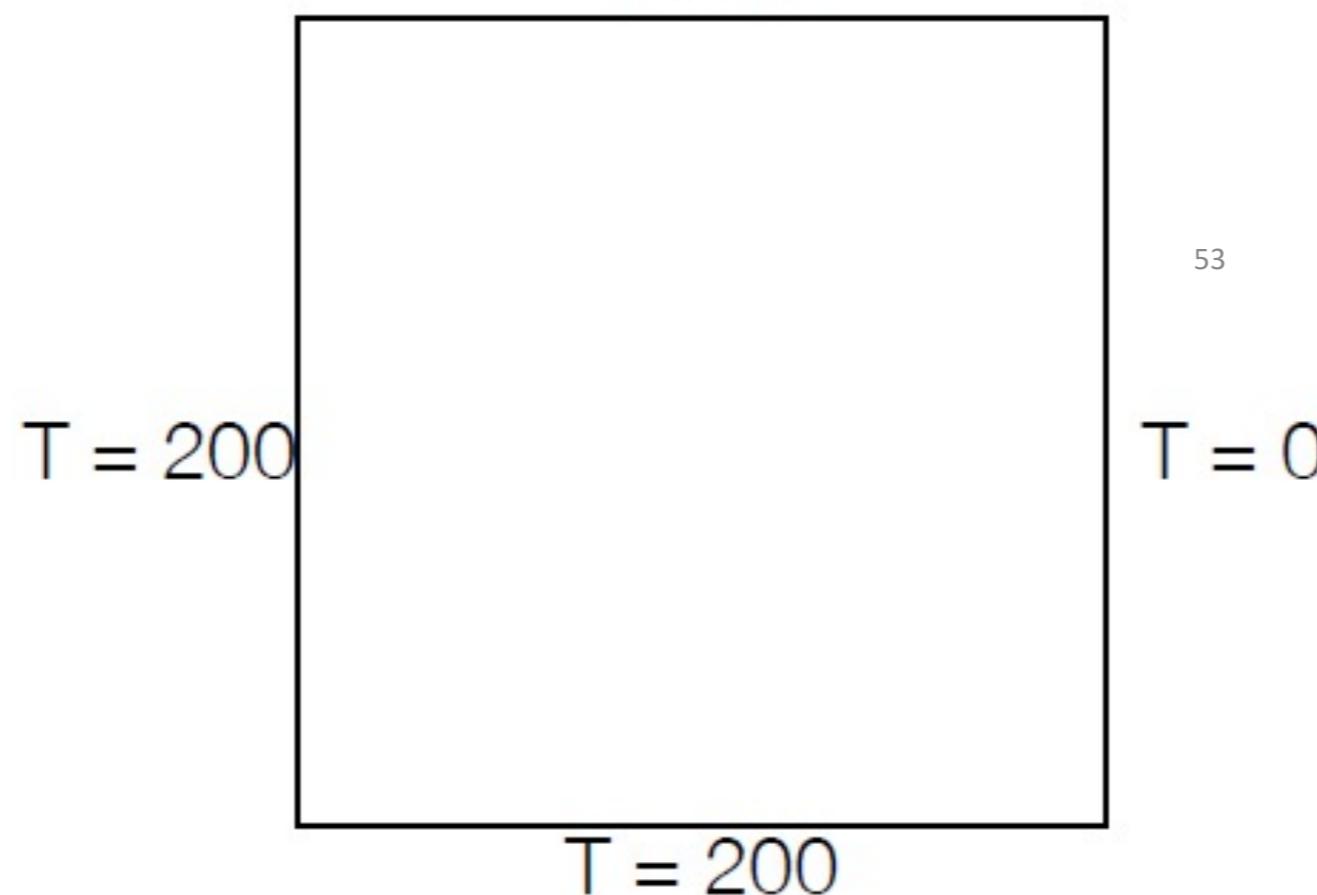
- ▶ Looking at the 2D heat diffusion problem from lab 2

$$\frac{\partial u}{\partial t} = \alpha \nabla^2 u$$

- ▶ Explicit scheme

$$u_{i,j}^{n+1} = u_{i,j}^n + \frac{\alpha k}{h^2} (u_{i,j+1}^n + u_{i,j-1}^n + u_{i+1,j}^n + u_{i-1,j}^n - 4u_{i,j}^n)$$

$T = 0$

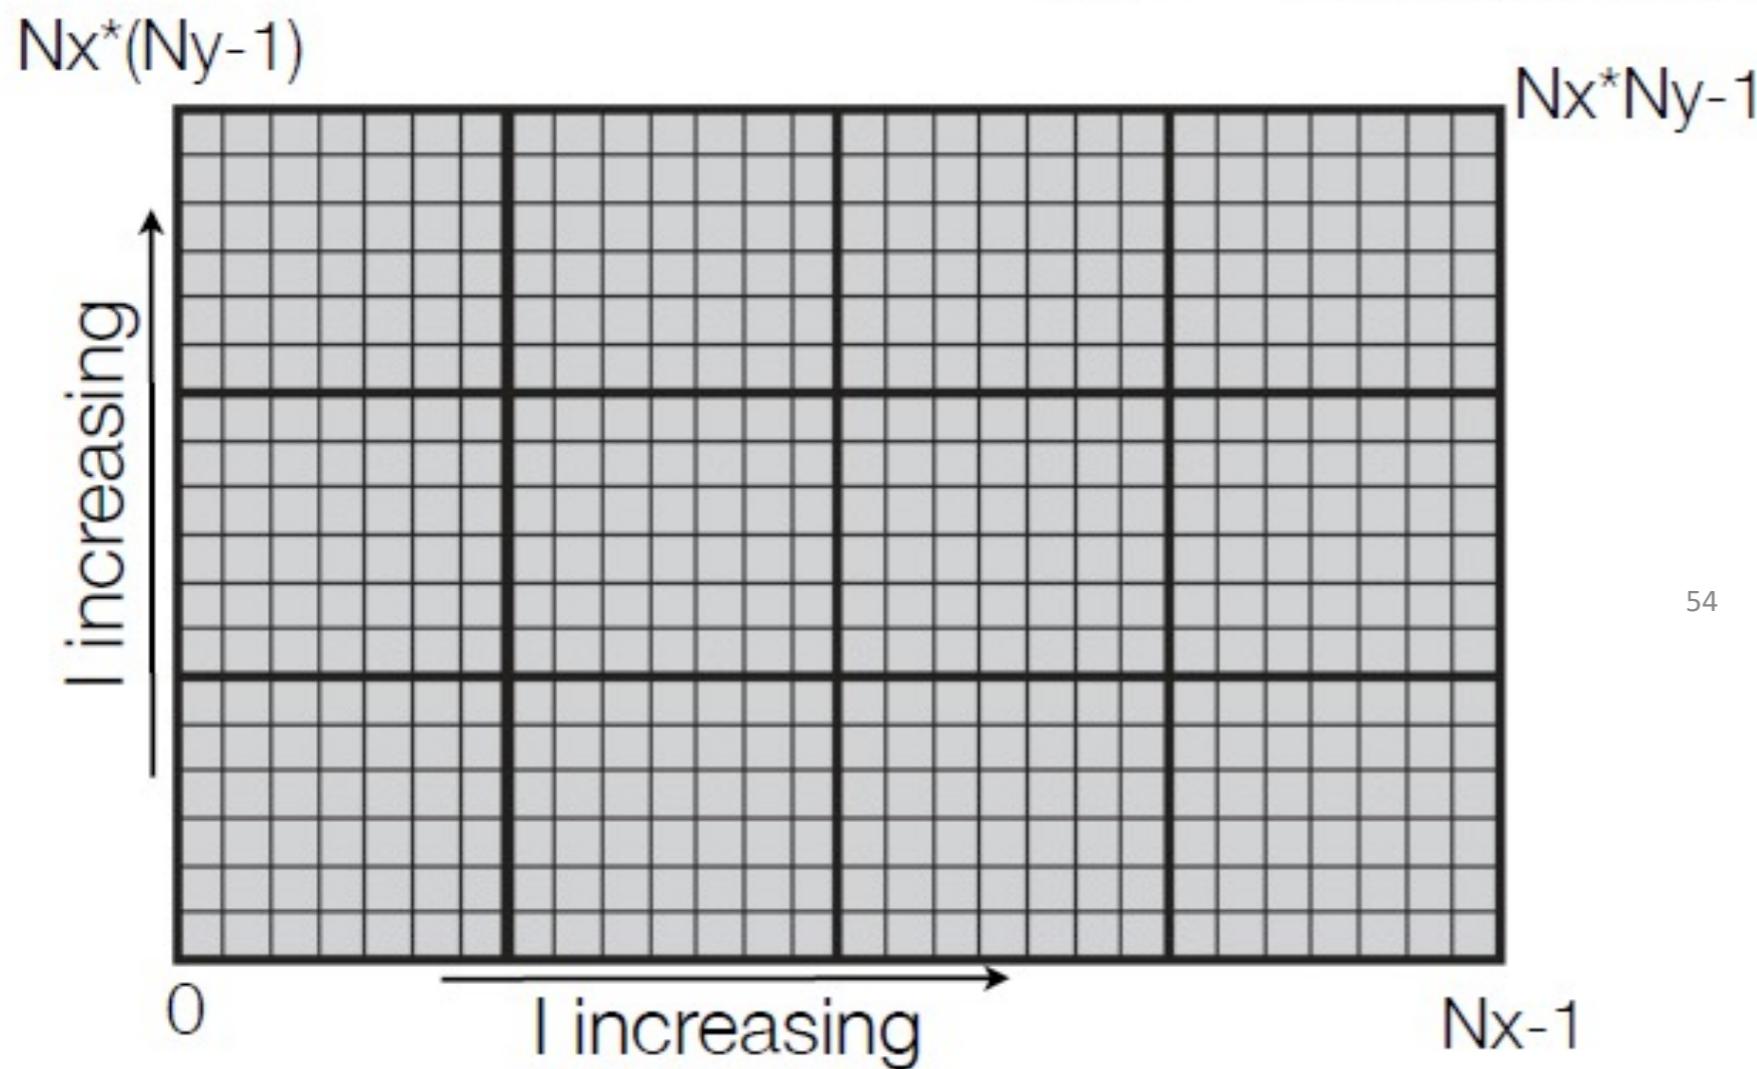


53

Shared Memory Implementation - Mapping Problem

- ▶ Using row major flattened array

```
int i = threadIdx.x;
int j = threadIdx.y;
int I = blockIdx.y*BSZ*N + blockIdx.x*BSZ + j*N + i;
```



54

Shared Memory Implementation - Global Memory

- ▶ This implementation has redundant loads to global memory → slow

```
__global__ void update (float *u, float *u_prev, int N, float h, float dt, float alpha, int BSZ)
{
    // Setting up indices
    int i = threadIdx.x;
    int j = threadIdx.y;
    int I = blockIdx.y*BSZ*N + blockIdx.x*BSZ + j*N + i;

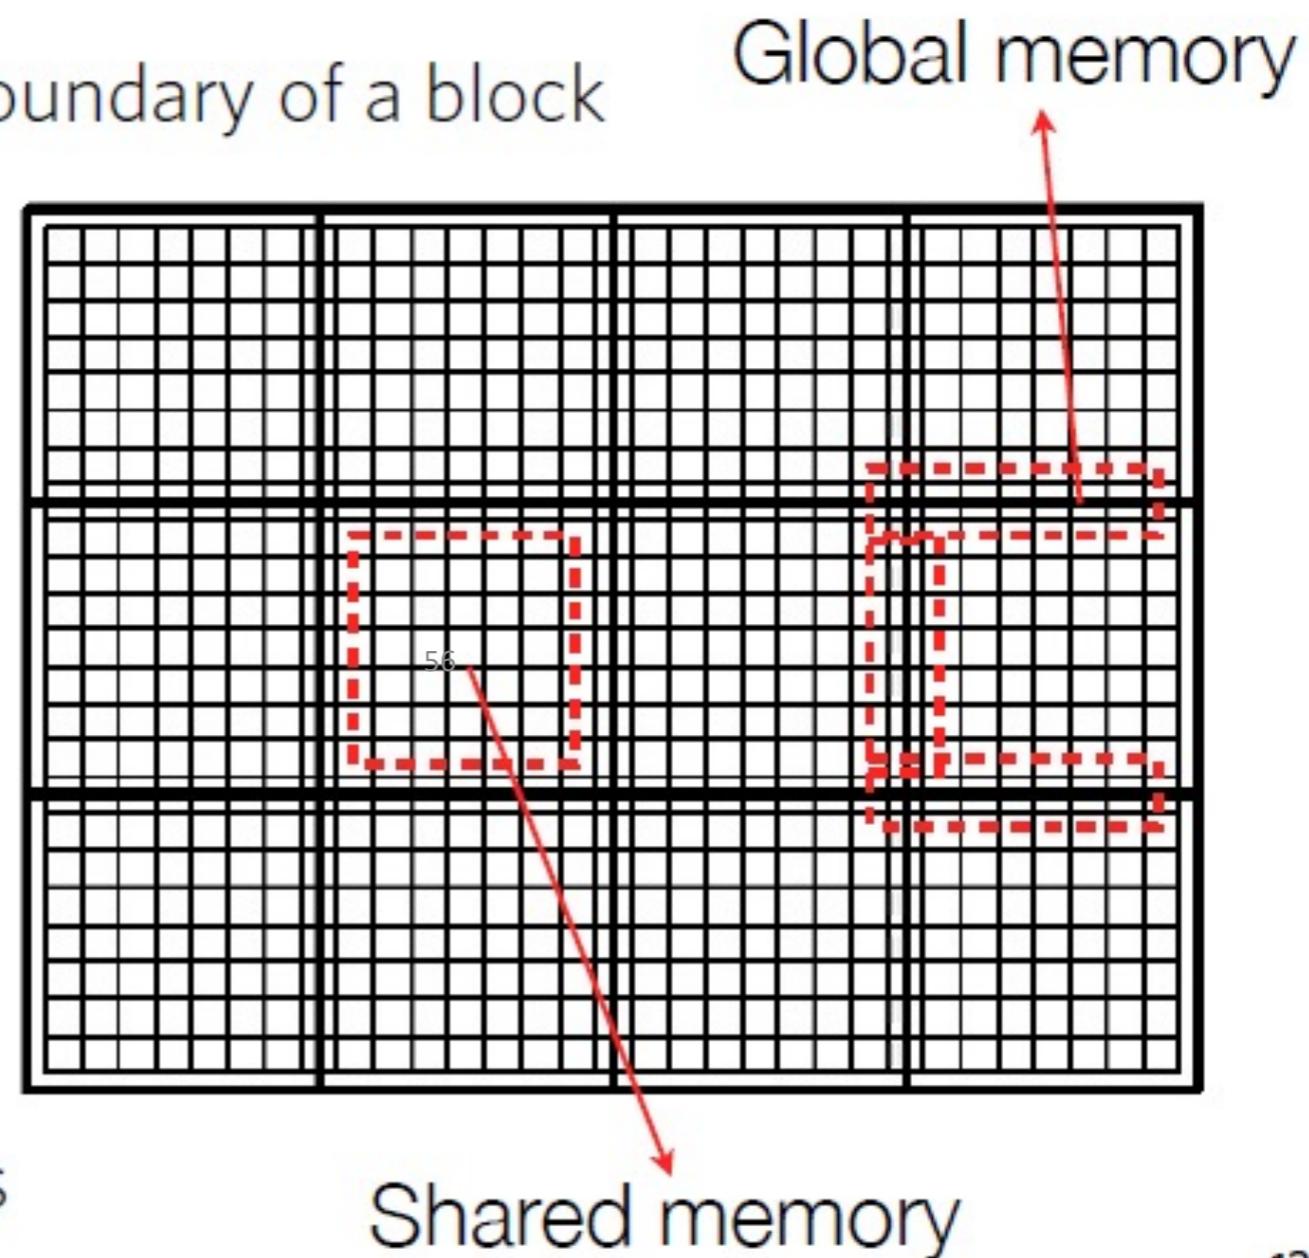
    if (I>=N*N){return;}

    // if not boundary do
    if ( (I>N) && (I< N*N-1-N) && (I%N!=0) && (I%N!=N-1))
        {      u[I] = u_prev[I] + alpha*dt/(h*h) * (u_prev[I+1] + u_prev[I-1] +
u_prev[I+N] + u_prev[I-N] - 4*u_prev[I]);}
}
```

55

Shared Memory Implementation - Solution 1

- ▶ Recast solution given earlier
 - Load to shared memory
 - Use shared memory if not on boundary of a block
 - Use global memory otherwise
- ▶ Advantage
 - Easy to implement
- ▶ Disadvantage
 - Branching statement
 - Still have some redundant loads



Shared Memory Implementation - Solution 1

```
__global__ void update (float *u, float *u_prev, int N, float h, float dt, float alpha)
{
    // Setting up indices
    int i = threadIdx.x;
    int j = threadIdx.y;
    int I = blockIdx.y*BSZ*N + blockIdx.x*BSZ + j*N + i;
    if (I>=N*N){return;}

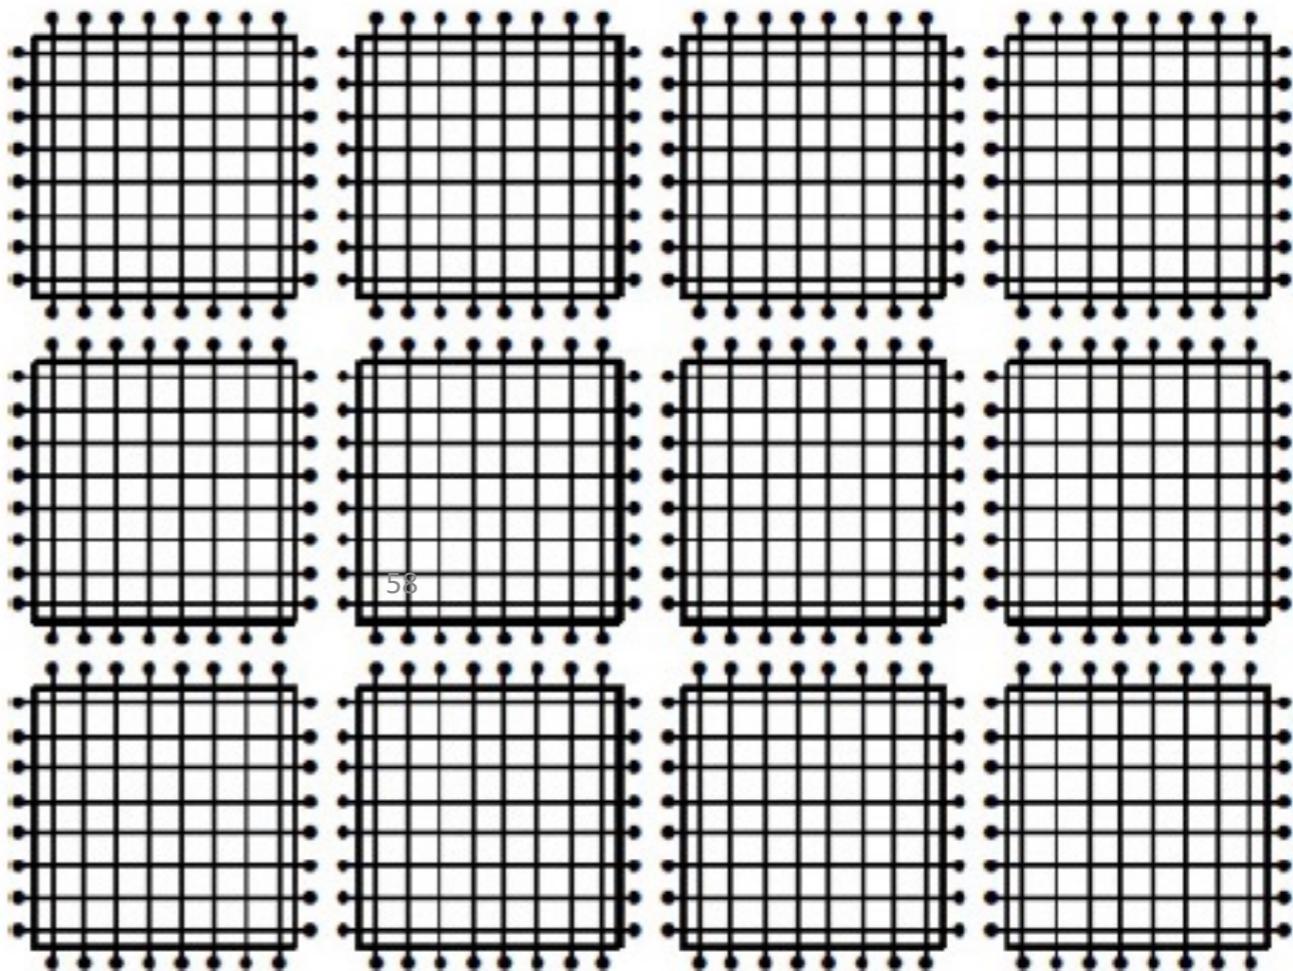
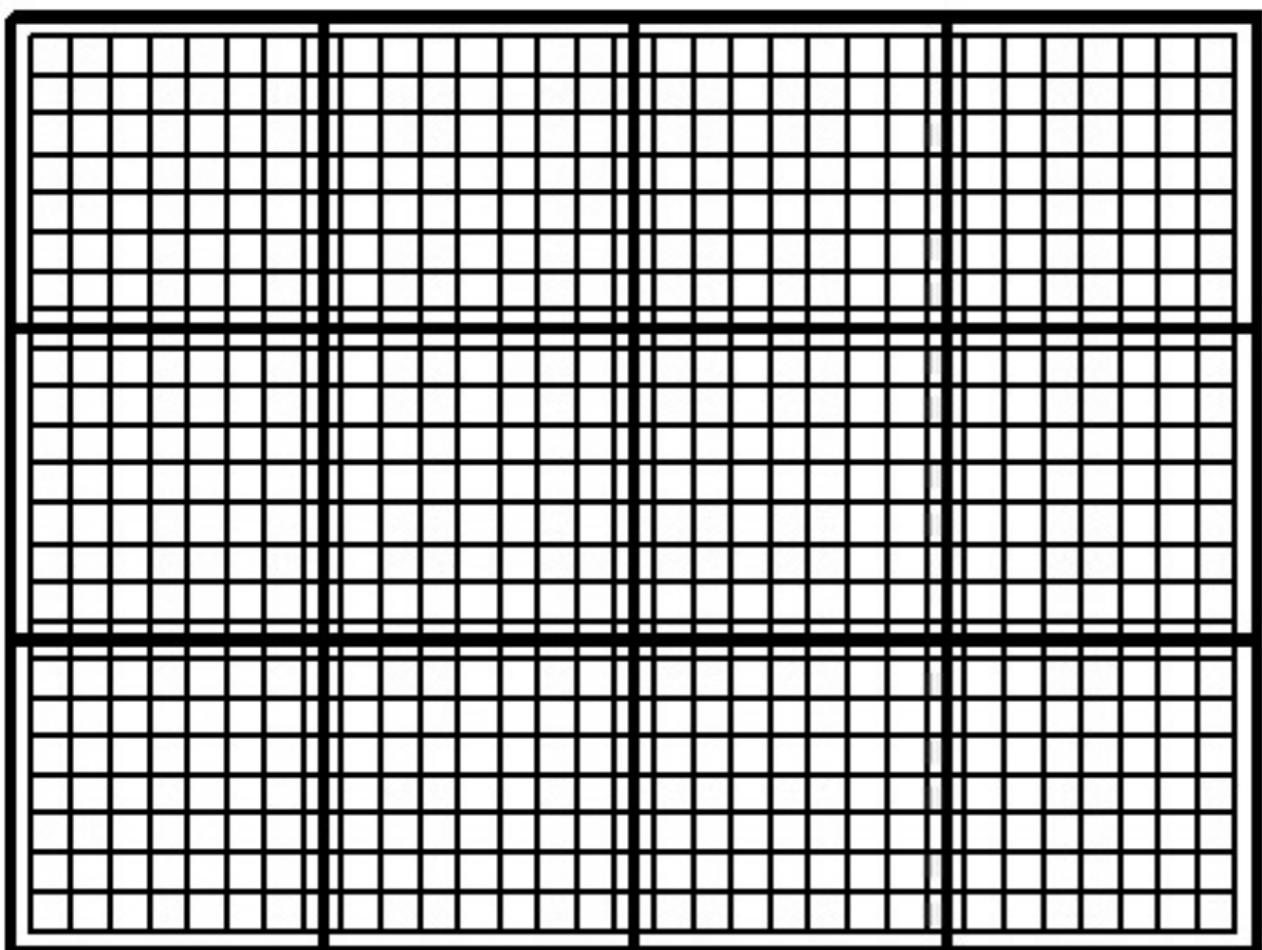
    __shared__ float u_prev_sh[BSZ][BSZ];
    u_prev_sh[i][j] = u_prev[I];

    __syncthreads();
    bool bound_check = ((I>N) && (I< N*N-1-N) && (I%N!=0) && (I%N!=N-1));
    bool block_check = ((i>0) && (i<BSZ-1) && (j>0) && (j<BSZ-1));

    // if not on block boundary do
    if (block_check)
    {
        u[I] = u_prev_sh[i][j] + alpha*dt/h/h * (u_prev_sh[i+1][j] + u_prev_sh[i-1]
[j] + u_prev_sh[i][j+1] + u_prev_sh[i][j-1] - 4*u_prev_sh[i][j]);
    }
    // if not on boundary
    else if (bound_check)
    {
        u[I] = u_prev[I] + alpha*dt/(h*h) * (u_prev[I+1] + u_prev[I-1] + u_prev[I+N]
+ u_prev[I-N] - 4*u_prev[I]);
    }
}
```

Shared Memory Implementation - Solution 2

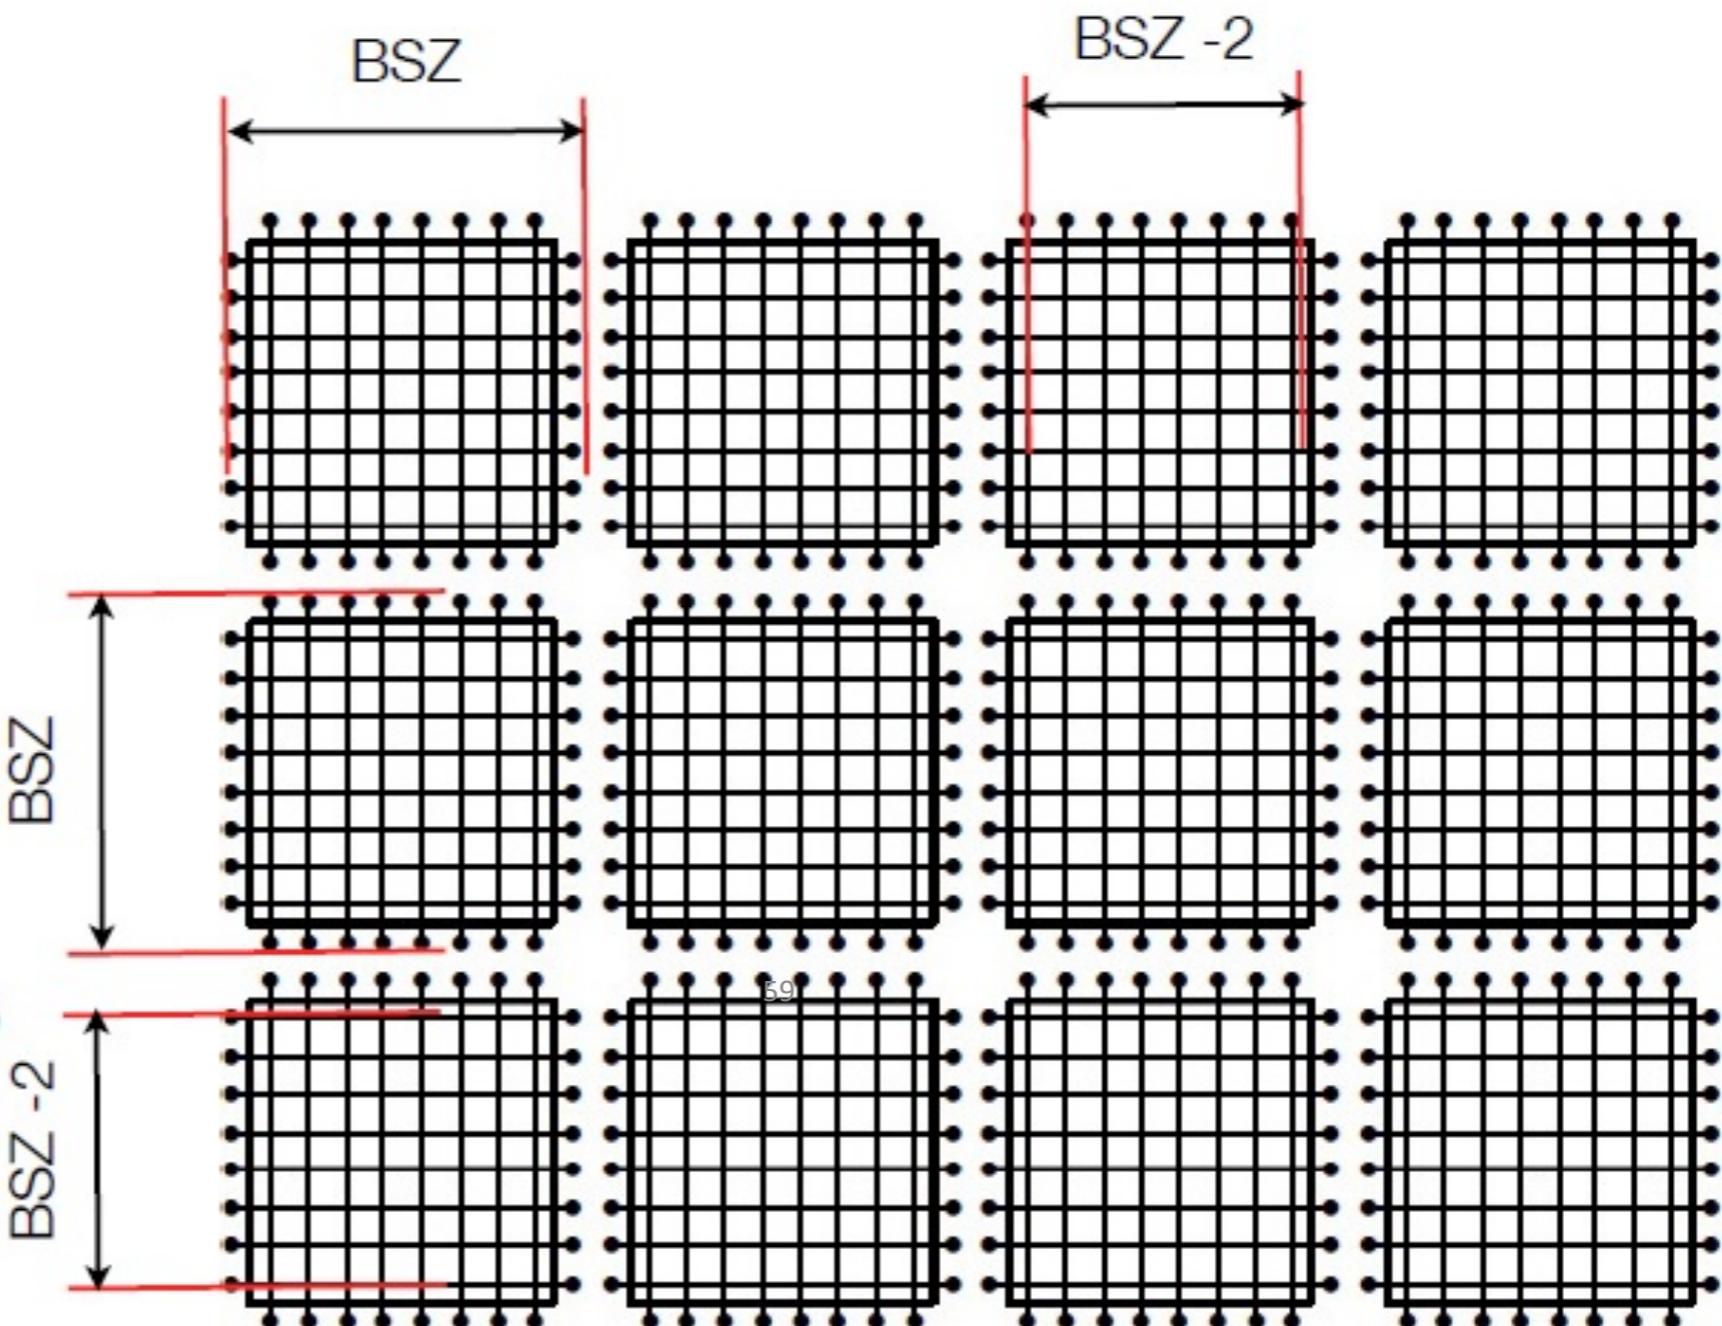
- ▶ We want to avoid the reads from global memory
 - Let's use halo nodes to compute block edges



Images: Mark Giles, Oxford, UK

Shared Memory Implementation - Solution 2

- ▶ Change indexing so to jump in steps of BSZ-2 instead of BSZ
- ▶ Load data to shared memory
- ▶ Operate on internal nodes
- ▶ We'll need $Nx/(BSZ-2)$ blocks per dimension, instead of Nx/BSZ



Shared Memory Implementation - Solution 2

```
__global__ void update (float *u, float *u_prev, int N, float h, float dt, float alpha)
{
    // Setting up indices
    int i = threadIdx.x, j = threadIdx.y, bx = blockIdx.x, by = blockIdx.y;

    int I = (BSZ-2)*bx + i, J = (BSZ-2)*by + j;
    int Index = I + J*N;

    if (I>=N || J>=N){return;}

    __shared__ float u_prev_sh[BSZ][BSZ];

    u_prev_sh[i][j] = u_prev[Index];

    __syncthreads();
```

60

```
    bool bound_check = ((I!=0) && (I<N-1) && (J!=0) && (J<N-1));
    bool block_check = ((i!=0) && (i<BSZ-1) && (j!=0) && (j<BSZ-1));

    if (bound_check && block_check)
    {
        u[Index] = u_prev_sh[i][j] + alpha*dt/h/h * (u_prev_sh[i+1][j] +
u_prev_sh[i-1][j] + u_prev_sh[i][j+1] + u_prev_sh[i][j-1] - 4*u_prev_sh[i][j]);
    }
}
```

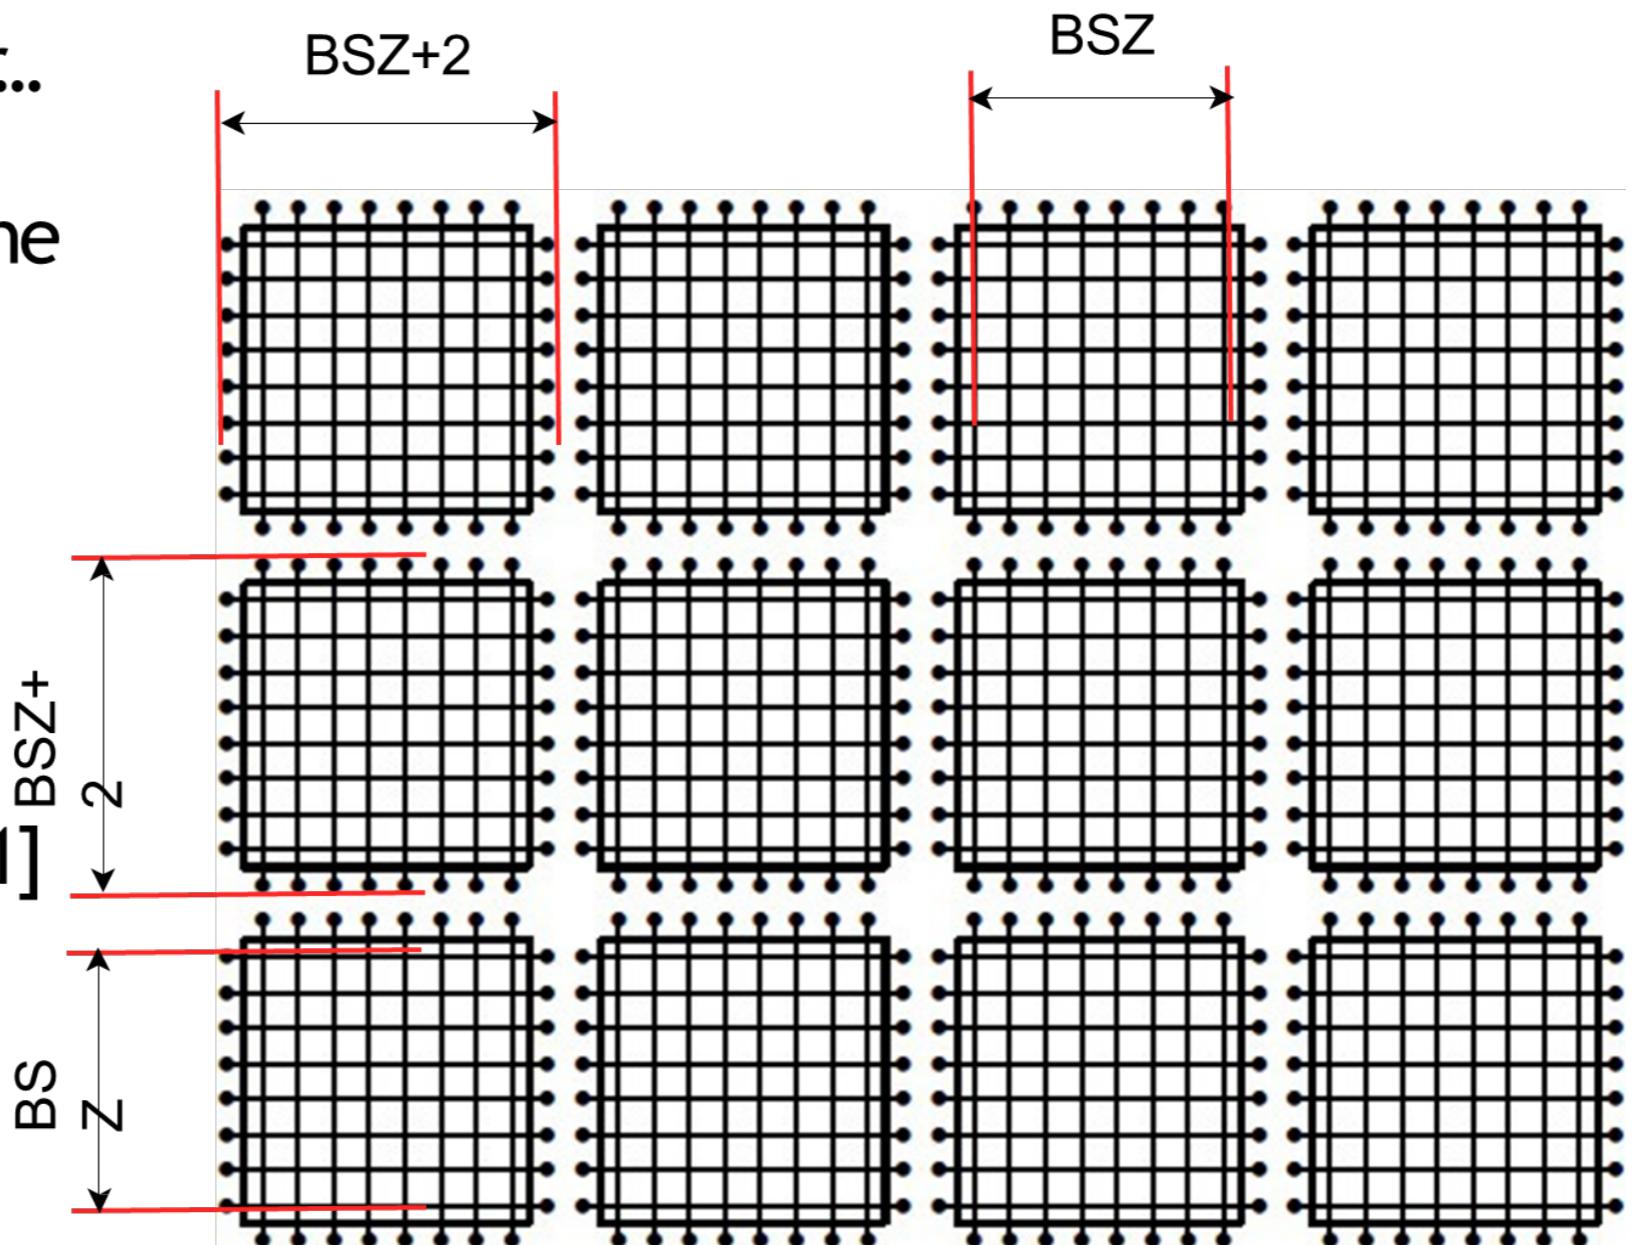
Shared Memory Implementation - Solution 2

- ▶ We've reduced global memory accesses!
- ▶ But...
 - There's still a heavy amount of branching
 - ▶ GPUs are not great at branching...
 - All threads read, but only some operate
 - ▶ We're underutilizing the device!61
 - ▶ If we have $16 \times 16 = 256$ threads, all read, but only $14 \times 14 = 196$ operate, and we're using only ~75% of the device. In 3D this number drops to ~40%!

Shared Memory Implementation - Solution 3

- We need to go further...

- To not underutilize the device, we need to load more data than threads
- Load in two stages
- Operate on $[i+1][j+1]$ threads



Shared Memory Implementation - Solution 3

› Loading in 2 steps

- Use the 64 available threads to load the 64 first values to shared

```
__shared__ float u_prev_sh[BSZ+2][BSZ+2];

int ii = j*BSZ + i, // Flatten thread indexing
    I = ii%(BSZ+2), // x-direction index including halo
    J = ii/(BSZ+2); // y-direction index including halo

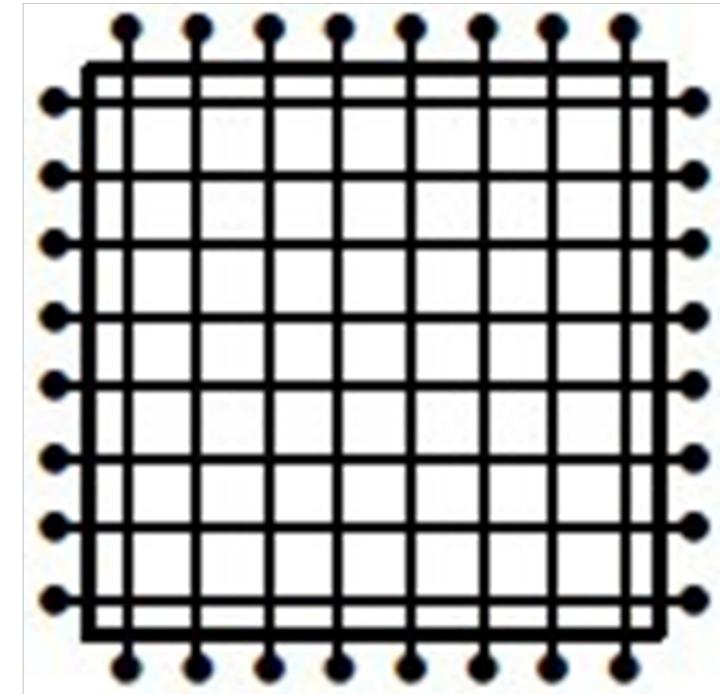
int I_n = I_0 + J*N + I; //General index
u_prev_sh[I][J] = u_prev[I_n];
```

- Load the remaining values

```
int ii2 = BSZ*BSZ + j*BSZ + i;
int I2  = ii2%(BSZ+2);
int J2  = ii2/(BSZ+2);

int I_n2 = I_0 + J2*N + I2; //General index

if ( (I2<(BSZ+2)) && (J2<(BSZ+2)) && (ii2 < N*N) )
    u_prev_sh[I2][J2] = u[I_n2];
```



63

8x8 threads
10x10 loads

Shared Memory Implementation - Solution 3

▶ Loading in 2 steps

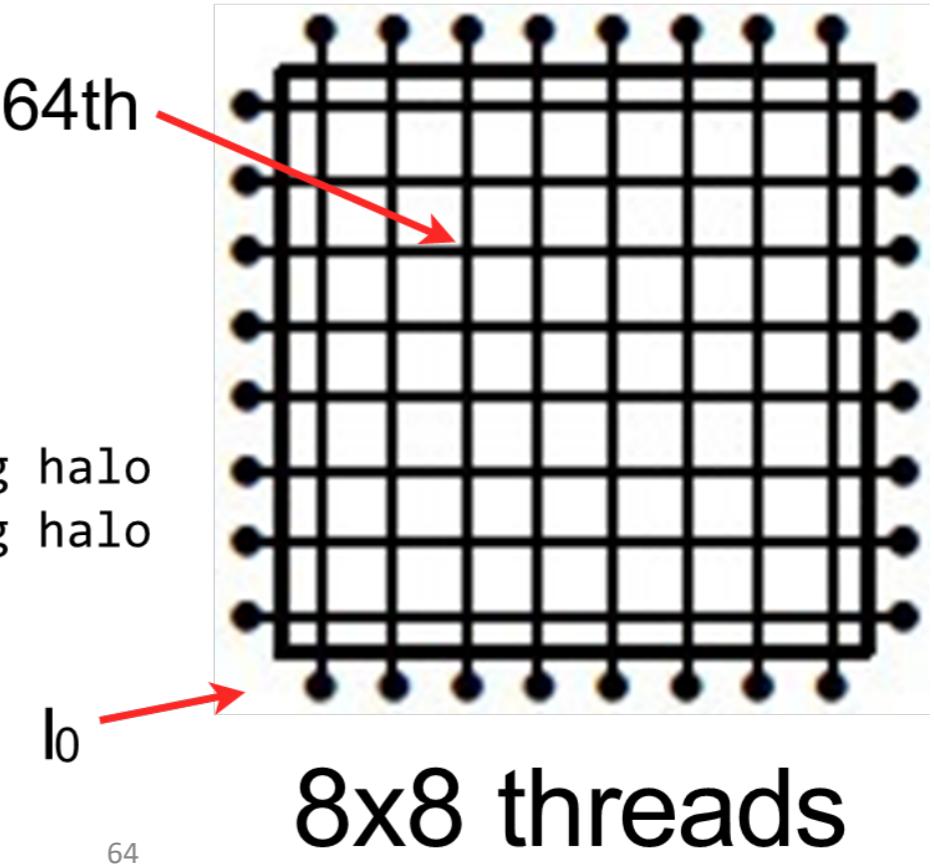
- Use the 64 available threads to load the 64 first values to shared

```
__shared__ float u_prev_sh[BSZ+2][BSZ+2];  
  
int ii = j*BSZ + i, // Flatten thread indexing  
    I = ii%(BSZ+2), // x-direction index including halo  
    J = ii/(BSZ+2); // y-direction index including halo
```

```
int I_n = I_0 + J*N + I; //General index  
u_prev_sh[I][J] = u_prev[I_n];
```

- Load the remaining values

```
int ii2 = BSZ*BSZ + j*BSZ + i;  
int I2  = ii2%(BSZ+2);  
int J2  = ii2/(BSZ+2);  
  
int I_n2 = I_0 + J2*N + I2; //General index  
  
if ( (I2<(BSZ+2)) && (J2<(BSZ+2)) && (ii2 < N*N) )  
    u_prev_sh[I2][J2] = u[I_n2];
```



8x8 threads
10x10 loads

Shared Memory Implementation - Solution 3

- ▶ Loading in 2 steps

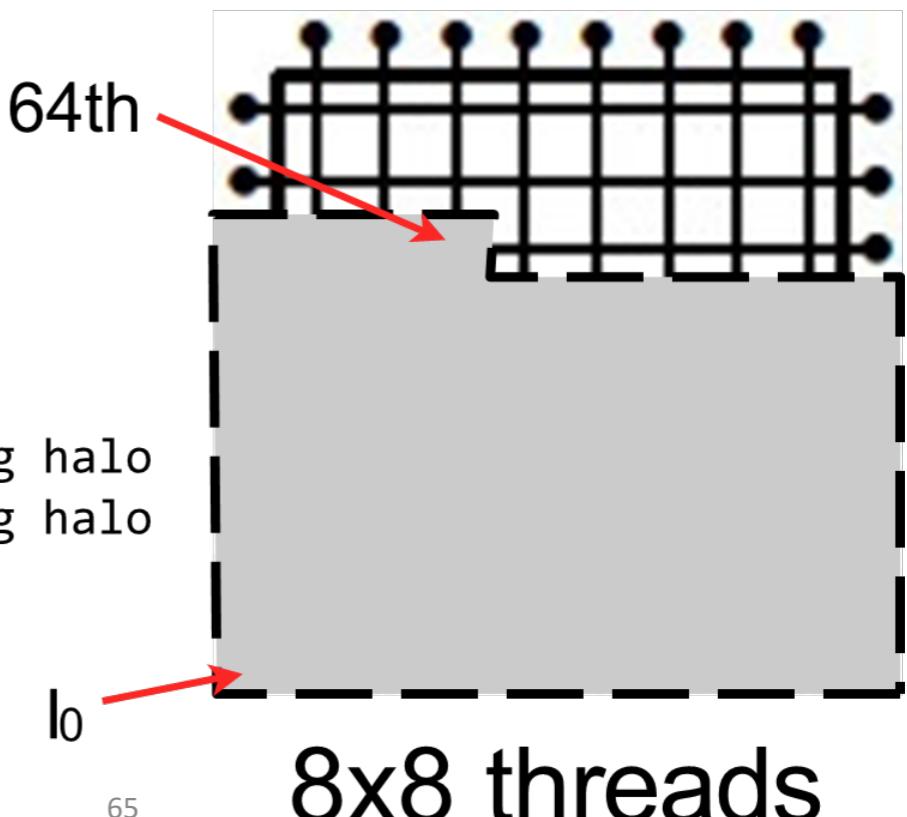
- Use the 64 available threads to load the 64 first values to shared

```
__shared__ float u_prev_sh[BSZ+2][BSZ+2];  
  
int ii = j*BSZ + i, // Flatten thread indexing  
    I = ii%(BSZ+2), // x-direction index including halo  
    J = ii/(BSZ+2); // y-direction index including halo
```

```
int I_n = I_0 + J*N + I; //General index  
u_prev_sh[I][J] = u_prev[I_n];
```

- Load the remaining values

```
int ii2 = BSZ*BSZ + j*BSZ + i;  
int I2  = ii2%(BSZ+2);  
int J2  = ii2/(BSZ+2);  
  
int I_n2 = I_0 + J2*N + I2; //General index  
  
if ( (I2<(BSZ+2)) && (J2<(BSZ+2)) && (ii2 < N*N) )  
    u_prev_sh[I2][J2] = u[I_n2];
```



Shared Memory Implementation - Solution 3

▶ Loading in 2 steps

- Use the 64 available threads to load the 64 first values to shared

```
__shared__ float u_prev_sh[BSZ+2][BSZ+2];
```

```

int ii = j*BSZ + i, // Flatten thread indexing
      I = ii%(BSZ+2), // x-direction index including halo
      J = ii/(BSZ+2); // y-direction index including halo

```

```
int I_n = I_0 + J*N + I; //General index  
u_prev_sh[I][J] = u_prev[I_n];
```

- Load the remaining values

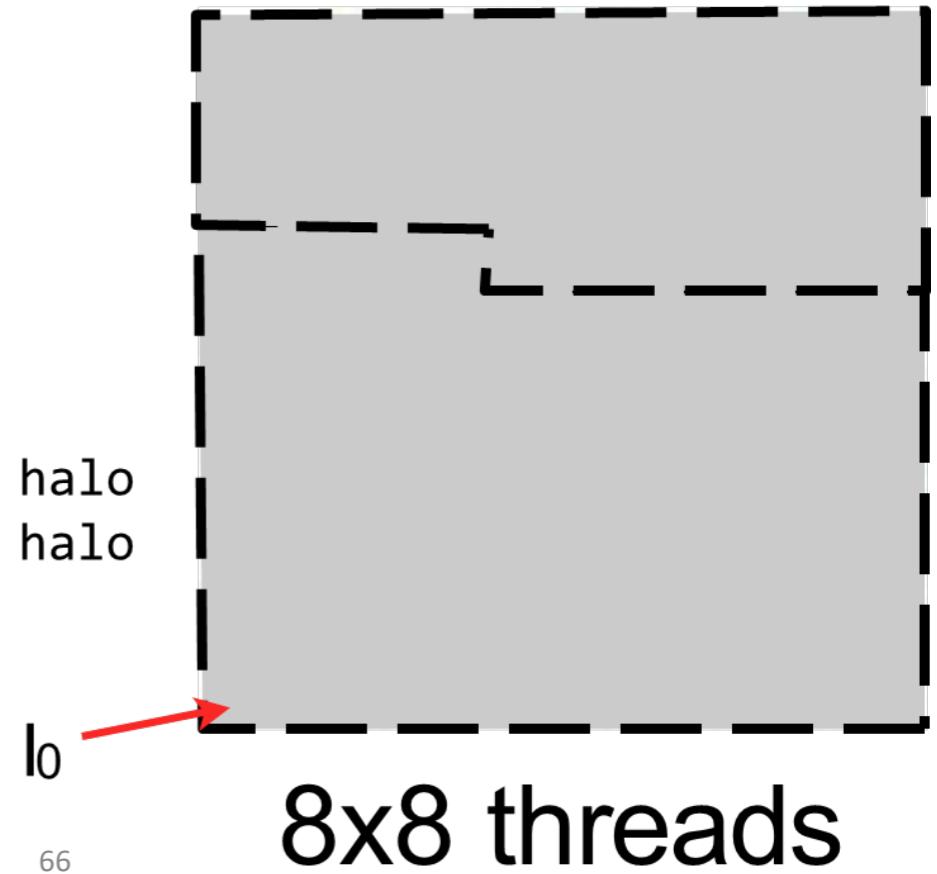
```
int ii2 = BSZ*BSZ + j*BSZ + i;
```

```
int I2 = ii%(BSZ+2);
```

```
int J2 = ii2/(BSZ+2);
```

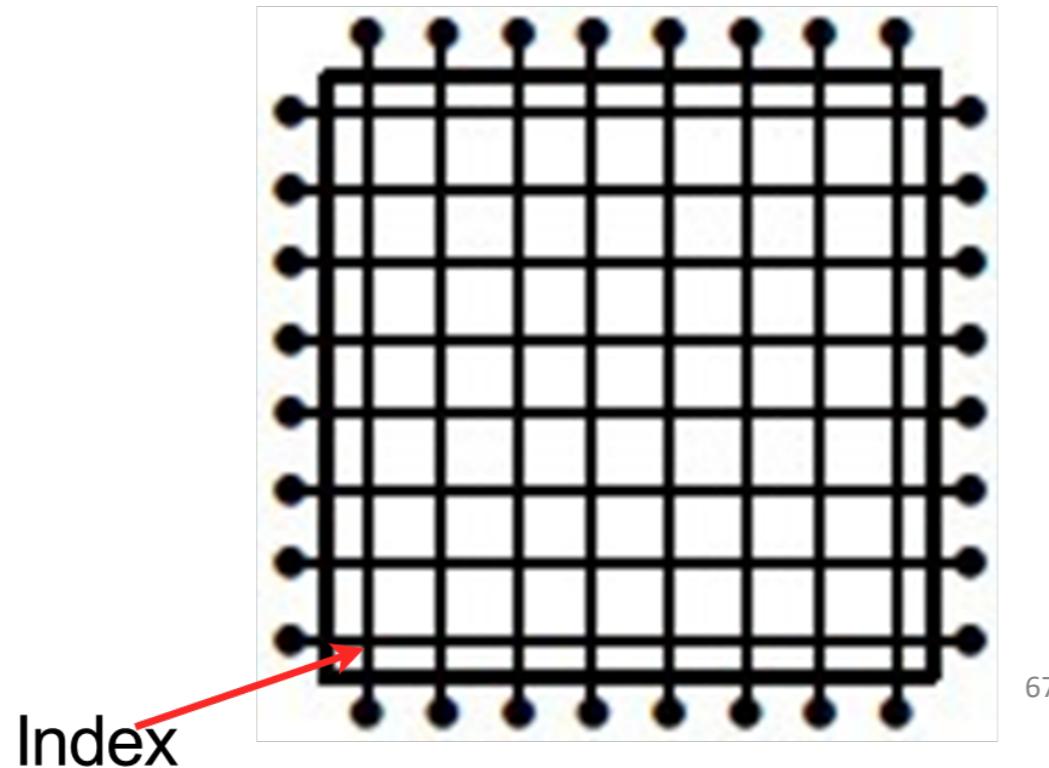
```
int I_n2 = I_0 + J2*N + I2; //General index
```

```
if ( (I2<(BSZ+2)) && (J2<(BSZ+2)) && (ii2 < N*N) ) ← Some threads won't load  
    u_prev_sh[I2][J2] = u[I_n2];
```



Shared Memory Implementation - Solution 3

- Compute on interior points: threads [i+1][j+1]



67

```
int Index = by*BSZ*N + bx*BSZ + (j+1)*N + i+1;  
  
u[Index] = u_prev_sh[i+1][j+1] + alpha*dt/h/h * (u_prev_sh[i+2][j+1] + u_prev_sh[i][j+1] +  
u_prev_sh[i+1][j+2] + u_prev_sh[i+1][j] - 4*u_prev_sh[i+1][j+1]);
```

SM Implementation

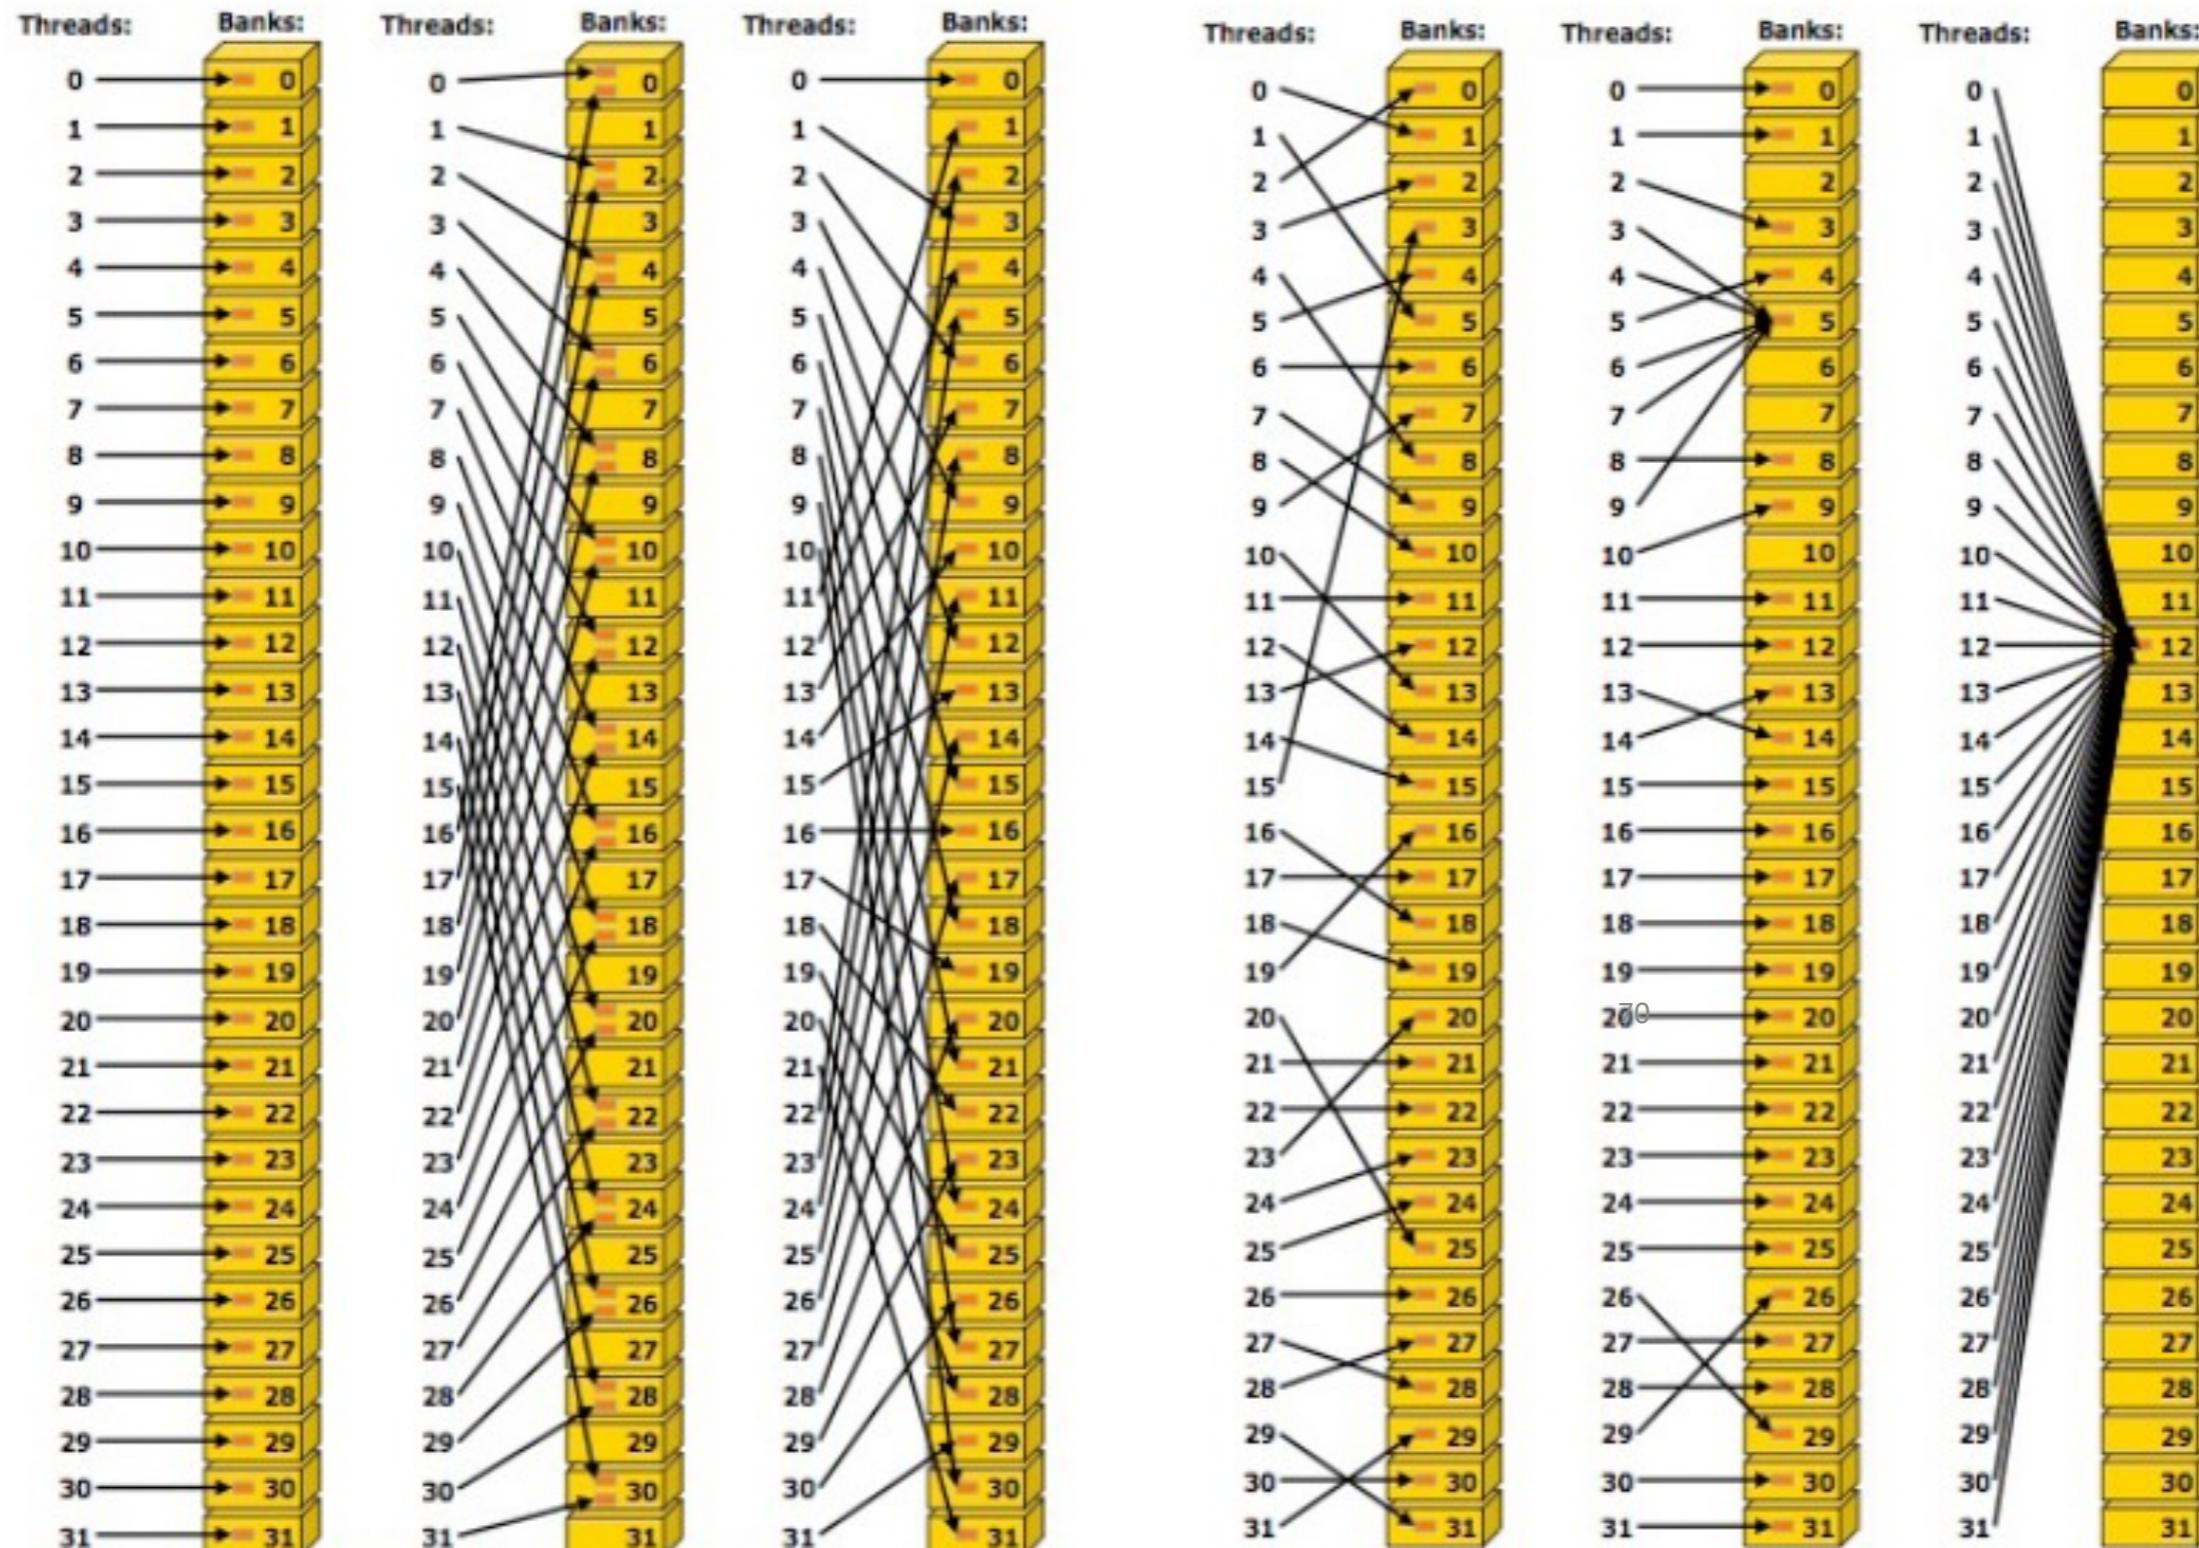
- The technique described is called **tiling**
 - Tiling means loading data to shared memory in tiles
 - Useful when shared memory is used as cache
 - Also used when all data is too large to fit in shared memory and you load it in smaller chunks

Shared Memory - Bank conflicts

- ▶ Shared memory arrays are subdivided into smaller subarrays called **banks**
- ▶ Shared memory has 32 (16) banks in 2.X (1.X). Successive 32-bit words are assigned to successive banks
- ▶ Different banks can be accessed simultaneously
- ▶ If two or more addresses of a memory request are in the same bank, the access is serialized
 - Bank conflicts exist only within a warp (half warp for 1.X)
- ▶ In 2.X there is no bank conflict if the memory request is for the same 32-bit word. This is not valid in 1.X.

69

Shared Memory - Bank conflicts



Left: Linear addressing with a stride of one 32-bit word (no bank conflict).

Middle: Linear addressing with a stride of two 32-bit words (2-way bank conflicts).

Right: Linear addressing with a stride of three 32-bit words (no bank conflict).

Left: Conflict-free access via random permutation.

Middle: Conflict-free access since threads 3, 4, 6, 7, and 9 access the same word within bank 5.

Right: Conflict-free broadcast access (all threads access the same word).

Shared Memory

► `__syncthreads()`

- Barrier that waits for all threads of the block before continuing
- Need to make sure all data is loaded to shared before access
- Avoids **race conditions**
- Don't over use!

```
u_shared[i] = u[I];
__syncthreads();

if (i>0 && i<BLOCKSIZE-1)
    u[I] = u_shared[i] - c*dt/dx*(u_shared[i] - u_shared[i-1]);
```

71

Race condition

- ▶ When two or more threads want to access and operate on a memory location without synchronization
- ▶ Example: we have the value 3 stored in global memory and two threads want to add one to that value.
 - Possibility 1:
 - ▶ Thread 1 reads the value 3 adds 1 and writes 4 back to memory
 - ▶ Thread 2 reads the value 4 and writes 5 back to memory
 - Possibility 2:
 - ▶ Thread 1 reads the value 3
 - ▶ Thread 2 reads the value 3
 - ▶ Both threads operate on 3 and write back the value 4 to memory
- ▶ Solutions:
 - `__syncthreads()` or atomic operations

Race condition

- ▶ When two or more threads want to access and operate on a memory location without synchronization
- ▶ Example: we have the value 3 stored in global memory and two threads want to add one to that value.
 - Possibility 1:
 - ▶ Thread 1 reads the value 3 adds 1 and writes 4 back to memory
 - ▶ Thread 2 reads the value 4 and writes 5 back to memory ✓
 - Possibility 2:
 - ▶ Thread 1 reads the value 3
 - ▶ Thread 2 reads the value 3
 - ▶ Both threads operate on 3 and write back the value 4 to memory ✗
- ▶ Solutions:
 - `__syncthreads()` or atomic operations

Atomic operations

- ▶ Atomic operations deal with race conditions
 - It guarantees that while the operation is being executed, that location in memory is not accessed
 - Still we can't rely on any ordering of thread execution!
 - Types
 - atomicAdd
 - atomicSub
 - atomicExch
 - atomicMin
 - atomicMax
 - etc...

Atomic operations

```
__global__ update (int *values, int *who)
{
    int i = threadIdx.x + blockDim.x*blockIdx.x;
    int I = who[i];

    atomicAdd(&values[I], 1);
}
```

Atomic operations

- ▶ Useful if you have a sparse access pattern
- ▶ Atomic operations are slower than “normal” function
- ▶ They can serialize your execution if many threads want to access the same memory location
 - Think about parallelizing your data, not only execution
 - Use hierarchy of atomic operations to avoid this
- ▶ Prefer `__syncthreads()` if you can use it instead
 - If you have a regular access pattern