
Lecture 8

CUDA Programming 3

张奇
复旦大学

Outline

Optimization & Example

CUDA Libraries

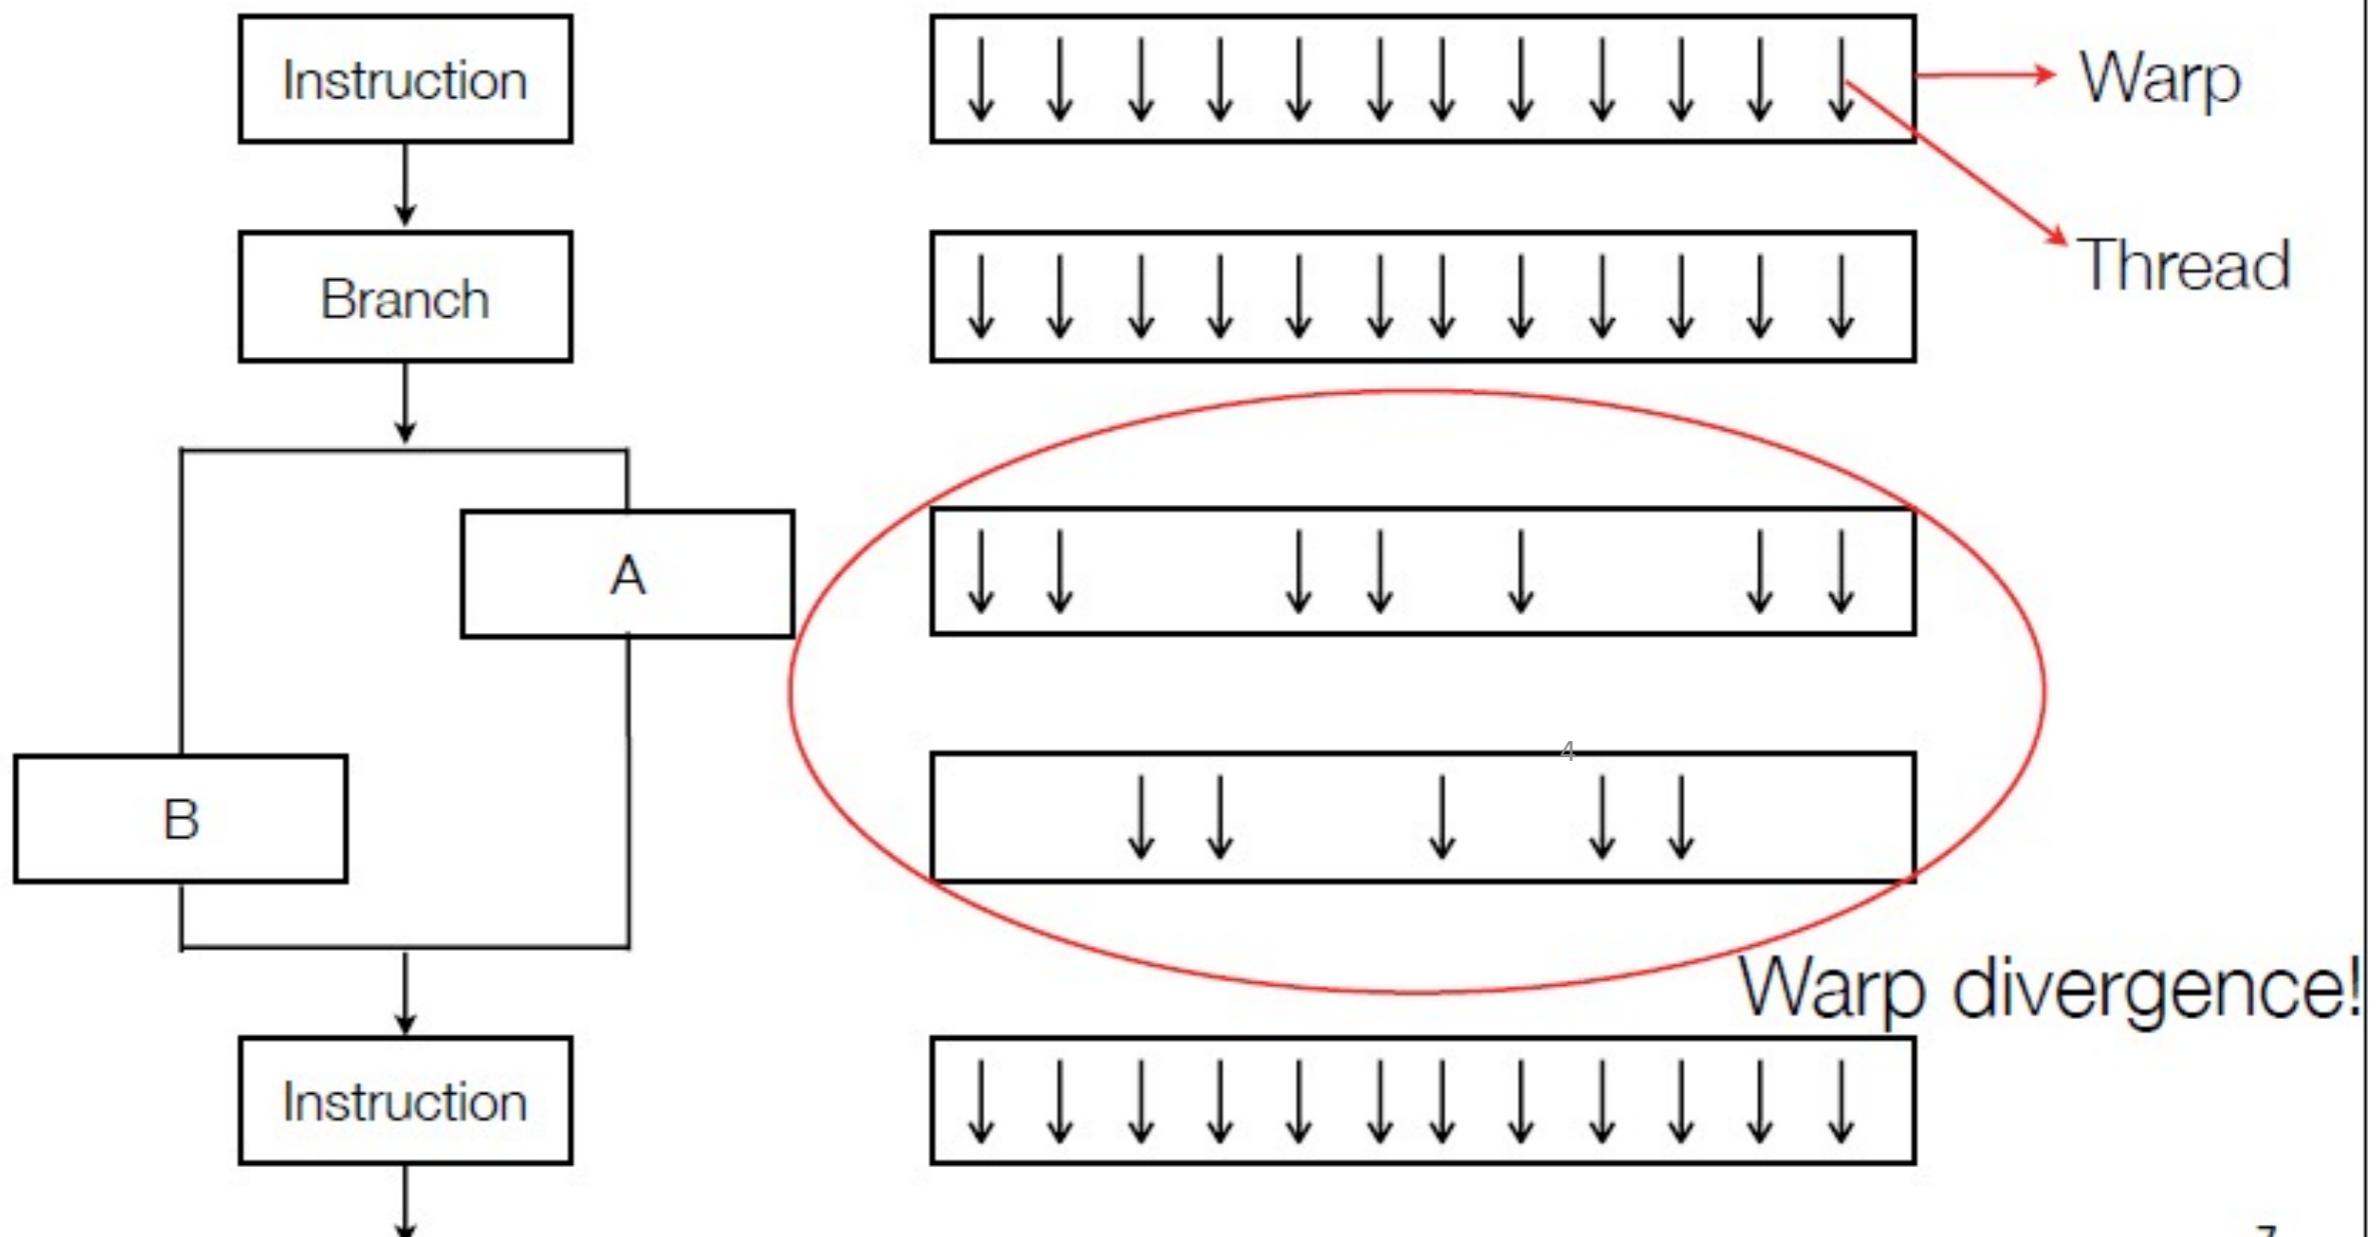
Control flow

- ▶ If statement
 - Threads are executed in warps
 - Within a warp, the hardware is not capable of executing *if* and *else* statements at the same time!

```
__global__ void function();  
{  
    ....  
  
    if (condition)            3  
    {  
        ...  
    }  
    else  
    {  
        ...  
    }  
}
```

Control flow

- ▶ How does the hardware deal with an if statement?



Control flow

- ▶ Hardware serializes the different execution paths
- ▶ Recommendations
 - Try to make every thread in the same warp do the same thing
 - ▶ If the if statement cuts at a multiple of the warp size, there is no warp divergence and the instruction can be done in one pass
 - Remember threads are placed consecutively in a warp (t0-t31, t32-t63, ...)
 - ▶ But we cannot rely on any execution order within warps
 - If you can't avoid branching, try to make as many consecutive threads as possible do the same thing

5

8

Control flow - An illustration

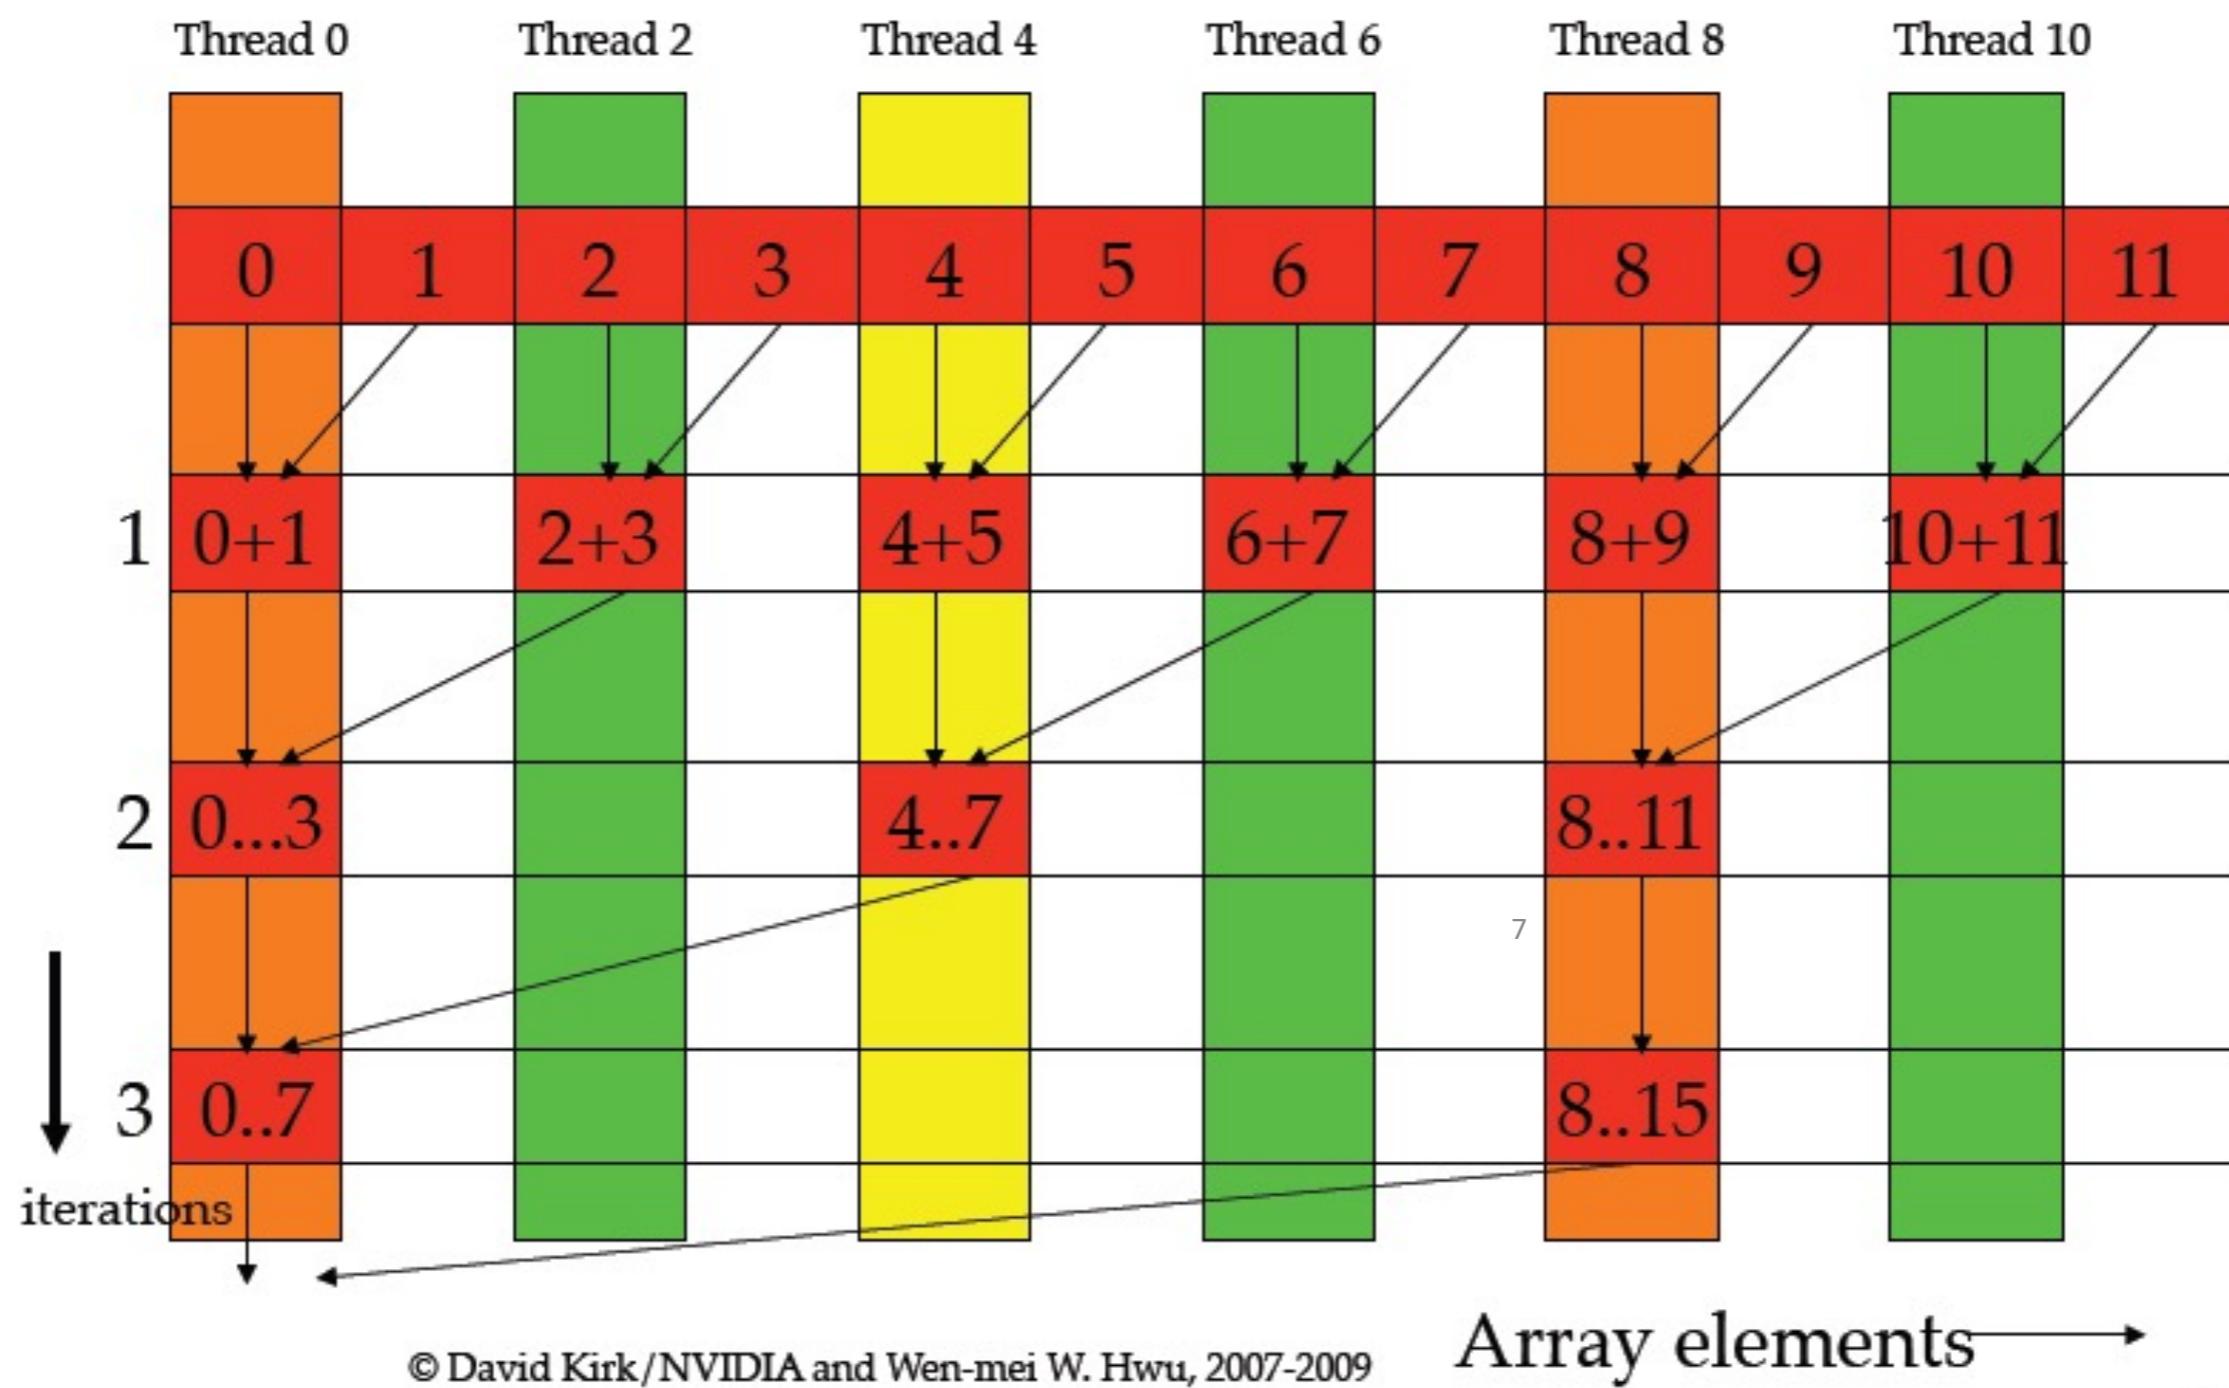
- ▶ Let's implement a sum reduction
- ▶ In a serial code, we would just loop over all elements and add
- ▶ Idea:
 - To implement in parallel, let's take every other value and add in place it to its neighbor
 - To the resulting array do the same thing until we have only one value

6

```
__shared__ float partialSum[]  
int t = threadIdx.x  
for(int stride = 1; stride<blockDim.x; stride*=2)  
{  
    __syncthreads();  
    if (t%(2*stride)==0)  
        partialSum[t] += partialSum[t+stride];  
}
```

9

Control flow - An illustration



Control flow - An illustration

- ▶ Advantages

- Right result
- Runs in parallel

- ▶ Disadvantages

- The number of threads decreases per iteration, but we're using much more warps than needed
 - ▶ There is warp divergence in every iteration⁸
- No more than half the threads per thread are being executed per iteration
- ▶ Let's change the algorithm a little bit...

Control flow - An illustration

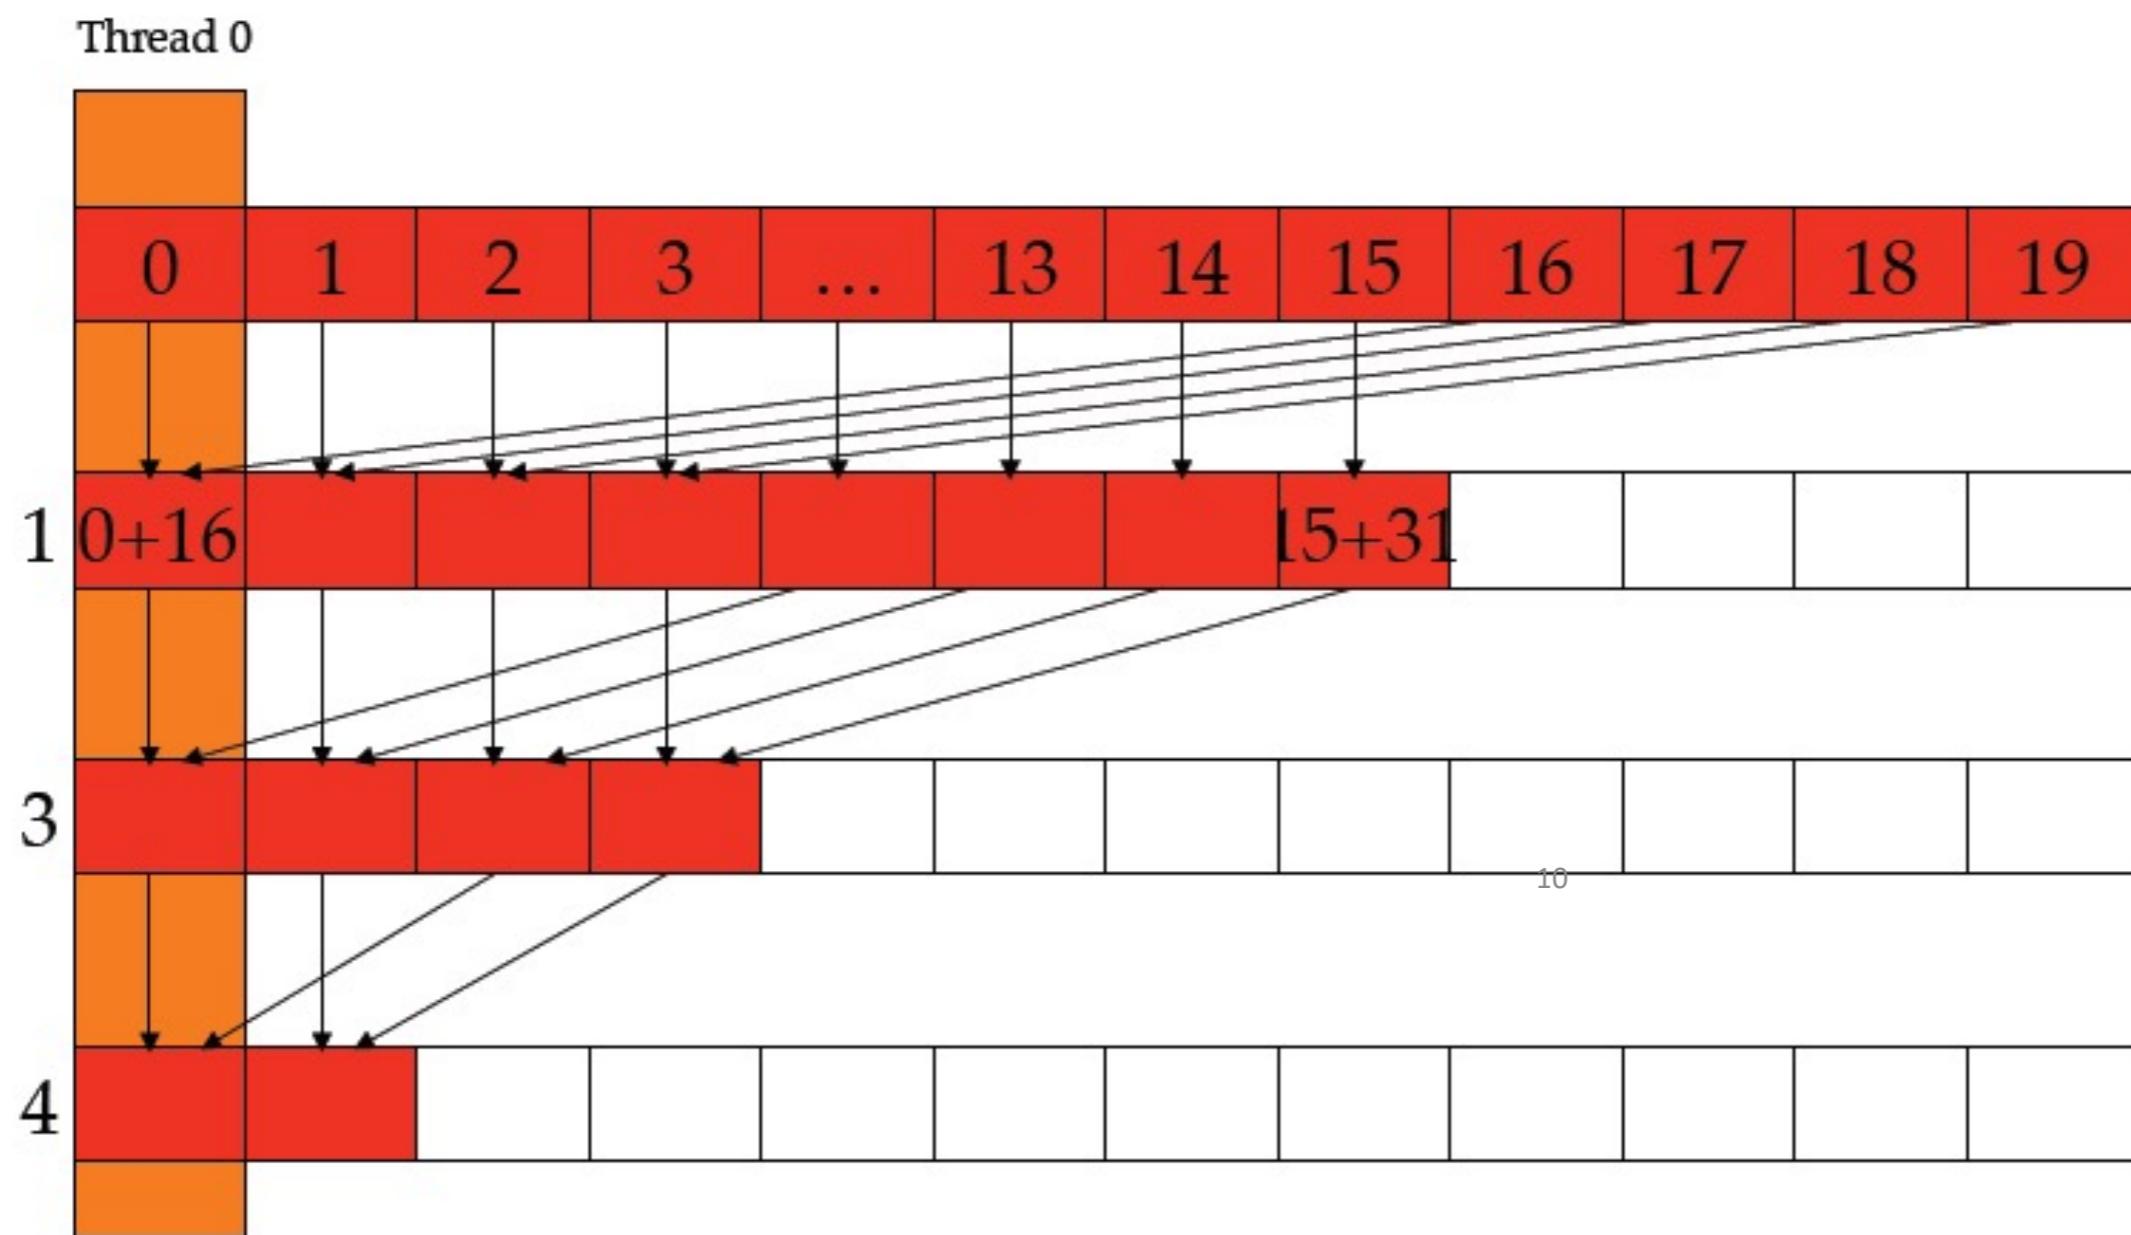
- ▶ Improved version

- Instead of adding neighbors, let's add values with stride half a section away
- Divide the stride by two after each iteration

```
__shared__ float partialSum[]  
int t = threadIdx.x  
for(int stride = blockDim.x; stride>1; stride>>=1)  
{  
    __syncthreads();  
    if (t<stride)  
        partialSum[t] += partialSum[t+stride];  
}
```

9

Control flow - An illustration



Control flow - An illustration

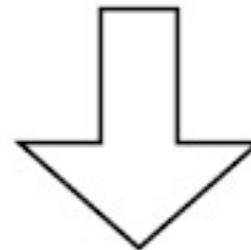
- ▶ 512 elements

Iteration	Exec. threads	Warp
1	256	16
2	128	8
3	64	4
4	32	2
5	16	1
6	8	1
7	4	1
8	2	1

Threads > Warp

Threads < Warp

11



Warp divergence!

Control flow - An illustration

- ▶ We get warp divergence only for the last 5 iterations
- ▶ Warps will be shut down as the iteration progresses
 - This will happen much faster than for the previous case
 - Resources are better utilized
 - For the last 5 iterations, only 1 warp is still active

Control flow - Loop divergence

- ▶ Work per thread data dependent

```
__global__ void per_thread_sum (int *indices, float *data, float *sums)
{
    ...
    for (int j=indices[i]; j<indices[i+1]; j++)
    {
        sum += data[j];
    }
    sums[i] = sum
}
```

13

Control flow - Loop divergence

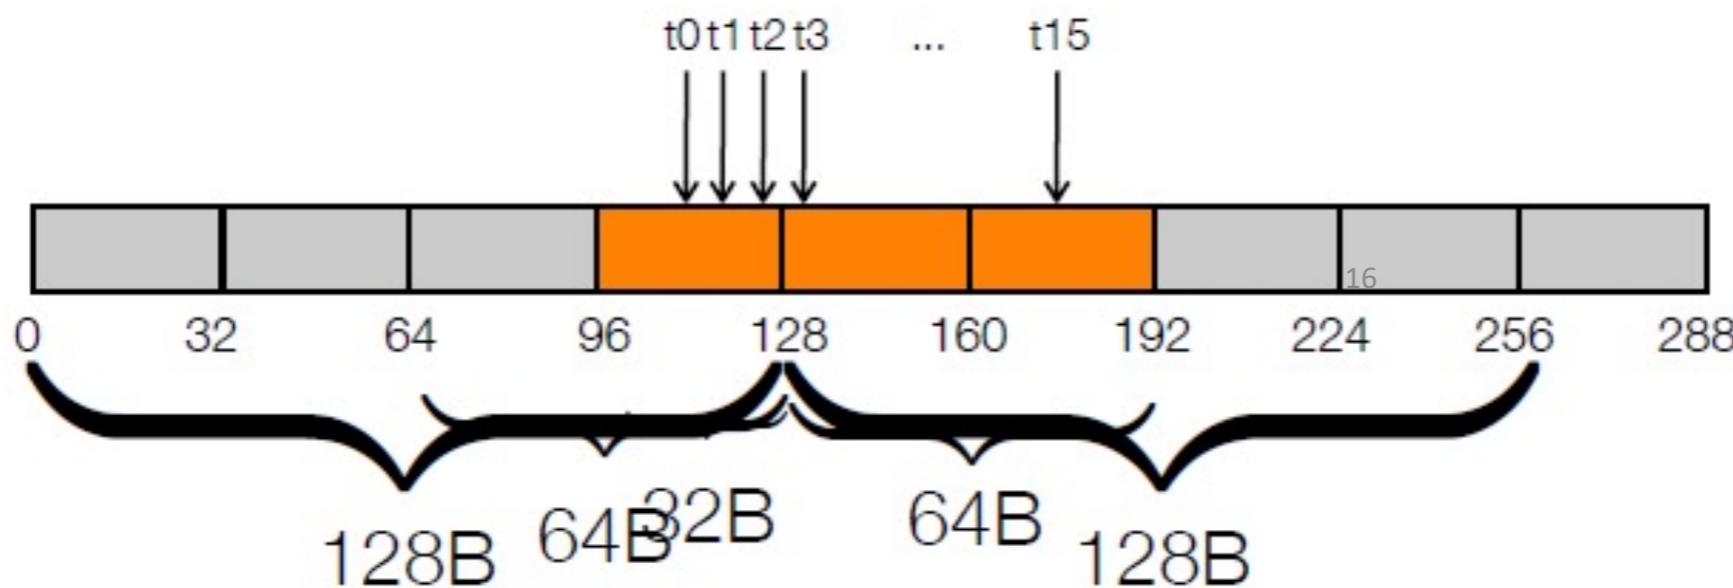
- ▶ Warp wont finish until the last thread finishes
 - Warp will be dragged
- ▶ Possible solution:
 - Try flatten peaks by making threads work in multiple data

Memory coalescing

- ▶ Global memory is accessed in chunks of aligned 32, 64 or 128 bytes
- ▶ Following protocol is used to issue a memory transaction of a half warp (valid for 1.X)
 - Find memory segment that contains address requested by the lowest numbered active thread
 - Find other active threads whose requested address lies in same segment, and reduce transaction size if possible
 - Do transaction. Mark serviced threads as inactive.¹⁵
 - Repeat until all threads are serviced
- ▶ Worse case: fetch 32 bytes, use only 4 bytes: 7/8 wasted bandwidth!

Memory coalescing

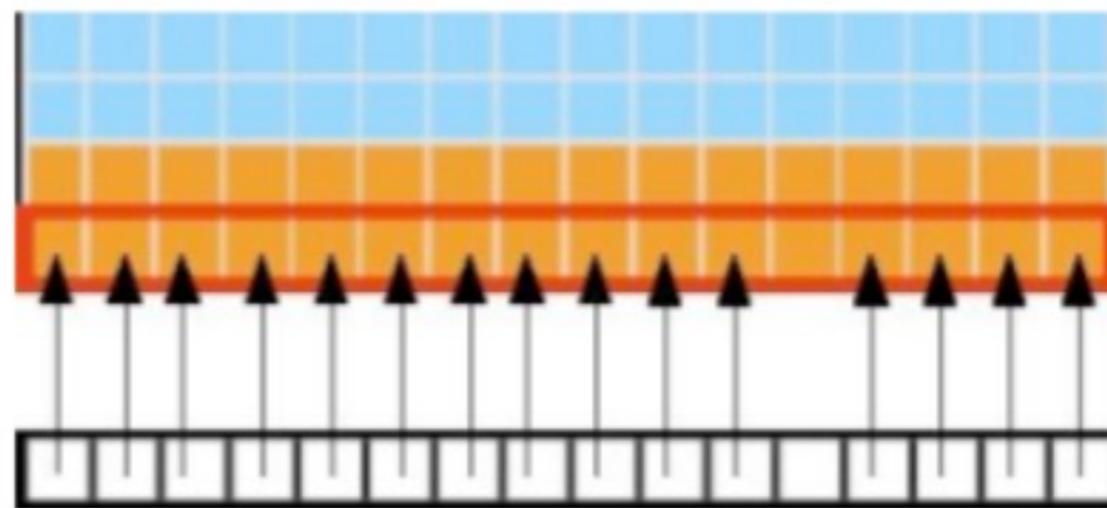
- ▶ Access pattern visualization
 - Thread 0 is lowest active, accesses address 116
 - ▶ Belongs to 128-byte segment 0-127



David Tarjan - NVIDIA

Memory coalescing

- ▶ Simple access pattern

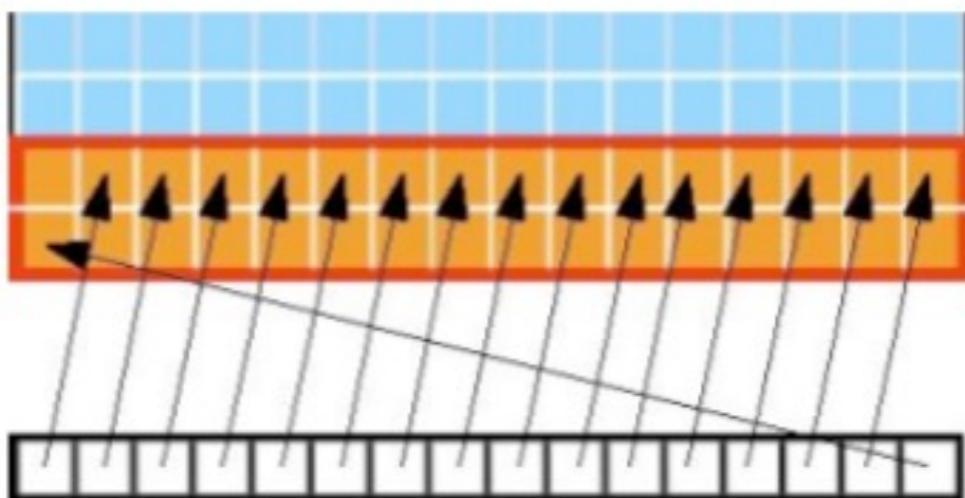


One 64 byte transaction¹⁷

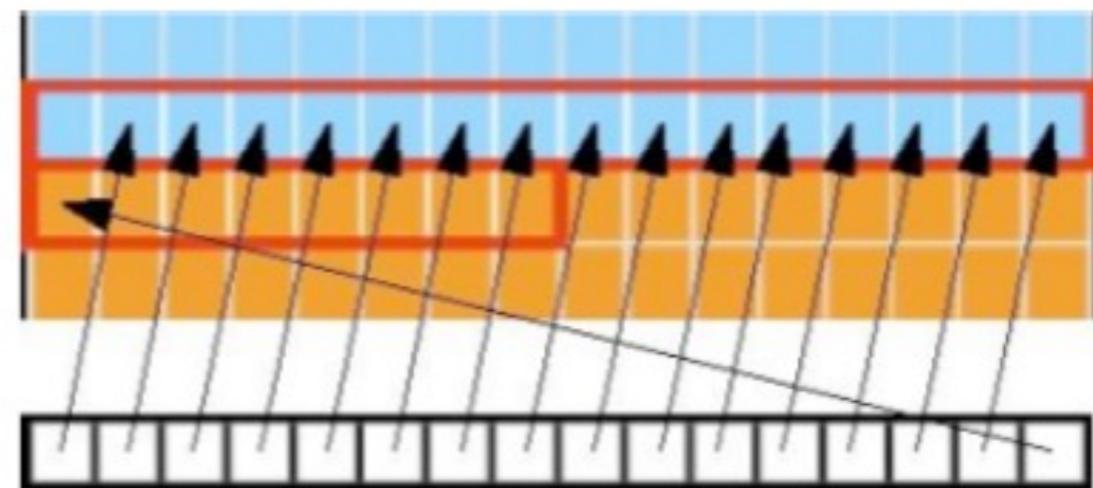
Will be looking at compute capability 1.X examples

Memory coalescing

- ▶ Sequential but misaligned



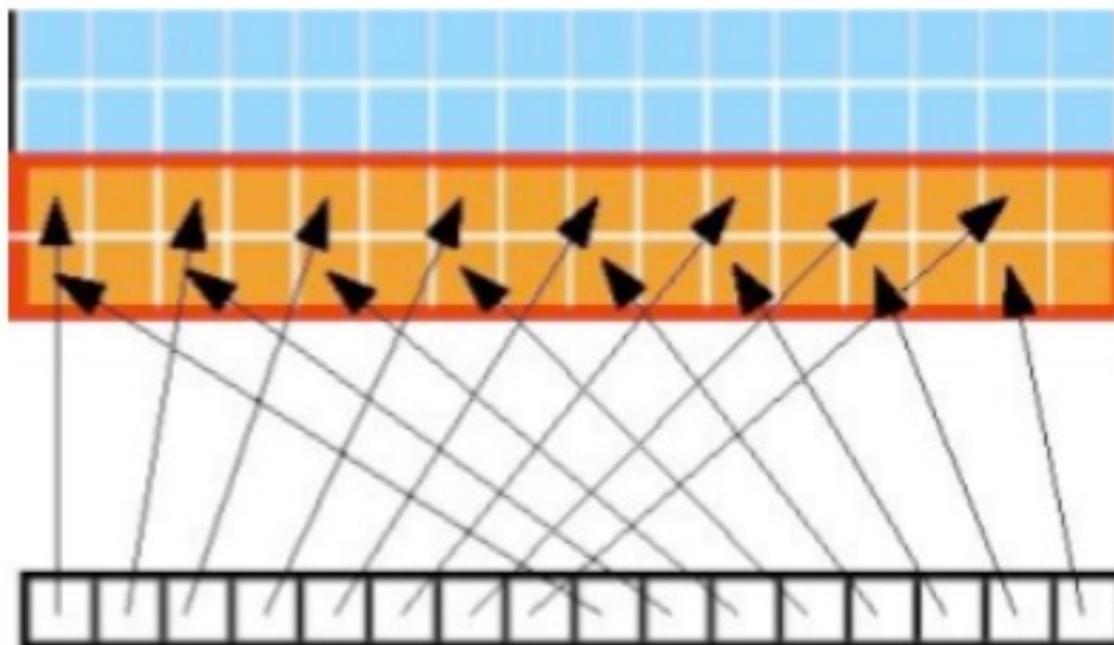
One 128 byte transaction



One 64 byte transaction
and one 32 byte transaction

Memory coalescing

- ▶ Strided access

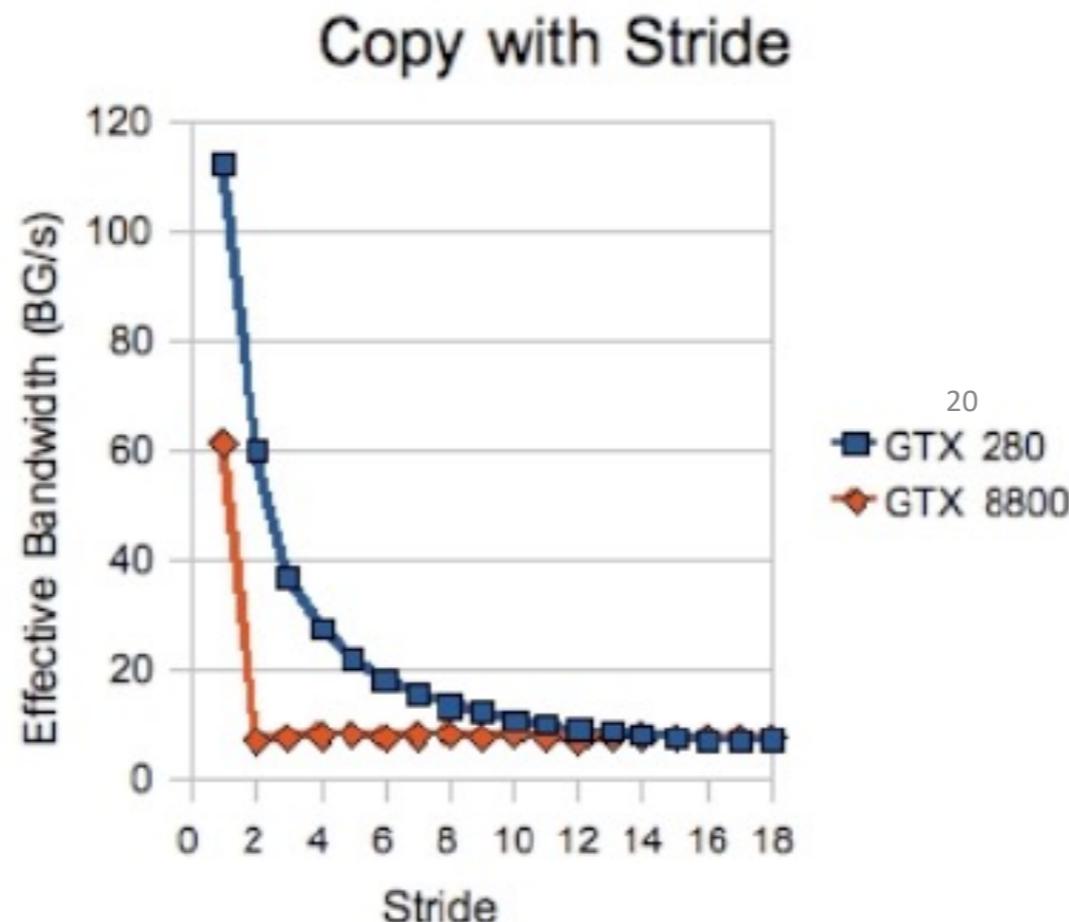


One 128 byte transaction, but half of
bandwidth is wasted

Memory coalescing

- ▶ Example: Copy with stride

```
__global__ void strideCopy(float *odata, float *idata, int stride)
{
    int xid = (blockIdx.x*blockDim.x+threadIdx.x)*stride;
    odata[xid] = idata[xid];
}
```



Memory coalescing

- ▶ 2.X architecture
 - Global memory is cached
 - ▶ Cached in both L1 and L2: 128 byte transaction
 - ▶ Cached only in L2: 32 byte transaction
 - Whole warps instead of half warps

Memory coalescing - SoA or AoS?

- ▶ Array of structures

```
struct record
{
    int key;
    int value;
    int flag;
};

record *d_record;
cudaMalloc((void**) &d_records, ...);
```

22

Memory coalescing - SoA or AoS?

- ▶ Structure of array

```
struct SoA
{
    int *key;
    int *value;
    int *flag;
};

SoA *d_AoA_data;
cudaMalloc((void**) &d_SoA_data.keys, ...);
cudaMalloc((void**) &d_SoA_data.value, ...);
cudaMalloc((void**) &d_SoA_data.flag,23 ...);
```

Memory coalescing - SoA or AoS?

- ▶ cudaMalloc guarantees aligned memory, then accessing the SoA will be much more efficient

```
__global__ void bar (record *AoS_data, SoA SoA_data)
{
    int i = blockDim.x*blockIdx.x + threadIdx.x;
    // AoS wastes bandwidth
    int key = AoS_data[i].key;
    // SoA efficient use of bandwidth
    int key_better = SoA_data.keys[i];
}
```

24

Parallel Reduction

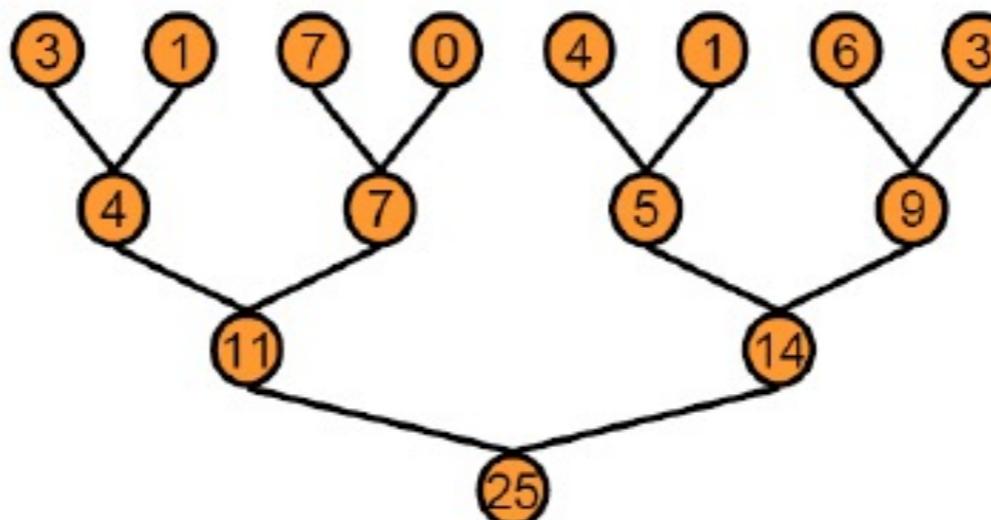
- Common and important data parallel primitive
 - Easy to implement in CUDA
 - Harder to get it right
- Serves as a great optimization example
 - Step by step through 7 different versions
 - Demonstrates several important optimization strategies

$$(3) + (1) + (7) + (0) + (4) + (1) + (6) + (3) = ?$$

25

Parallel Reduction

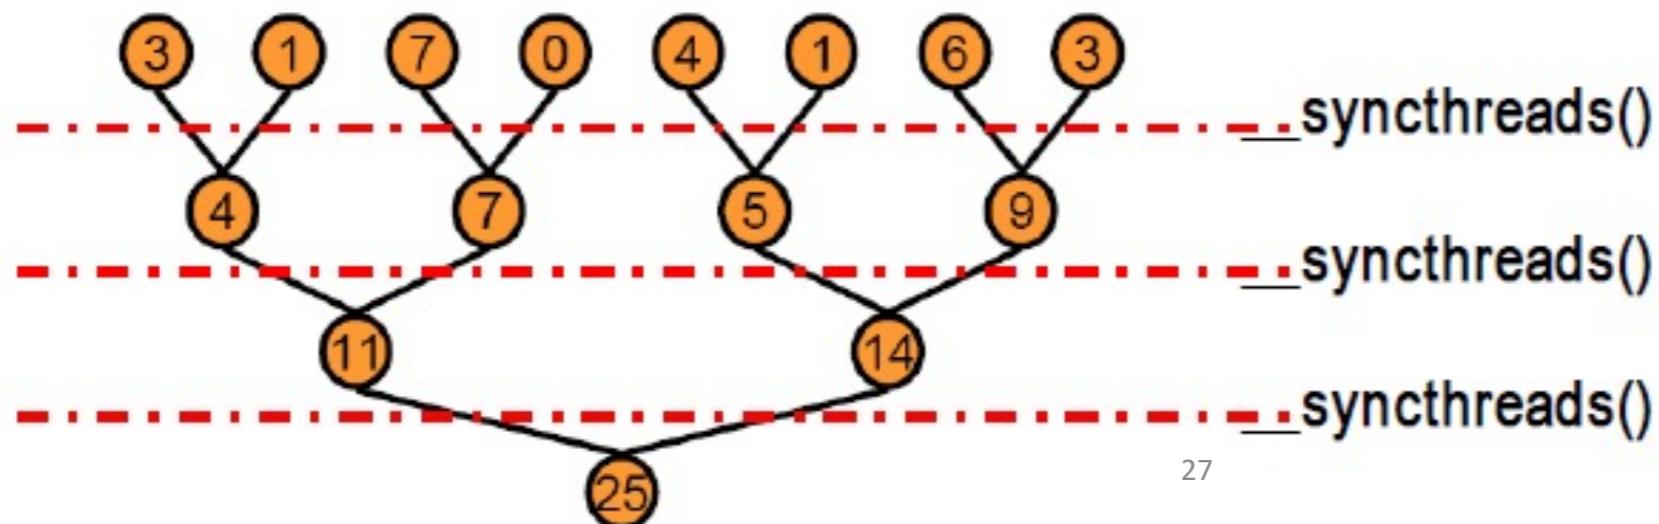
- Tree-based approach used within each thread block



- Multiple thread blocks
 - Very large array
 - Each thread block processes a portion of the array
- How to manage communication between thread blocks?

Parallel Reduction: Global Synchronization

- Thread synchronization after pair-wise reduction



27

Parallel Reduction: Global Synchronization

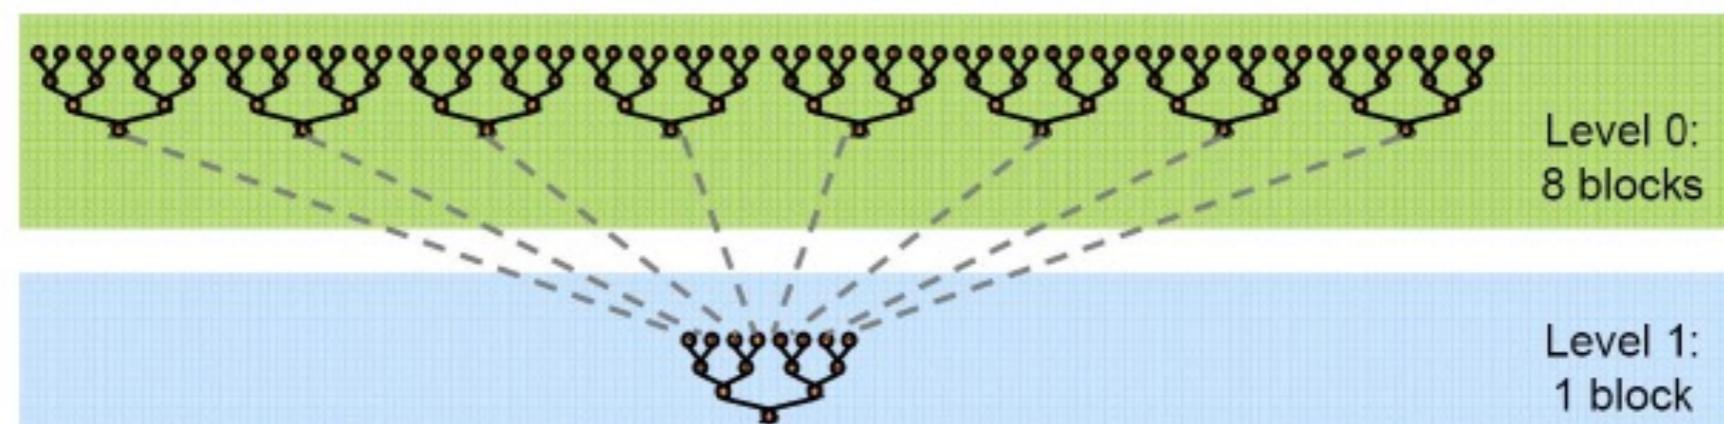
- Problem of CUDA — not support global synchronization
 - Too many kernels, hardware synchronization is costly
 - Dead-lock

- Solution
 - Thread synchronization within thread block

28

Parallel Reduction: Decomposition

- Partition the data into blocks/kernels



- Code in each thread is identical
- Iterative

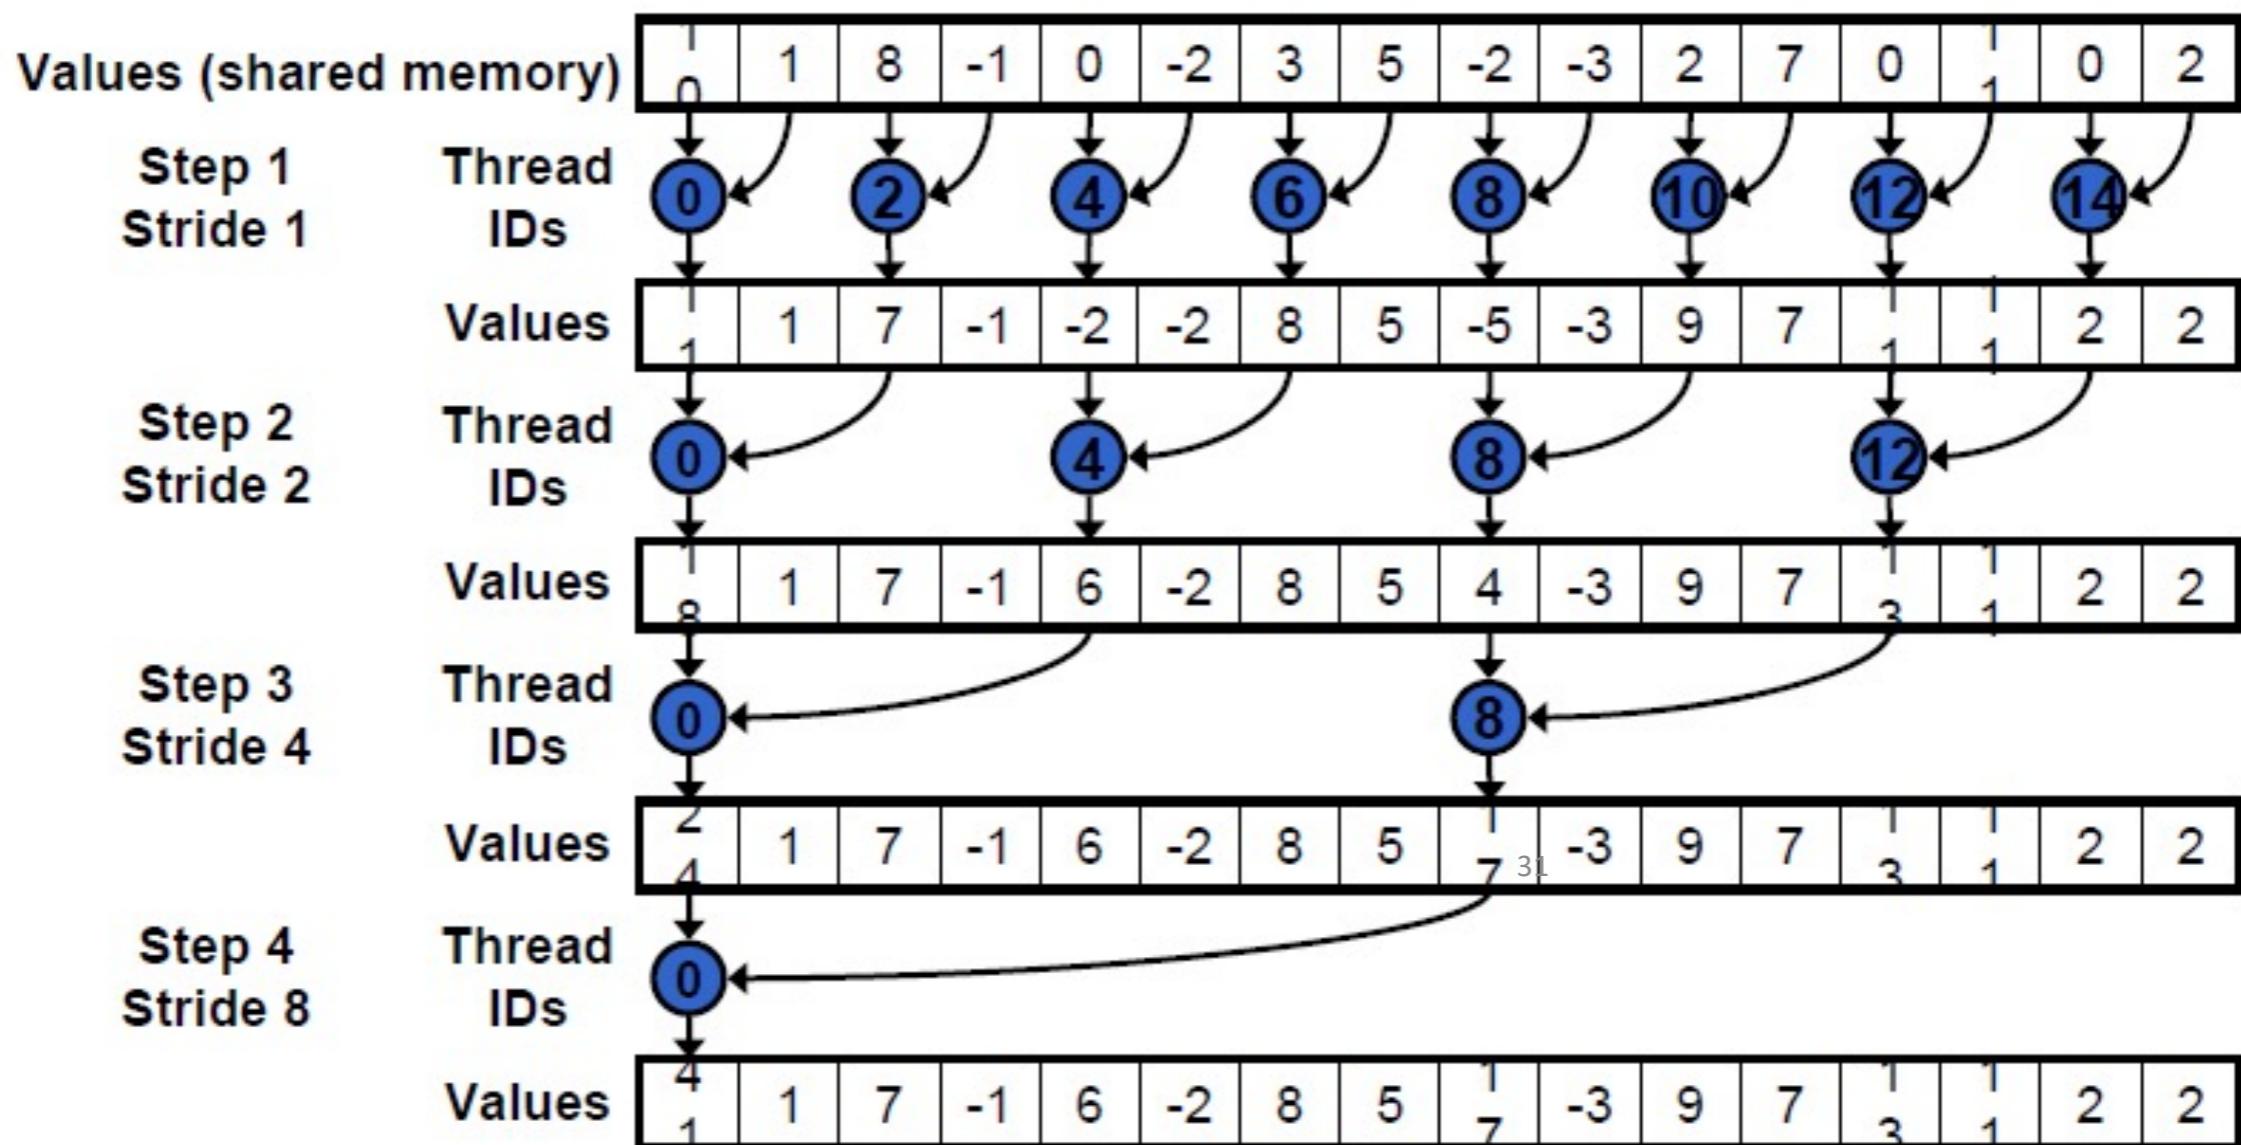
29

Optimization Goal

- Philosophy
 - GFLOP/s: for compute-bound kernels
 - QR factorization, convolution, FIR filter, e.g.
 - Bandwidth: for memory-bound kernels
 - Database, video playback, ..., e.g.
 - Both
 - Pattern matching, singular value decomposition, ..., e.g.
- Reduction has very low computation intensity
 - 1 floating point operation / 2 elements
 - **Maximize bandwidth!**
- G80
 - 384-bit memory interface, 900 MHz DDR
 - $384 * 900/4 = 86.4 \text{ GB/s}$

30

Reduction #1: Interleaved Addressing



Reduction #1: Interleaved Addressing

```
__global__ void reduce0(int *g_idata, int *g_odata) {
    extern __shared__ int sdata[];
    // each thread loads one element from global to shared mem
    unsigned int tid = threadIdx.x;
    unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;
    sdata[tid] = g_idata[i];
    __syncthreads();

    // do reduction in shared mem
    for(unsigned int s=1; s < blockDim.x; s *= 2) {
        if (tid % (2*s) == 0)
            sdata[tid] += sdata[tid + s];
        __syncthreads();
    }
    // write result for this block to global mem
    if (tid == 0) g_odata[blockIdx.x] = sdata[0];
}
```

32

Reduction #1: Interleaved Addressing

```
__global__ void reduce0(int *g_idata, int *g_odata) {
    extern __shared__ int sdata[];
    // each thread loads one element from global to shared mem
    unsigned int tid = threadIdx.x;
    unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;
    sdata[tid] = g_idata[i];
    __syncthreads();
    // do reduction in shared mem
    for(unsigned int s=1; s < blockDim.x; s *= 2) {
        if (tid % (2*s) == 0) { ←
            sdata[tid] += sdata[tid + s];
        }
        __syncthreads();
    }
    // write result for this block to global mem
    if (tid == 0) g_odata[blockIdx.x] = sdata[0];
}
```

Problem: highly divergent
branching results in very poor
performance!
23
% operator is very slow

Performance for 4M Element Reduction

	Time (2^{22} ints)	Bandwidth
Kernel 1: interleaved addressing with divergent branching	8.054 ms	2.083 GB/s

Note: Block Size = 128 threads for all tests

Reduction #2: Interleaved Addressing

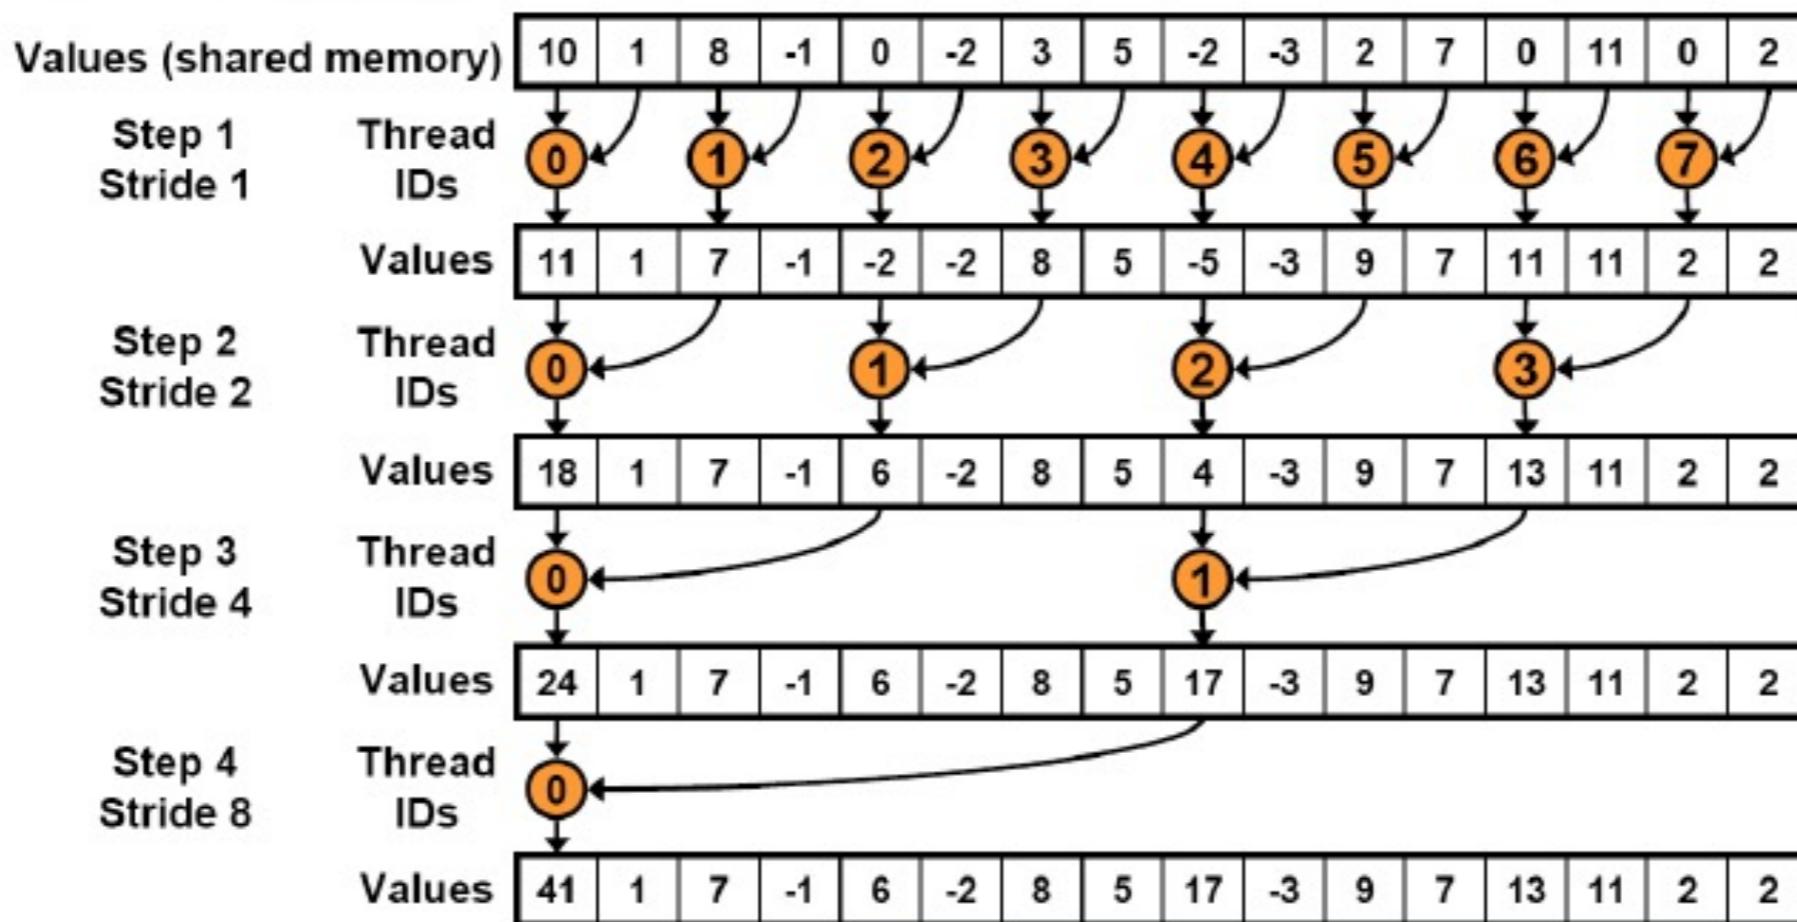
- Divergence

```
for (unsigned int s=1; s < blockDim.x; s *= 2) {  
    if (tid % (2*s) == 0) {  
        sdata[tid] += sdata[tid + s];  
    }  
    __syncthreads();  
}
```

- Avoid divergence

```
for (unsigned int s=1; s < blockDim.x; s *= 2) {  
    int index = 2 * s * tid; 35  
    if (index < blockDim.x) {  
        sdata[index] += sdata[index + s];  
    }  
    __syncthreads();  
}
```

Reduction #2: Interleaved Addressing



```
for (unsigned int s=1; s < blockDim.x; s *= 2) {  
    int index = 2 * s * tid;  
    if (index < blockDim.x) {  
        sdata[index] += sdata[index + s];  
    }  
    __syncthreads();  
}
```

Performance for 4M Element Reduction

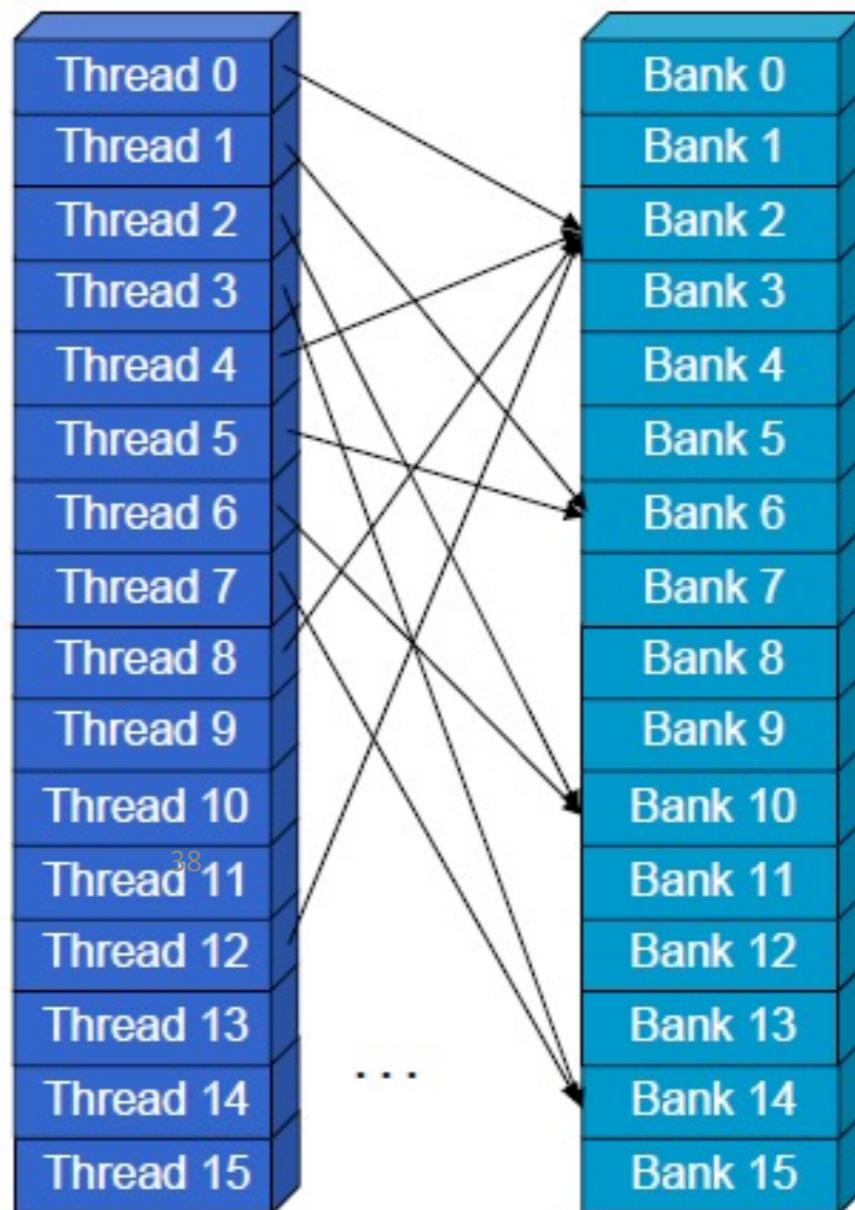
	Time (2^{22} ints)	Bandwidth	Step Speedup	Cumulative Speedup
Kernel 1: interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
Kernel 2: interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x

Reduction #2: Interleaved Addressing

■ Bank conflict!

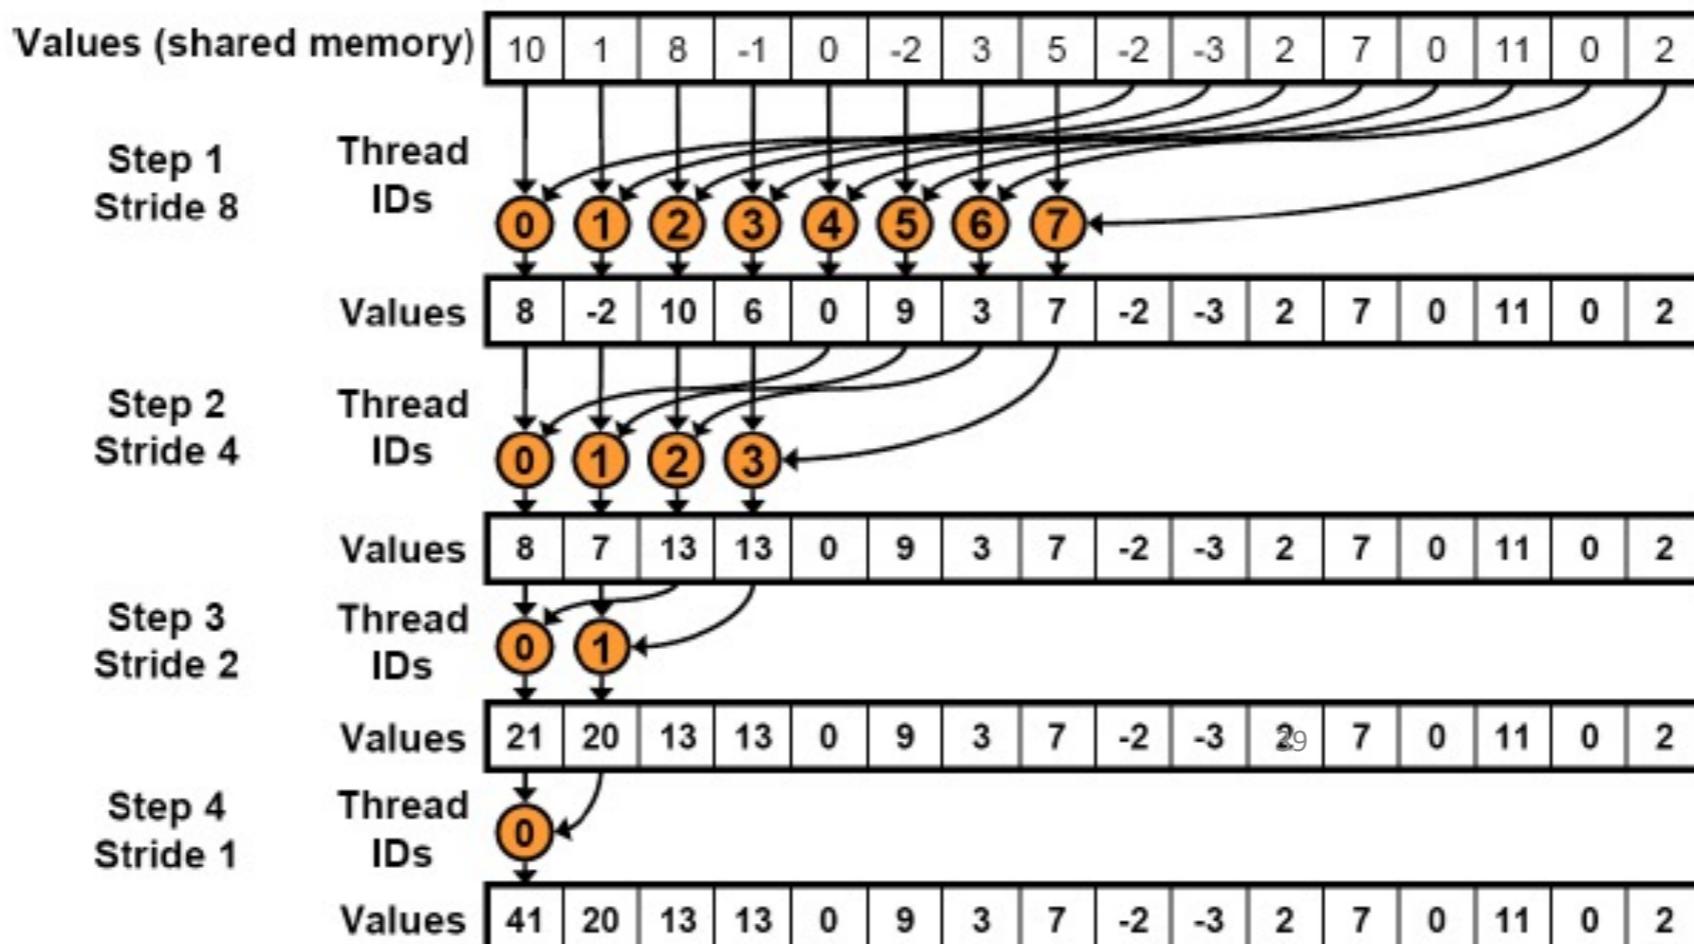
```
for (unsigned int s=1; s < blockDim.x; s *= 2) {  
    int index = 2 * s * tid;  
    if (index < blockDim.x) {  
        sdata[index] += sdata[index + s];  
    }  
    __syncthreads();  
}
```

s=2



Parallel Reduction: Sequential Addressing

■ Bank conflict free



Reduction #3: Sequential Addressing

- Just replace strided indexing in inner loop:

```
for (unsigned int s=1; s < blockDim.x; s *= 2) {  
    int index = 2 * s * tid;  
    if (index < blockDim.x) {  
        sdata[index] += sdata[index + s];  
    }  
    __syncthreads();  
}
```

- With reversed loop and threadID-based indexing:

```
for (unsigned int s=blockDim.x/2; s>0; s>>=1) {  
    if (tid < s) {  
        sdata[tid] += sdata[tid + s];  
    }  
    __syncthreads();  
}
```

Performance for 4M Element Reduction

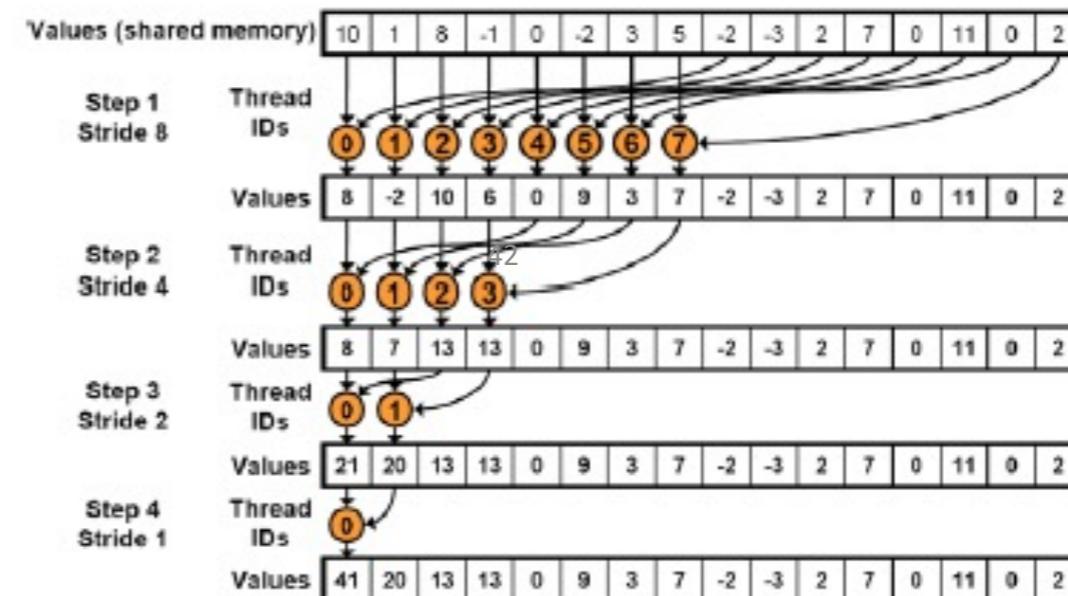
	Time (2 ²² ints)	Bandwidth	Step Speedup	Cumulative Speedup
Kernel 1: interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
Kernel 2: interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x
Kernel 3: sequential addressing	1.722 ms	9.741 GB/s	2.01x	4.68x

41

Observation: Idle Threads

```
for (unsigned int s=blockDim.x/2; s>0; s>>=1) {  
    if (tid < s) {  
        sdata[tid] += sdata[tid + s];  
    }  
    __syncthreads();  
}
```

- In the first iteration, half threads are idle!
 - Waste half resources ...



Reduction #4: First Add During Load

```
// each thread loads one element from global to shared mem
unsigned int tid = threadIdx.x;
unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;
sdata[tid] = g_idata[i];
__syncthreads();
```

- Reduce # blocks by half
- 2 global memory loads in the first add

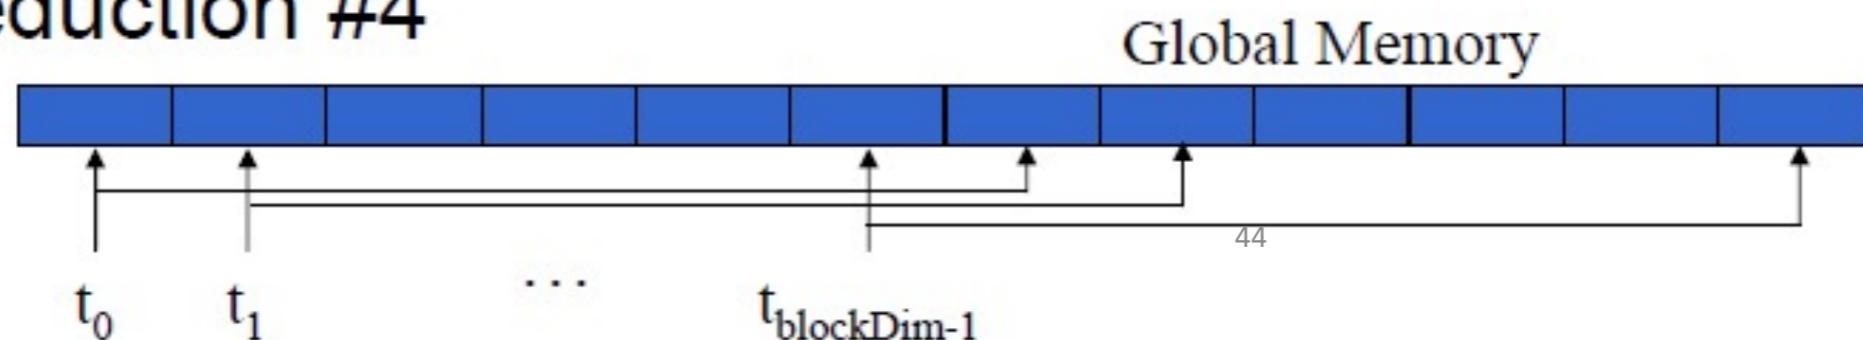
```
// perform first level of reduction
// reading from global memory, writing to shared memory
unsigned int tid = threadIdx.x;
unsigned int i = blockIdx.x*(blockDim.x*2) + threadIdx.x; 43
sdata[tid] = g_idata[i] + g_idata[i+blockDim.x];
__syncthreads();
```

Reduction #4: First Add During Load

- Reduction #3



- Reduction #4



Performance for 4M Element Reduction

	Time (2 ²² ints)	Bandwidth	Step Speedup	Cumulative Speedup
Kernel 1: interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
Kernel 2: interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x
Kernel 3: sequential addressing	1.722 ms	9.741 GB/s	2.01x	4.68x
Kernel 4: first add during global load	0.965 ms	17.377 GB/s	1.78x	8.34x

45

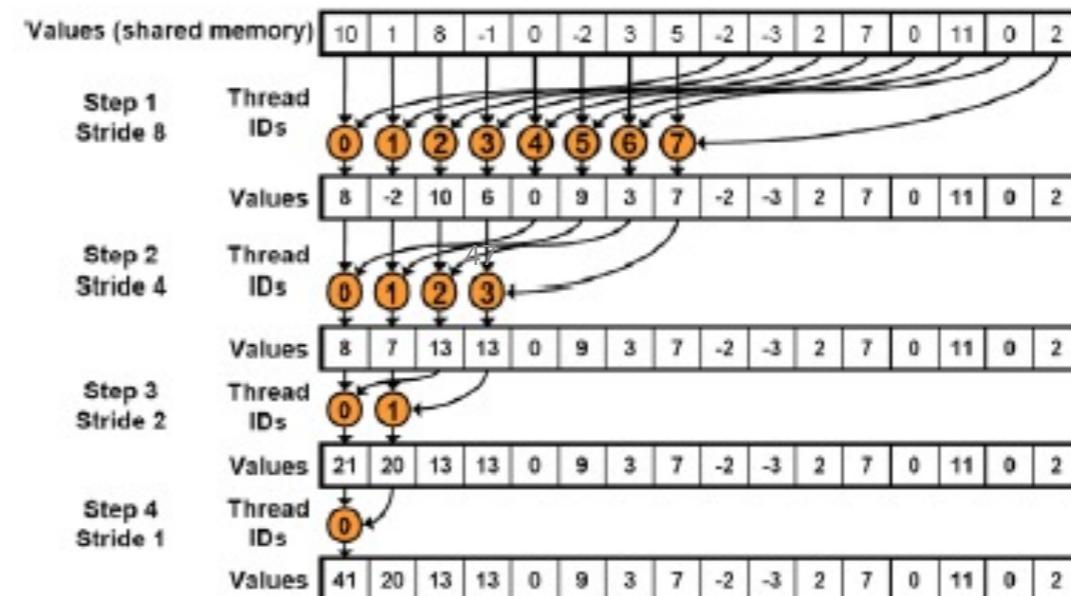
Performance Bottleneck

- 17 GB/s << 86.4 GB/s
 - Algorithm has been optimized
 - Other operations may incur the cost
 - loads, stores
 - loops
- Strategy
 - Code optimization

46

Loop Unrolling

- # threads drops
 - When # threads ≤ 32 , only one warp is working
 - Other warps are idle
 - We don't need "if (tid < s)" because it doesn't save any work
- Instructions are SIMD synchronous within a warp
- No synchronization is required in a warp
 - Scoreboarding automatically maintain synchronization
- Unroll the last 6 loops



Reduction #5: Unroll the Last Warp

```
for (unsigned int s=blockDim.x/2; s>0; s>>=1) {  
    if (tid < s) {  
        sdata[tid] += sdata[tid + s];  
    }  
    __syncthreads();  
}
```



```
for (unsigned int s=blockDim.x/2; s>32; s>>=1){  
    if (tid < s)  
        sdata[tid] += sdata[tid + s];  
    __syncthreads();  
}  
if (tid < 32){  
    sdata[tid] += sdata[tid + 32]; sdata[tid] += sdata[tid + 16];  
    sdata[tid] += sdata[tid + 8]; sdata[tid] += sdata[tid + 4];  
    sdata[tid] += sdata[tid + 2]; sdata[tid] += sdata[tid + 1];  
}
```

48

Performance for 4M Element Reduction

	Time (2^{22} ints)	Bandwidth	Step Speedup	Cumulative Speedup
Kernel 1: interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
Kernel 2: interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x
Kernel 3: sequential addressing	1.722 ms	9.741 GB/s	2.01x	4.68x
Kernel 4: first add during global load	0.965 ms	17.377 GB/s	1.78x	8.34x
Kernel 5: unroll last warp	0.536 ms	31.289 GB/s	1.8x	15.01x

49

Completely Unrolled

- Loops may be completely unrolled
 - Assumption: the number of iterations is known!
 - Max # block threads = 512
 - 2^m threads in thread block
 - Unroll loops for 2^m threads, where $m = 0, 1, 2, \dots$
- So we can easily unroll for a fixed block size
 - How can we unroll for block sizes that we don't know at compile time?
- How to make the code generic?
 - How to know # threads in a block when compiling?
- **Template!**
 - CUDA supports C++ template parameter

Unrolling with Templates

- Specify blocksize as a function template parameter:

```
template <unsigned int blockSize>
__global__ void reduce5(int *g_idata, int *g_odata)
```

Reduction #6: Completely Unrolled

```
if (blockSize >= 512) {
    if (tid < 256) { sdata[tid] += sdata[tid + 256]; } __syncthreads();
}
if (blockSize >= 256) {
    if (tid < 128) { sdata[tid] += sdata[tid + 128]; } __syncthreads();
}
if (blockSize >= 128) {
    if (tid < 64) { sdata[tid] += sdata[tid + 64]; } __syncthreads();
}
if (tid < 32) {
    if (blockSize >= 64) sdata[tid] += sdata[tid + 32];
    if (blockSize >= 32) sdata[tid] += sdata[tid + 16];
    if (blockSize >= 16) sdata[tid] += sdata[tid + 8];
    if (blockSize >= 8) sdata[tid] += sdata[tid + 4];
    if (blockSize >= 4) sdata[tid] += sdata[tid + 2];
    if (blockSize >= 2) sdata[tid] += sdata[tid + 1];
}
```

52

All code in RED will be evaluated at compile time

Invoking Template Kernels

- Don't we still need block size at compile time?
 - No! just a switch statement for 10 possible block sizes:

```
switch (threads){  
    case 512: reduce5<512><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;  
    case 256: reduce5<256><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;  
    case 128: reduce5<128><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;  
    case 64: reduce5< 64><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;  
    case 32: reduce5< 32><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;  
    case 16: reduce5< 16><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;  
    case 8: reduce5< 8><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;  
    case 4: reduce5< 4><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;  
    case 2: reduce5< 2><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;  
    case 1: reduce5< 1><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;  
}
```

Performance for 4M Element Reduction

	Time (2 ²² ints)	Bandwidth	Step Speedup	Cumulative Speedup
Kernel 1: interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
Kernel 2: interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x
Kernel 3: sequential addressing	1.722 ms	9.741 GB/s	2.01x	4.68x
Kernel 4: first add during global load	0.965 ms	17.377 GB/s	1.78x	8.34x
Kernel 5: unroll last warp	0.536 ms	31.289 GB/s	1.8x	15.01x
Kernel 6: completely unrolled	0.381 ms	43.996 GB/s	1.41x	21.16x

54

Reduction #7: Multiple Adds / Thread

- Replace adding 2 global memory loads in the first add

```
unsigned int tid = threadIdx.x;
unsigned int i = blockIdx.x*(blockDim.x*2) + threadIdx.x;
sdata[tid] = g_idata[i] + g_idata[i+blockDim.x];
__syncthreads();
```

- By adding multiple global memory loads

```
unsigned int tid = threadIdx.x;
unsigned int i = blockIdx.x*(blockSize*2) + threadIdx.x;
unsigned int gridSize = blockSize*2*gridDim.x;
sdata[tid] = 0;                                55
while (i < n) {
    sdata[tid] += g_idata[i] + g_idata[i+blockSize];
    i += gridSize;
}
__syncthreads();
```

Reduction #7: Multiple Adds / Thread

- gridSize, incremental global memory address
 - gridSize is multiples of 16
 - Aligned with next global memory load
 - Memory coalescing!**

```
unsigned int tid = threadIdx.x;
unsigned int i = blockIdx.x*(blockSize*2) + threadIdx.x;
unsigned int gridSize = blockSize*2*gridDim.x;
sdata[tid] = 0;
while (i < n) {
    sdata[tid] += g_idata[i] + g_idata[i+blockSize];
    i += gridSize;
}
__syncthreads();
```

Performance for 4M Element Reduction

	Time (2^{22} ints)	Bandwidth	Step Speedup	Cumulative Speedup
Kernel 1: interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
Kernel 2: interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x
Kernel 3: sequential addressing	1.722 ms	9.741 GB/s	2.01x	4.68x
Kernel 4: first add during global load	0.965 ms	17.377 GB/s	1.78x	8.34x
Kernel 5: unroll last warp	0.536 ms	31.289 GB/s	1.8x	15.01x
Kernel 6: completely unrolled	0.381 ms	43.996 GB/s	⁵⁷ 1.41x	21.16x
Kernel 7: multiple elements per thread	0.268 ms	62.671 GB/s	1.42x	30.04x

Kernel 7 on 16M elements: 72 GB/s!

```

template <unsigned int blockSize>
__global__ void reduce6(int *g_idata, int *g_odata, unsigned int n)
{
    extern __shared__ int sdata[];
    unsigned int tid = threadIdx.x;
    unsigned int i = blockIdx.x*(blockSize*2) + tid;
    unsigned int gridSize = blockSize*2*gridDim.x;
    sdata[tid] = 0;
    do { sdata[tid] += g_idata[i] + g_idata[i+blockSize]; i += gridSize; }
        while (i < n);
    __syncthreads();
    if (blockSize >= 512) { if (tid < 256) { sdata[tid] += sdata[tid + 256]; }
        __syncthreads(); }
    if (blockSize >= 256) { if (tid < 128) { sdata[tid] += sdata[tid + 128]; }
        __syncthreads(); }
    if (blockSize >= 128) { if (tid < 64) { sdata[tid] += sdata[tid + 64]; }
        __syncthreads(); }
    if (tid < 32) {
        if (blockSize >= 64) sdata[tid] += sdata[tid + 32];
        if (blockSize >= 32) sdata[tid] += sdata[tid + 16];
        if (blockSize >= 16) sdata[tid] += sdata[tid + 8];
        if (blockSize >= 8) sdata[tid] += sdata[tid + 4];
        if (blockSize >= 4) sdata[tid] += sdata[tid + 2];
        if (blockSize >= 2) sdata[tid] += sdata[tid + 1];
    }
    if (tid == 0) g_odata[blockIdx.x] = sdata[0];
}

```

58

final version

Parallel Reduction Summary

- Algorithm optimization
 - Changes to addressing, algorithm cascading
 - 8.34x speedup, combined!
- Code optimization
 - Loops Unrolling
 - 3.6x speedup, combined

59

GPU Computing with CUDA

Lecture 6 - CUDA Libraries - Thrust

*Christopher Cooper
Boston University*

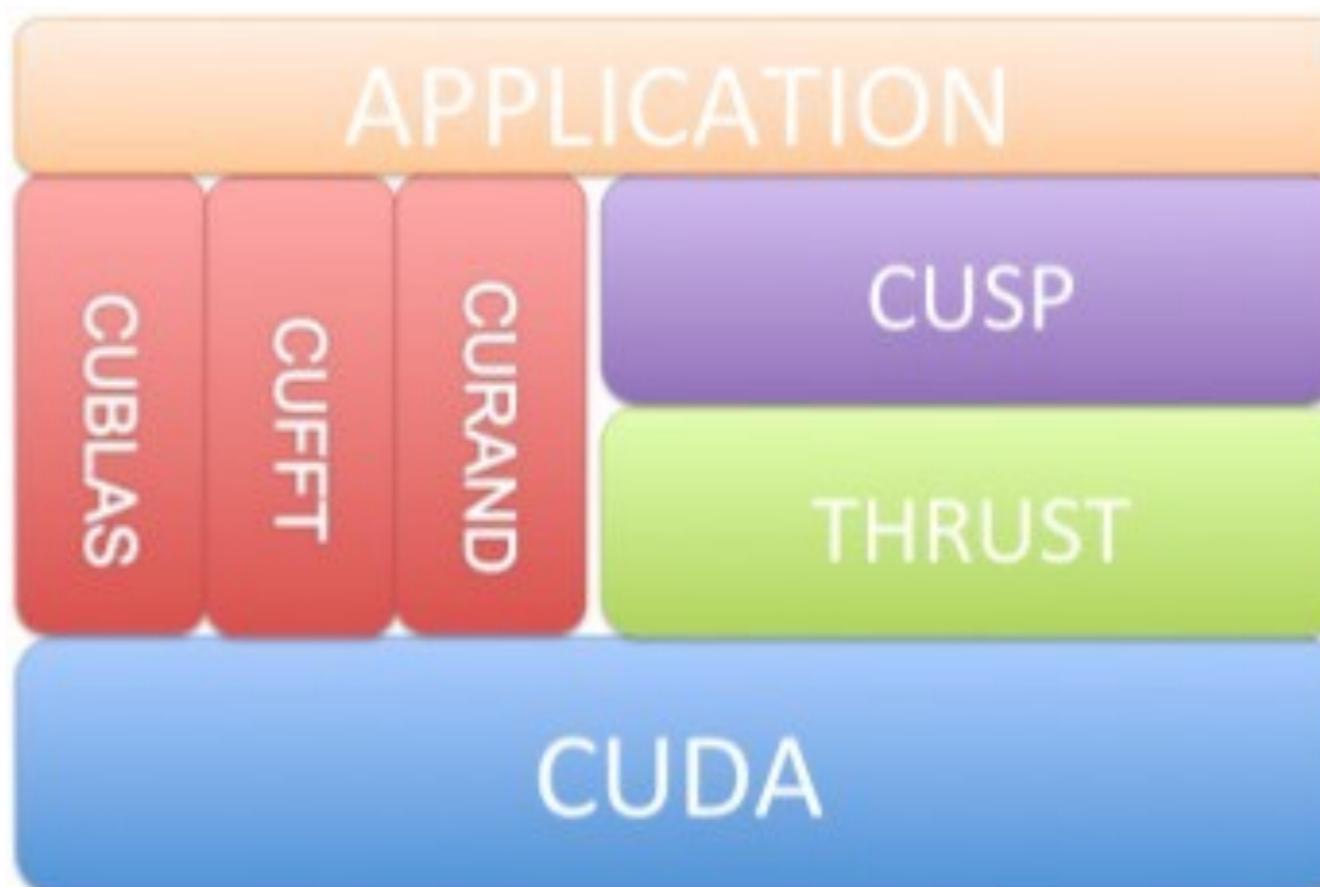
*August, 2011
UTFSM, Valparaíso, Chile*

Outline of lecture

- ▶ CUDA Libraries
- ▶ What is Thrust?
- ▶ Features of thrust
- ▶ Best practices in thrust

CUDA Libraries

- NVIDIA has developed several libraries to abstract the user from CUDA



CUDA Libraries



**DEVELOPER
ZONE**

DEVELOPER CENTERS TECHNOLOGIES TOOLS RESOURCES COMMUNITY

► CUDPP

Libraries

CUBLAS, CUSP, CUFFT, Thrust, and many other CUDA based libraries can be found here.

► MAGMA

► IMSL

► VSIPL

► NVML

► CUPTI

►

GPU AI - PATH FINDING

Technology preview that includes libraries and samples applicaitons CUDA-accelerated path finding.



NVIDIA PERFORMANCE PRIMITIVES

NVIDIA NPP is a library of functions for performing CUDA accelerated processing. The initial set of functionality in the library focuses on imaging and video processing and is widely applicable for...



THRUST

Standard Template Library for CUDA, featuring many highly optimized implementations



CUBLAS

CUDA Basic Linear Algebra Library



CUFFT

CUDA Fast Fourier Transform Library



CUSP

Cusp is a library for sparse linear algebra and graph computations on CUDA. Cusp provides a flexible, high-level interface for manipulating sparse matrices and solving sparse linear systems.



<http://developer.nvidia.com/technologies/libraries>

CUDA Libraries

 **DEVELOPER
ZONE**

DEVELOPER CENTERS TECHNOLOGIES TOOLS RESOURCES COMMUNITY

▶ CUDPP

Libraries

CUBLAS, CUSP, CUFFT, Thrust, and many other CUDA based libraries can be found here.

▶ MAGMA

▶ IMSL

▶ VSIPL

▶ NVML

▶ CUPTI

▶

GPU AI - PATH FINDING

Technology preview that includes libraries and samples applicaitons CUDA-accelerated path finding.



NVIDIA PERFORMANCE PRIMITIVES

NVIDIA NPP is a library of functions for performing CUDA accelerated processing. The initial set of functionality in the library focuses on imaging and video processing and is widely applicable for...



THRUST

Standard Template Library for CUDA, featuring many highly optimized implementations



CUBLAS

CUDA Basic Linear Algebra Library



CUFFT

CUDA Fast Fourier Transform Library



CUSP

Cusp is a library for sparse linear algebra and graph computations on CUDA. Cusp provides a flexible, high-level interface for manipulating sparse matrices and solving sparse linear systems.



<http://developer.nvidia.com/technologies/libraries>

Thrust - Introduction

- ▶ Template library for CUDA
 - Resembles C++ Standard Template Library (STL)
 - Collection of data parallel primitives
- ▶ Objectives
 - Programmer productivity
 - Encourage generic programming
 - High performance
 - Interoperability
- ▶ Comes with CUDA 4.0



Thrust - Introduction

- ▶ Containers
 - `thrust::host_vector<T>`
 - `thrust::device_vector<T>`
- ▶ Algorithms
 - `thrust::sort()`
 - `thrust::reduce()`
 - `thrust::inclusive_scan()`
- ▶ <http://code.google.com/p/thrust/>
- ▶ Slides from Nathan Bell and Jared Hoberock - NVIDIA

Thrust - Containers

- ▶ Thrust provides two vector containers
 - `host_vector`: resides on CPU
 - `device_vector`: resides on GPU
- ▶ Hides `cudaMalloc` and `cudaMemcpy`

```
// allocate host vector with two elements
thrust::host_vector<int> h_vec(2);

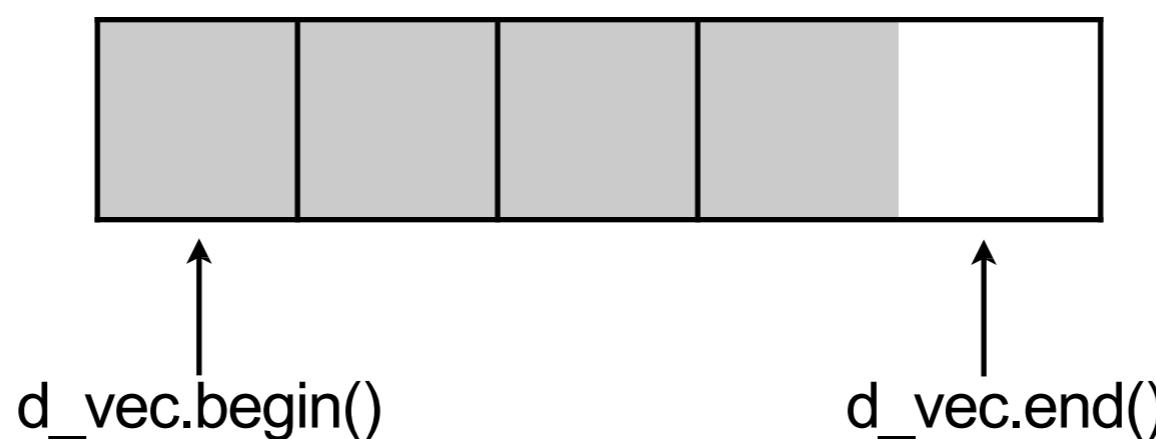
// copy host vector to device
thrust::device_vector<int> d_vec = h_vec;
// manipulate device values from the host
d_vec[0] = 13;
d_vec[1] = 27;                                + d_vec[1] << std::endl;

std::cout << "sum: " << d_vec[0]
// vector memory automatically released w/ free() or cudaFree()
```

Thrust - Iterators

- › Iterators can be thought as pointers to array elements
 - They carry other information

```
// allocate device vector
thrust::device_vector<int> d_vec(4);
d_vec.begin(); // returns iterator at first element of
d_vec
d_vec.end()// returns iterator one past the last
element of d_vec
// [begin, end] pair defines a sequence of 4 elements
```



Thrust - Iterators

- Use iterators like pointers

```
// allocate device vector
thrust::device_vector<int> d_vec(4);
thrust::device_vector<int>::iterator begin = d_vec.begin();
*begin = 13;           // same as d_vec[0] = 13;
int temp = *begin;    // same as temp = d_vec[0];
begin++;              // advance iterator one position
*begin = 25;           // same as d_vec[1] = 25;
```

Thrust - Iterators

- Important: keep track of your memory space

```
// initialize random values on host
thrust::host_vector<int> h_vec(1000);
thrust::generate(h_vec.begin(), h_vec.end(), rand);

// copy values to device
thrust::device_vector<int> d_vec = h_vec;

// compute sum on host
int h_sum = thrust::reduce(h_vec.begin(), h_vec.end());

// compute sum on device
int d_sum = thrust::reduce(d_vec.begin(), d_vec.end());
```

Thrust - Iterators

- Convertible to raw pointers

```
// allocate device vector
thrust::device_vector<int> d_vec(4);

// obtain raw pointer to device vector's memory
int * ptr = thrust::raw_pointer_cast(&d_vec[0]);

// use ptr in a CUDA C kernel
my_kernel<<<N/256, 256>>>(N, ptr);

// Note: ptr cannot be dereferenced on the host!
```

Thrust - Iterators

- Wrap raw pointers to use in thrust

```
int N = 10;

// raw pointer to device memory
int * raw_ptr;
cudaMalloc((void **) &raw_ptr, N * sizeof(int));

// wrap raw pointer with a device_ptr
thrust::device_ptr<int> dev_ptr(raw_ptr);

// use device_ptr in thrust algorithms
thrust::fill(dev_ptr, dev_ptr + N, (int) 0);

// access device memory through device_ptr
dev_ptr[0] = 1;

// free memory
cudaFree(raw_ptr);
```

Thrust - Algorithms

- ▶ Standard algorithms
 - Reductions
 - Transformations
 - Prefix sums
 - Sorting
- ▶ Many have straight analog to STL
- ▶ You can use your own user defined types

Thrust - Algorithms

‣ Reductions

```
#include <thrust/reduce.h>

// declare storage
device_vector<int> i_vec = ...
device_vector<float> f_vec = ...

// sum of integers (equivalent calls)
reduce(i_vec.begin(), i_vec.end());
reduce(i_vec.begin(), i_vec.end(), 0, plus<int>());

// sum of floats (equivalent calls)
reduce(f_vec.begin(), f_vec.end());
reduce(f_vec.begin(), f_vec.end(), 0.0f, plus<float>());

// maximum of integers
reduce(i_vec.begin(), i_vec.end(), 0, maximum<int>());
```

Thrust - Algorithms

- Thrust comes with lots of important built in transformations

```
#include<thrust/device_vector.h>
#include<thrust/transform.h>
#include<thrust/sequence.h>
#include<thrust/copy.h>
#include<thrust/fill.h>
#include<thrust/replace.h>
#include<thrust/functional.h>
#include<iostream>

//allocate three device_vectors with 10 elements
thrust::device_vector<int>X(10);
thrust::device_vector<int> Y(10);
thrust::device_vector<int>Z(10);
//initialize X to 0, 1, 2, 3,....
thrust::sequence(X.begin(),X.end());
//compute Y=-X
thrust::transform(X.begin(),X.end(),Y.begin(),thrust::negate<int>());
//fill Z with two s
thrust::fill(Z.begin(),Z.end(),2);
//compute Y = X mod 2
thrust::transform(X.begin(),X.end(),Z.begin(),Y.begin(),thrust::modulus<int>());
//replace all the ones in Y with tens
thrust::replace(Y.begin(),Y.end(),1,10);
//print Y
thrust::copy(Y.begin(),Y.end(),std::ostream_iterator<int>(std::cout,"\\n"));
return 0;
```

Thrust - Algorithms

▶ Prefix sums

```
#include<thrust/scan.h>
int data [6] = {1,0,2,2,1,3};
thrust::inclusive_scan(data, data + 6, data); //in-place scan
//data is now {1,1,3,5,6,9}

        data[2] = data[0] + data[1] + data[2]
```

```
#include<thrust/scan.h>
int data [6] = {1,0,2,2,1,3};
thrust::exclusive_scan(data, data + 6, data); //in-place scan
//data is now {0,1,1,3,5,6}

        data[2] = data[0] + data[1]
```

Thrust - Algorithms

‣ Sorting

```
#include<thrust/sort.h>
. . .
const int N=6;
int A [N] = {1,4,2,8,5,7};
thrust::sort(A,A+N);
// A is now {1,2,4,5,7,8}
```

Thrust - Fancy iterators

- ▶ Behave like “normal” iterators
 - Also they can be seen as pointers
- ▶ Examples
 - `constant_iterator`
 - `counting_iterator`
 - `transform_iterator`
 - `permutation_iterator`
 - `zip_iterator`

Thrust - Fancy iterators

- ▶ constant_iterator
 - Mimics an infinite array with constant values

```
// create iterators
constant_iterator<int> begin(10);
constant_iterator<int> end = begin + 3;

begin[0]    // returns 10
begin[1]    // returns 10
begin[100]  // returns 10

// sum of (begin, end)
reduce(begin, end);    // returns 30 (i.e. 3 * 10)
```

Thrust - Fancy iterators

- ▶ counting_iterator
 - Mimics an infinite array with sequential values

```
// create iterators
counting_iterator<int> begin(10);
counting_iterator<int> end = begin + 3;

begin[0]    // returns 10
begin[1]    // returns 11
begin[100]  // returns 110

// sum of (begin, end)
reduce(begin, end);    // returns 33 (i.e. 10 +
11 + 12)
```

Thrust - Fancy iterators

- ▶ `transform_iterator`
 - Allows us to fuse separate algorithms into one

```
// initialize vector
device_vector<int> vec(3);
vec[0] = 10; vec[1] = 20; vec[2] = 30;

// create iterator (type omitted)
begin = make_transform_iterator(vec.begin(), negate<int>());
end   = make_transform_iterator(vec.end(),   negate<int>());

begin[0] // returns -10
begin[1] // returns -20
begin[2] // returns -30

// sum of [begin, end)
reduce(begin, end); // returns -60 (i.e. -10 + -20 + -30)
```

Thrust - Fancy iterators

- ▶ `permutation_iterator`

- Allows to fuse gather and scatter operations

```
#include<thrust/iterator/permutation_iterator.h>
...
//gather locations
thrust::device_vector<int> map(4);
map[0] = 3;
map[1] = 1;
map[2] = 0;
map[3] = 5;

//array to gather from
thrust::device_vector<int> source(6);
source[0] = 10;
source[1] = 20;
source[2] = 30;
source[3] = 40;
source[4] = 50;
source[5] = 60;

//fuse gather with reduction: sum = source[map[0]] + source[map[1]]+...
int sum = thrust::reduce(thrust::make_permutation_iterator(source.begin(),map.begin()),
thrust::make_permutation_iterator(source.begin(),map.end()));
```

Thrust - Fancy iterators

- ▶ **zip_iterator**

- Looks like an array of structs

- Stored in structure of arrays

```
// initialize vectors
device_vector<int> A(3);
device_vector<char> B(3);
A[0] = 10; A[1] = 20; A[2] = 30;
B[0] = 'x'; B[1] = 'y'; B[2] = 'z';

// create iterator (type omitted)
begin = make_zip_iterator(make_tuple(A.begin(), B.begin()));
end   = make_zip_iterator(make_tuple(A.end(),   B.end()));

begin[0] // returns tuple(10, 'x')
begin[1] // returns tuple(20, 'y')
begin[2] // returns tuple(30, 'z')

// maximum of [begin, end)
maximum< tuple<int,char> > binary_op;
reduce(begin, end, begin[0], binary_op); // returns tuple(30, 'z')
```

Thrust - Best practices

- ▶ Fusion
 - Combine related operations together
- ▶ Structure of Arrays
 - Ensure memory coalescing
- ▶ Implicit Sequences
 - Eliminate memory accesses

Thrust - Fusion

- ▶ Combine related operations together
 - Conserves memory bandwidth
- ▶ Example: norm of a vector
 - Square each element
 - Compute the sum of squares and take sqrt()

Thrust - Fusion

› Unoptimized example

```
// define transformation f(x) -> x^2
struct square
{
    __host__ __device__
    float operator()(float x)
    {
        return x * x;
    }
};

float snrm2_slow(device_vector<float>& x)
{
    // without fusion
    device_vector<float> temp(x.size());
    transform(x.begin(), x.end(), temp.begin(), square());

    return sqrt( reduce(temp.begin(), temp.end()) );
}
```

Thrust - Fusion

- ▶ Optimized implementation (3.8x)

```
// define transformation f(x) -> x^2
struct square
{
    __host__ __device__
    float operator()(float x)
    {
        return x * x;
    }
};

float snrm2_fast(device_vector<float>& x)
{
    // with fusion
    return sqrt( transform_reduce(x.begin(), x.end(), square(), 0.0f, plus<float>()));
}
```

Thrust - Structure of Arrays (SoA)

- ▶ Array of structures (AoS)
 - Often does not obey coalescing rules

```
device_vector<float3>
```
- ▶ Structure of arrays (SoA)
 - Obeys coalescing rules
 - Components stored in separate arrays

```
device_vector<float> x,y,z;
```
- ▶ Example: rotate 3D vectors

Thrust - Structure of Arrays (SoA)

```
struct rotate_float3
{
    __host__ __device__
    float3 operator()(float3 v)
    {
        float x = v.x;
        float y = v.y;
        float z = v.z;
        float rx = 0.36f*x + 0.48f*y + -0.80f*z;
        float ry = -0.80f*x + 0.60f*y + 0.00f*z;
        float rz = 0.48f*x + 0.64f*y + 0.60f*z;
        return make_float3(rx, ry, rz);
    }
};

...
device_vector<float3> vec(N);
transform(vec.begin(), vec.end, vec.begin(), rotate_float3());
```

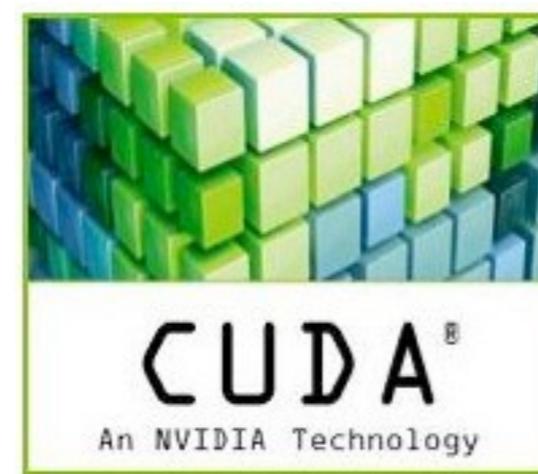
Thrust - Structure of Arrays (SoA)

```
struct rotate_tuple
{
    __host__ __device__
    tuple<float, float, float> operator()(tuple<float, float, float> v)
    {
        float x = get<0>(v);
        float y = get<1>(v);
        float z = get<2>(v);
        float rx = 0.36f*x + 0.48f*y + -0.80f*z;
        float ry = -0.80f*x + 0.60f*y + 0.00f*z;
        float rz = 0.48f*x + 0.64f*y + 0.60f*z;
        return make_tuple(rx, ry, rz);
    }
};

...
device_vector<float> x(N), y(N), z(N);
transform(make_zip_iterator(make_tuple(x.begin(), y.begin(), z.begin())),
         make_zip_iterator(make_tuple(x.end(), y.end(), z.end())),
         make_zip_iterator(make_tuple(x.begin(), y.begin(), z.begin())),
         rotate_tuple());
```

PyCUDA

- Python + CUDA = PyCUDA
- Python: scripting language → easy to code, but slow
- CUDA → difficult to code, but fast!
- PyCUDA wants to take the best of both worlds
- <http://mathematician.de/software/pycuda>



PyCUDA

- ▶ Scripting language
 - High level programming language that is interpreted by another program at runtime rather than compiled
 - Advantages: ease on programmer
 - Disadvantages: slow (specially inner loops)
- ▶ PyCUDA
 - CUDA codes does not need to be a constant at compile time
 - Machine generated code: automatic manage of resources

PyCUDA

```
import pycuda.autoinit
import pycuda.driver as drv
import numpy

from pycuda.compiler import SourceModule
mod = SourceModule("""
__global__ void multiply_them(float *dest, float *a, float *b)
{
    const int i = threadIdx.x;
    dest[i] = a[i] * b[i];
}
""")

multiply_them = mod.get_function("multiply_them")

a = numpy.random.randn(400).astype(numpy.float32)
b = numpy.random.randn(400).astype(numpy.float32)

dest = numpy.zeros_like(a)
multiply_them(
    drv.Out(dest), drv.In(a), drv.In(b),
    block=(400,1,1), grid=(1,1))

print dest-a*b
```

PyCUDA

▶ Transferring data

```
import numpy
a = a.astype(numpy.float32)
a_gpu = cuda.mem_alloc(a.nbytes)

cuda.memcpy_htod(a_gpu, a)
```

▶ Executing a kernel

```
from pycuda.compiler import SourceModule
mod = SourceModule("""
__global__ void doublify(float *a)
{
    int idx = threadIdx.x + threadIdx.y*4;
    a[idx] *= 2;
}
""")
... # Allocate, generate and transfer
func = mod.get_function("doublify")
func(a_gpu, block=(4,4,1))
```

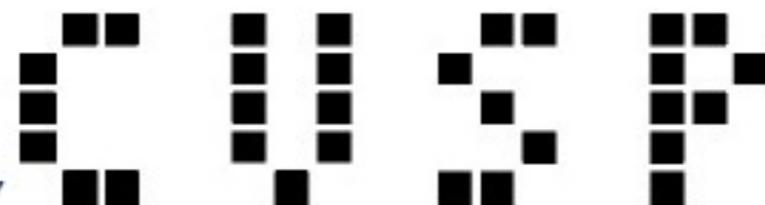
```
a_doubled = numpy.empty_like(a)
cuda.memcpy_dtoh(a_doubled, a_gpu)
print a_doubled
print a
```

CUFFT

- ▶ CUFFT: CUDA library for FFTs on the GPU
- ▶ Supported by NVIDIA
- ▶ Features:
 - 1D, 2D, 3D transforms for complex and real data
 - Batch execution for multiple transforms
 - Up to 128 million elements (limited by memory)
 - In-place or out-of-place transforms
 - Double precision on GT200 or later
 - Allows streamed execution: simultaneous computation and data movement

Cusp

- ▶ Library for sparse linear algebra
- ▶ <http://code.google.com/p/cusp-library/>



Project Information

Starred by 72 users
Activity Medium
[Project feeds](#)

Code license
[Apache License 2.0](#)

Labels
Sparse, Iterative, CUDA, GPU, ConjugateGradient, Graph, Matrix, SpMV

Members
wrb...@gmail.com, mgerl...@gmail.com, 5 committers

Featured

Downloads
[cusp-v0.2.0.zip](#), [examples-v0.2.zip](#), [Show all](#)

Wiki pages
[Documentation](#), [Frequently Asked Questions](#), [QuickStartGuide](#), [Show all](#)

Links

External links
[Threads](#), [CUDA](#)

Groups
[Cusp User Discussion List](#)

What is Cusp?

C U S P

Cusp is a library for **sparse linear algebra** and **graph computations** on CUDA. Cusp provides a flexible, high-level interface for manipulating sparse matrices and solving sparse linear systems. [Get Started](#) with Cusp today!

News

- Cusp v0.2.0 has been [released!](#) See [CHANGELOG](#) for release information.
- Cusp v0.1.2 has been [released!](#) v0.1.2 contains compatibility fixes for Thrust v1.3.0.
- Cusp v0.1.1 has been [released!](#) v0.1.1 contains compatibility fixes for CUDA 5.1.
- Cusp v0.1.0 has been [released!](#).
- Added [QuickStartGuide](#) page.

Examples

The following example loads a matrix from disk, transparently converts the matrix to the highly-efficient HYB format, and transfers the matrix to the GPU device. The linear system $A \cdot x = b$ is then solved on the device using the Conjugate Gradient method. A more detailed version of this example is also available.

```
#include <cusp/hyb_matrix.h>
#include <cusp/io/matrix_market.h>
#include <cusp/krylov/cg.h>

int main(void)
{
    // create an empty sparse matrix structure (HYB format)
    cusp::hyb_matrix<float, cusp::device_memory> A;

    // load a matrix stored in MatrixMarket format
    cusp::io::read_matrix_market_file(A, "3pt_10x10.mtx");

    // allocate storage for solution (x) and right hand side (b)
    cusp::array1d<float, cusp::device_memory> x(A.num_rows, 0);
    cusp::array1d<float, cusp::device_memory> b(A.num_rows, 1);

    // solve the linear system A * x = b with the Conjugate Gradient method
    cusp::krylov::cg(A, x, b);

    return 0;
}
```